

Implementation of Multi Layer Perceptron Algorithm from scratch

```
import torch
print(torch.__version__)
import torch.nn as nn
import numpy as np

import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

import matplotlib.pyplot as plt

2.8.0+cu126

import torchvision
import torchvision.datasets as datasets
from torchvision import transforms

transform = transforms.Compose([
    transforms.ToTensor(),                # Converts to float tensor
    [0,1] and adds channel dimension
    transforms.Lambda(lambda x: x.view(-1)) # Flatten tensor of
dimension 28x28 to vector of size 784
])

mnist_trainset = datasets.MNIST(root='./data', train=True,
download=True, transform=transform)
mnist_testset = datasets.MNIST(root='./data', train=False,
download=True, transform=transform)

print(f"The number of images in training dataset are
{len(mnist_trainset)}")
print(f"The number of images in test dataset are
{len(mnist_testset)}")

print(mnist_trainset.data.shape)
print(mnist_testset.data.shape)

The number of images in training dataset are 60000
The number of images in test dataset are 10000
torch.Size([60000, 28, 28])
torch.Size([10000, 28, 28])

print(mnist_trainset[0][0].shape)

torch.Size([784])
```

```
batch_size = 64
train_loader = DataLoader(mnist_trainset, batch_size=batch_size,
                           shuffle=True)
test_loader = DataLoader(mnist_testset, batch_size=batch_size,
                          shuffle=False)
```

#Question 1 Part A:

#Implementing Multi Layer Perceptron using Sigmoid Activation function

Model architecture:

- Input layer (784 neurons)
- Hidden layers with 500, 250, and 100 neurons using Sigmoid activations
- Output layer with 10 neurons (linear, raw logits)
- Use of softmax for probability assignment
- Glorot uniform weight initialization and zeros biases per assignment

```
class MLP(nn.Module):
    def __init__(self, input_dim = 784, hidden_dims = [500, 250, 100],
                  output_dim = 10 ):

        layer_sizes = [input_dim] + hidden_dims + [output_dim]

        # Initialize weights and biases (Glorot initialization, zero
        biases)
        self.weights = []
        self.biases = []
        for i in range(len(layer_sizes) - 1):
            ip_layer_size = layer_sizes[i]
            op_layer_size = layer_sizes[i + 1]
            limit = np.sqrt(6 / (ip_layer_size + op_layer_size))
            W = np.random.uniform(-limit, limit, size=(ip_layer_size,
        op_layer_size))
            b = np.zeros((1, op_layer_size))
            self.weights.append(W)
            self.biases.append(b)

        def sigmoid(self, x):
            return 1 / (1 + np.exp(-x))

        def sigmoid_derivative(self, x):
            return x * (1 - x)

        def softmax(self, x):
            exp_x = np.exp(x - np.max(x))
            return exp_x / exp_x.sum(axis=1, keepdims=True)

        def forward(self, X):
            # Storing the inputs in arrays for backpropogation
```

```

self.inputs = []
self.z_values = []

# Input layer to first hidden layer
z1 = X @ self.weights[0] + self.biases[0]
a1 = self.sigmoid(z1)

# First hidden layer to second hidden layer
z2 = a1 @ self.weights[1] + self.biases[1]
a2 = self.sigmoid(z2)

# Second hidden layer to third one
z3 = a2 @ self.weights[2] + self.biases[2]
a3 = self.sigmoid(z3)

# Third hidden layer to raw output
z4 = a3 @ self.weights[3] + self.biases[3]

# Outut layer
output = self.softmax(z4)

self.inputs = [X, a1, a2, a3]
self.z_values = [z1, z2, z3, z4]

return output

def compute_loss(self, y_true, y_pred):
    m = y_true.shape[0] #number of samples in the batch
    eps = 1e-9 #value to avoid log(0)
    log_preds = np.log(y_pred + eps) #log of prediction
    loss = -np.sum(y_true * log_preds) / m
    return loss

def backward(self, y_true, y_pred, learning_rate):
    m = y_true.shape[0]

    delta4 = (y_pred - y_true) / m # output error

    dW4 = self.inputs[3].T @ delta4
    db4 = np.sum(delta4, axis=0, keepdims=True)

    delta3 = delta4 @ self.weights[3].T *
self.sigmoid_derivative(self.inputs[3])
    dW3 = self.inputs[2].T @ delta3
    db3 = np.sum(delta3, axis=0, keepdims=True)

    delta2 = delta3 @ self.weights[2].T *
self.sigmoid_derivative(self.inputs[2])
    dW2 = self.inputs[1].T @ delta2
    db2 = np.sum(delta2, axis=0, keepdims=True)

```

```

        delta1 = delta2 @ self.weights[1].T *
self.sigmoid_derivative(self.inputs[1])
        dW1 = self.inputs[0].T @ delta1
        db1 = np.sum(delta1, axis=0, keepdims=True)

        # Update weights and biases
        self.weights[3] -= learning_rate * dW4
        self.biases[3] -= learning_rate * db4

        self.weights[2] -= learning_rate * dW3
        self.biases[2] -= learning_rate * db3

        self.weights[1] -= learning_rate * dW2
        self.biases[1] -= learning_rate * db2

        self.weights[0] -= learning_rate * dW1
        self.biases[0] -= learning_rate * db1

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

cpu

def one_hot_encode(labels, num_classes=10):
    return np.eye(num_classes)[labels]

def train(model, train_loader, epochs=15, learning_rate=0.01,
log_interval=200):
    loss_history = []      # epoch losses
    iter_losses = []      # losses every 200 updates
    iters = []            # iteration numbers

    iteration = 0

    for epoch in range(epochs):
        epoch_loss = 0.0
        batch_count = 0

        for images, labels in train_loader:
            iteration += 1
            X = images.numpy()
            y_indices = labels.numpy()
            y = one_hot_encode(y_indices)

            output = model.forward(X)
            loss = model.compute_loss(y, output)

            epoch_loss += float(loss)
            batch_count += 1
            model.backward(y, output, learning_rate)

```

```

        # log every 200 updates
        if iteration % log_interval == 0:
            avg_loss_so_far = epoch_loss / batch_count
            iter_losses.append(avg_loss_so_far)
            iters.append(iteration)

    avg_loss = epoch_loss / batch_count
    loss_history.append(avg_loss)
    print(f"Epoch {epoch+1}/{epochs}, Loss: {avg_loss:.4f}")

    # plot epoch curve so far
    plt.figure(figsize=(6,4))
    plt.plot(range(1, len(loss_history)+1), loss_history,
marker='o')
    plt.title("Training Loss per Epoch")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.grid(True)
    plt.show()

    # plot iteration-based curve once at the end
    plt.figure(figsize=(6,4))
    plt.plot(iters, iter_losses, marker='o')
    plt.title(f"Training Loss (every {log_interval} updates)")
    plt.xlabel("Iterations (batch updates)")
    plt.ylabel("Loss")
    plt.grid(True)
    plt.show()

    return loss_history, iters, iter_losses

from sklearn.metrics import confusion_matrix, classification_report

def predict(model, X):
    probs = model.forward(X)
    preds = np.argmax(probs, axis=1)
    return preds

def evaluate(model, X, y_true):
    preds = predict(model, X)
    accuracy = np.mean(preds == y_true)
    print(f"Accuracy: {accuracy*100:.4f}%")

    print("Confusion Matrix:")
    print(confusion_matrix(y_true, preds))

    print("\nClassification Report:")
    print(classification_report(y_true, preds))

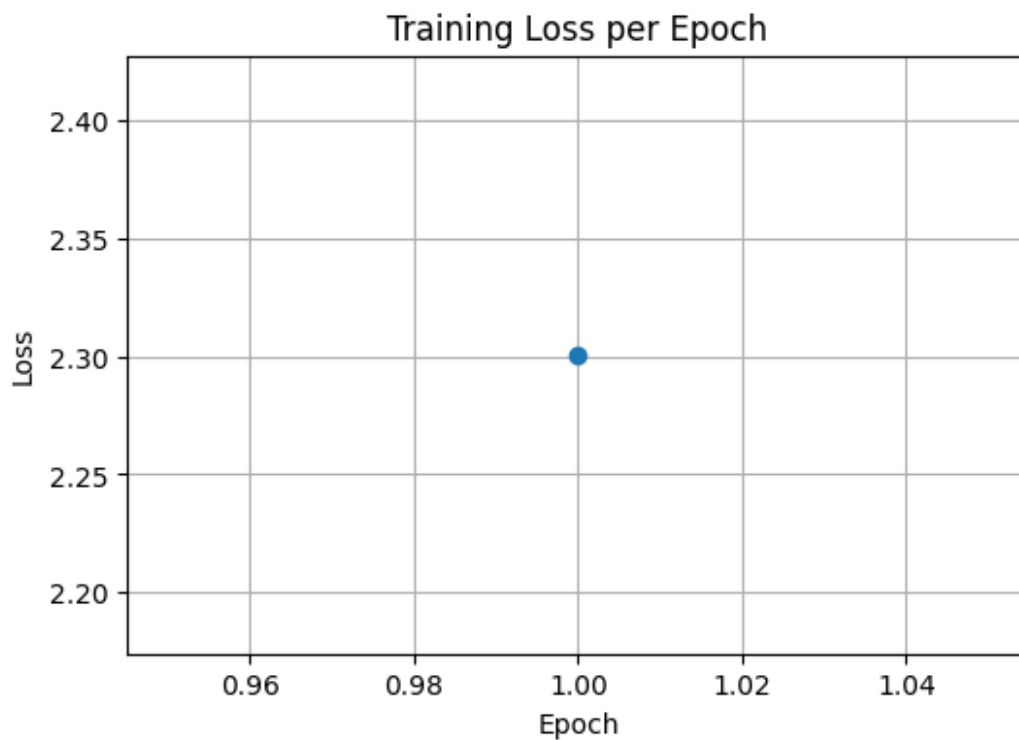
    return accuracy

```

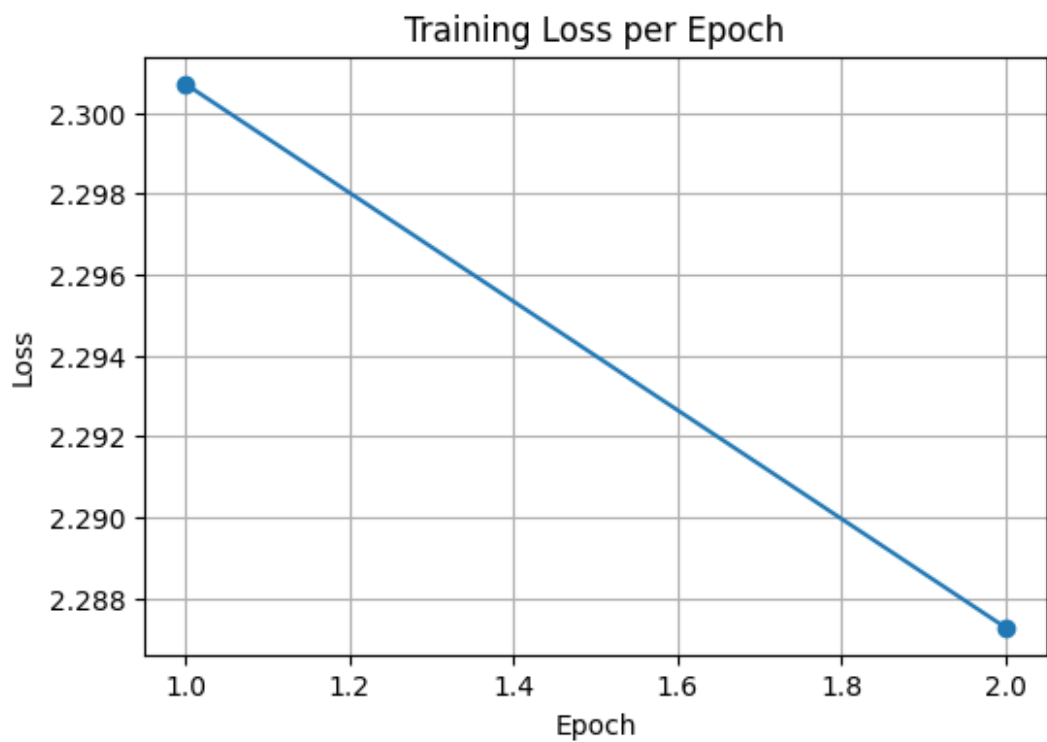
```
def plot_training_curves(train_losses):  
    plt.figure(figsize=(8,5))  
    plt.plot(range(1, len(train_losses)+1), train_losses, marker='o')  
    plt.title("Training Loss over Epochs")  
    plt.xlabel("Epoch")  
    plt.ylabel("Loss")  
    plt.grid(True)  
    plt.show()
```

```
model = MLP()  
loss_history, iter_num, iter_loss = train(model, train_loader, 15,  
0.01)
```

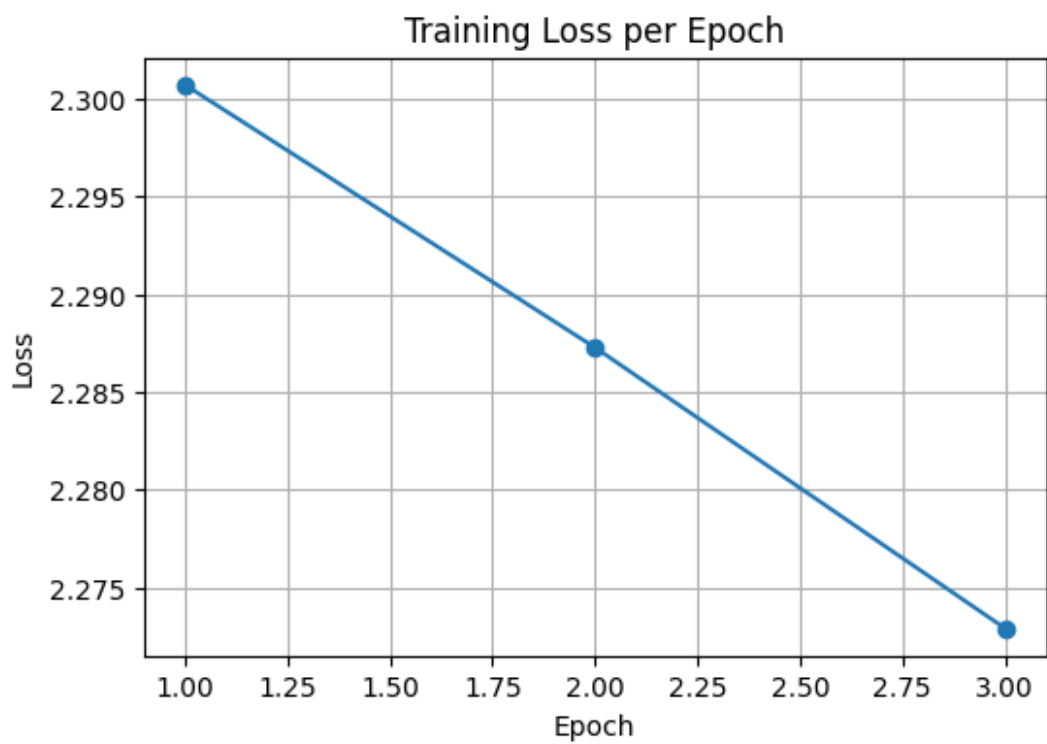
Epoch 1/15, Loss: 2.3007



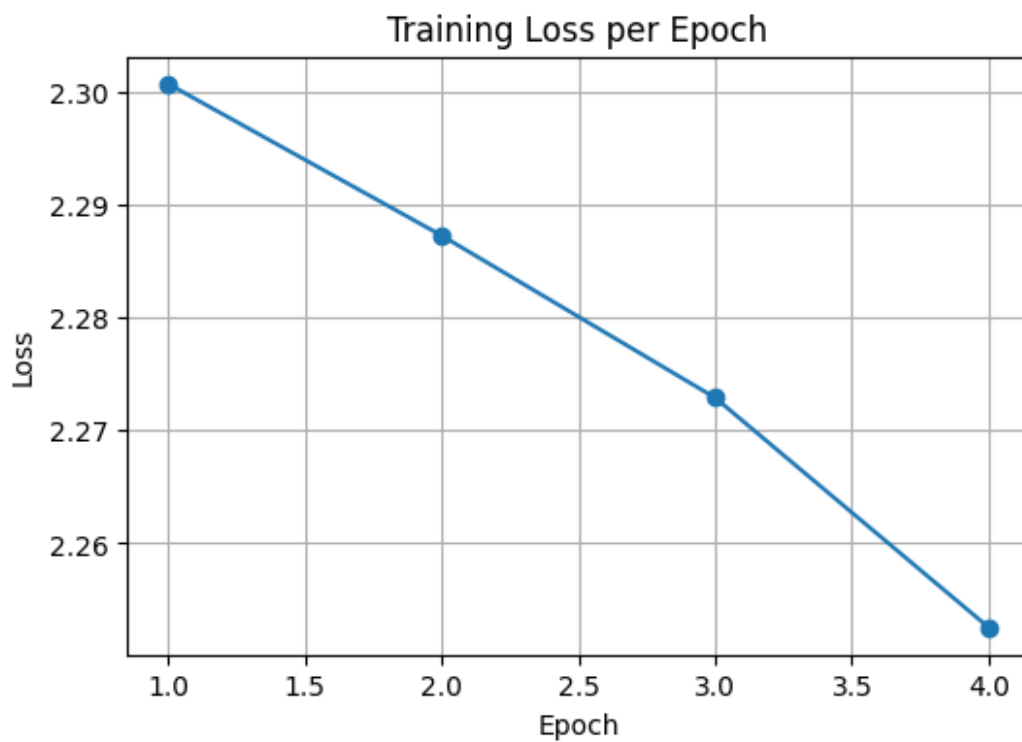
Epoch 2/15, Loss: 2.2873



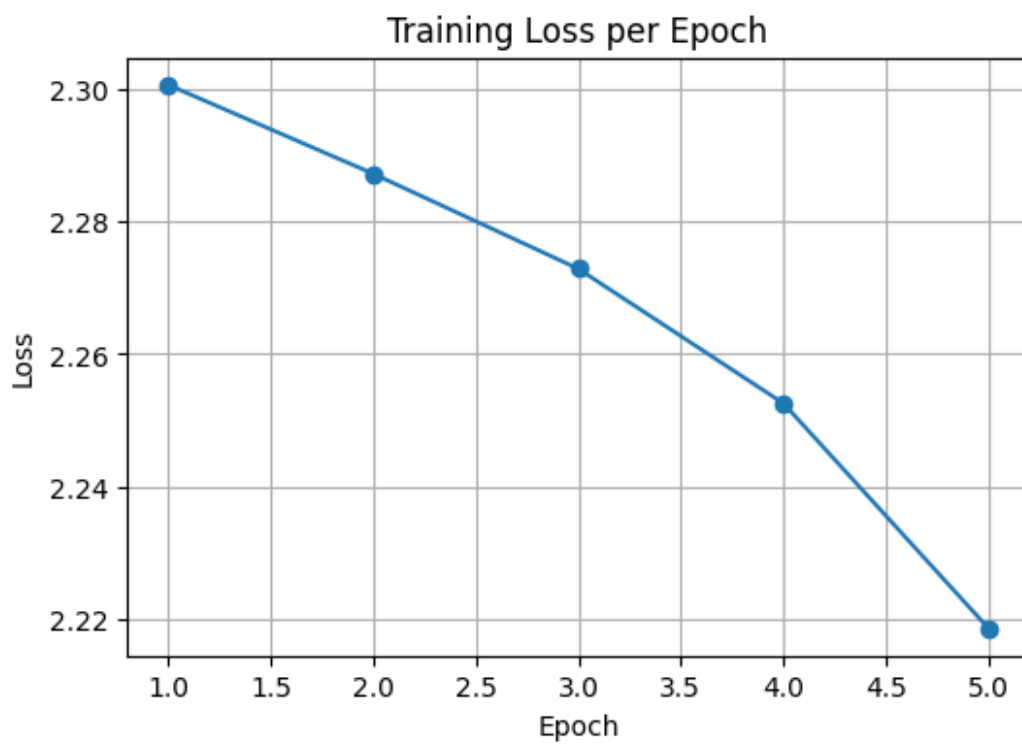
Epoch 3/15, Loss: 2.2729



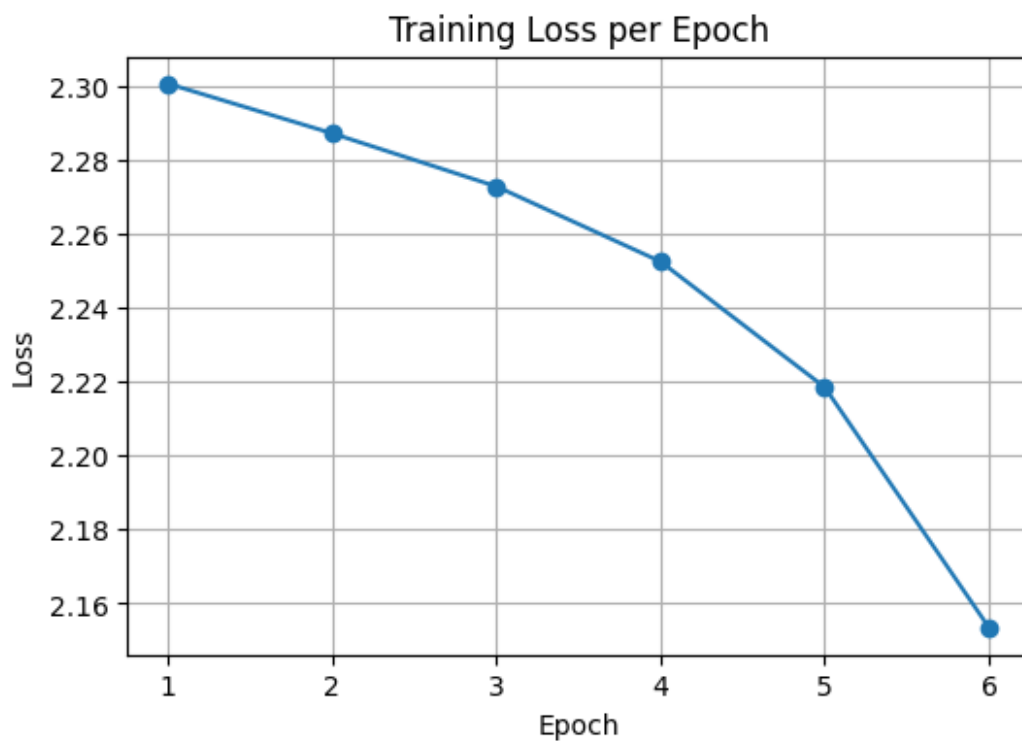
Epoch 4/15, Loss: 2.2526



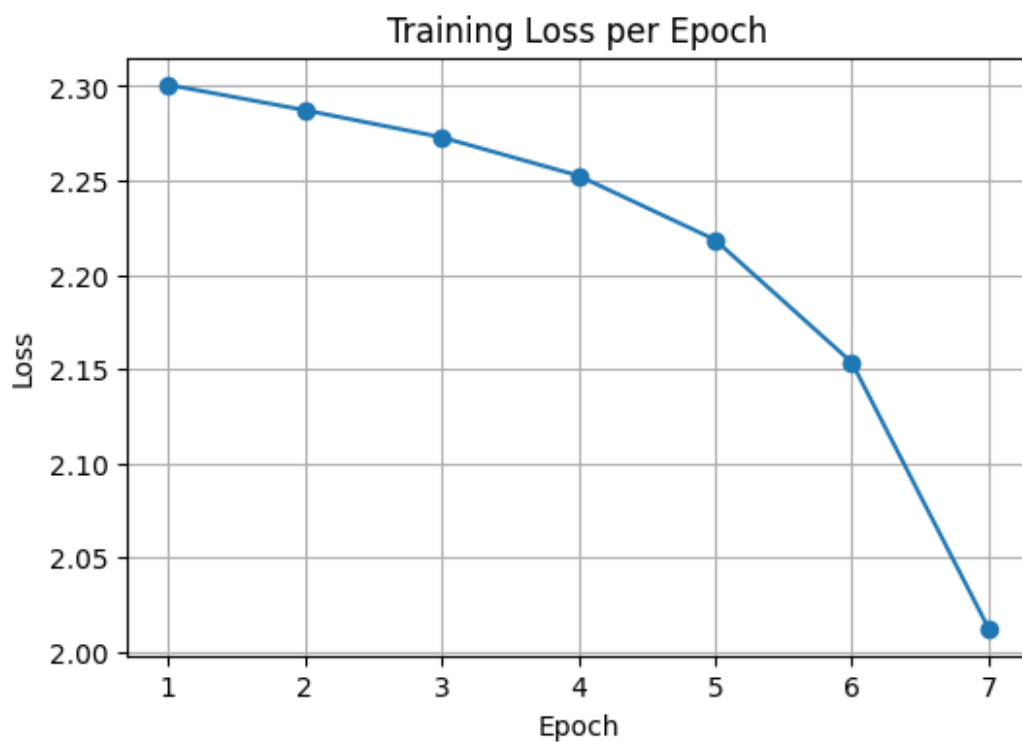
Epoch 5/15, Loss: 2.2186



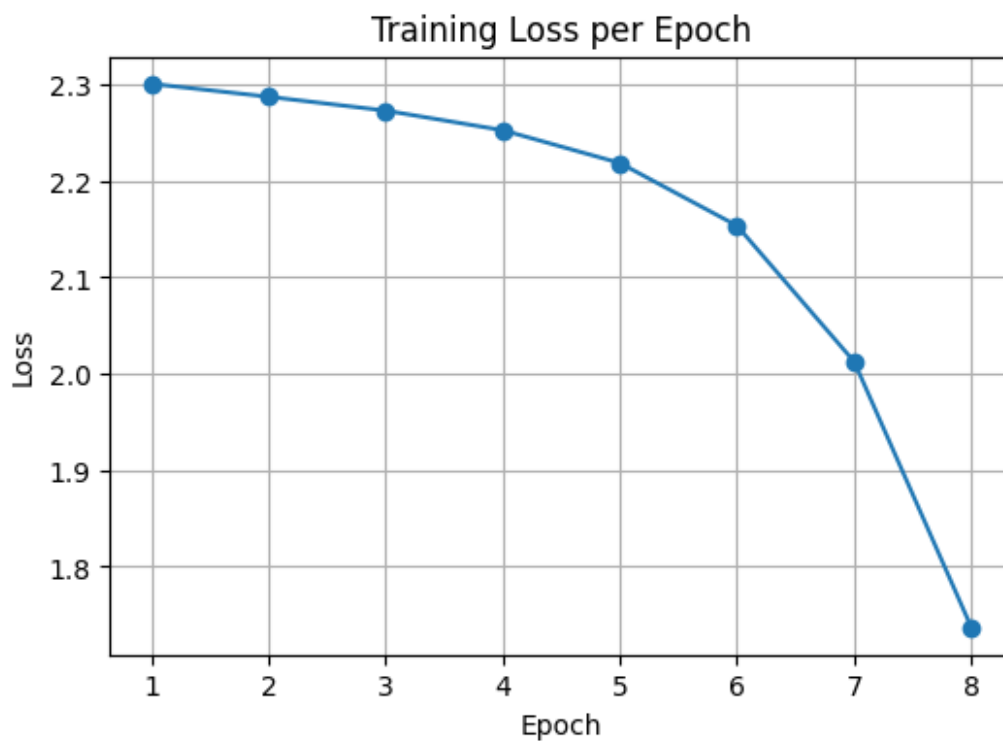
Epoch 6/15, Loss: 2.1535



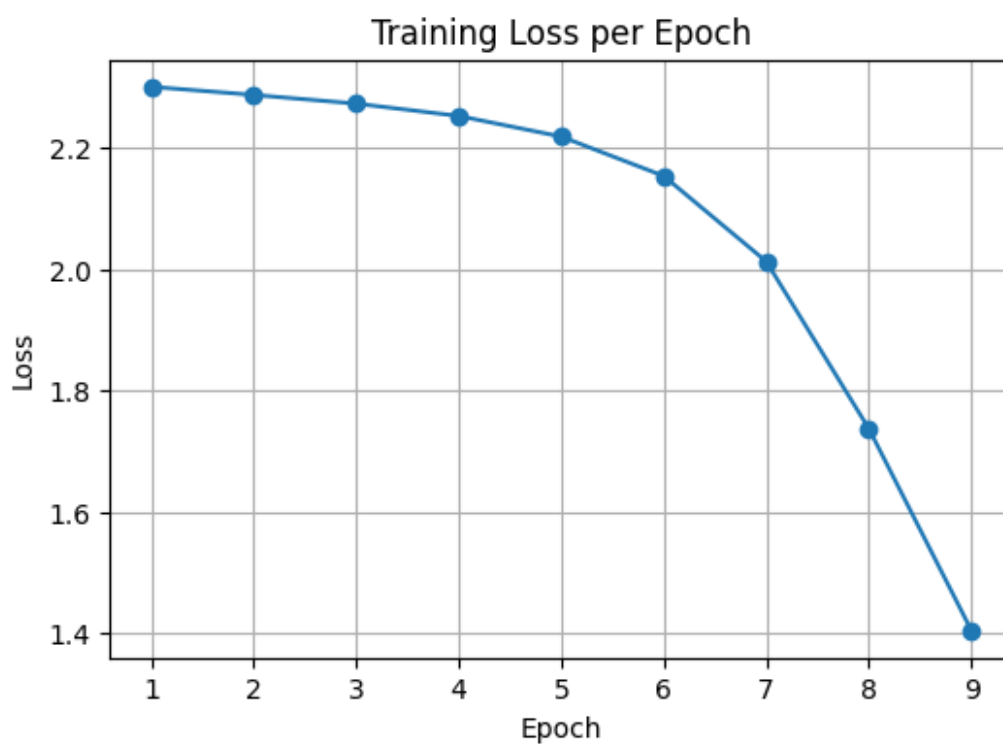
Epoch 7/15, Loss: 2.0126



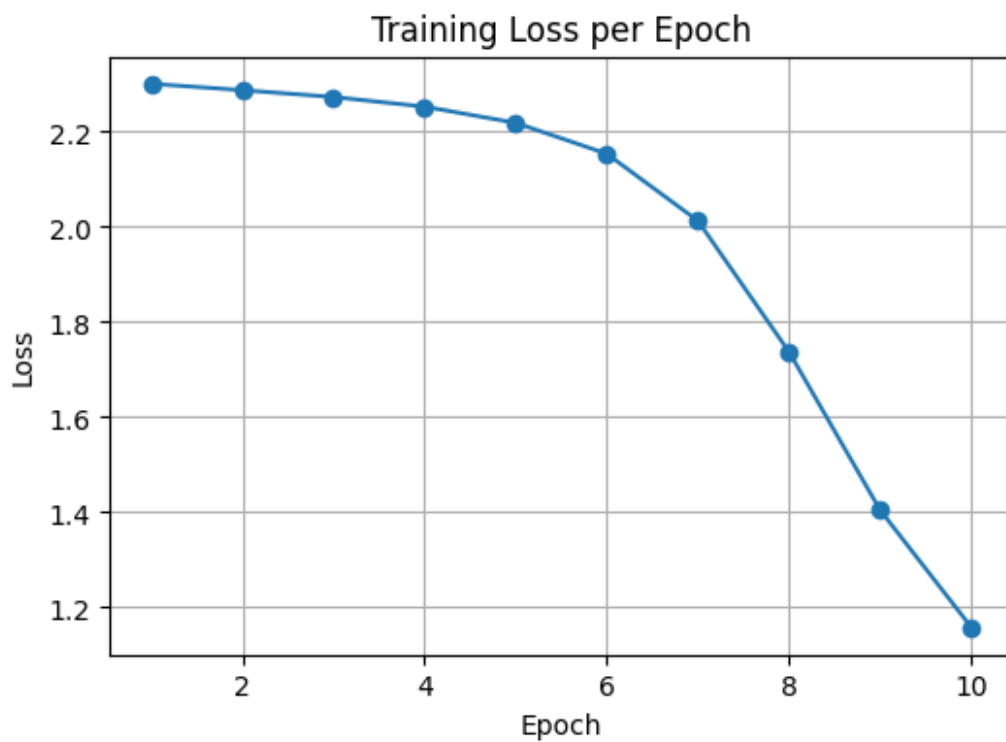
Epoch 8/15, Loss: 1.7371



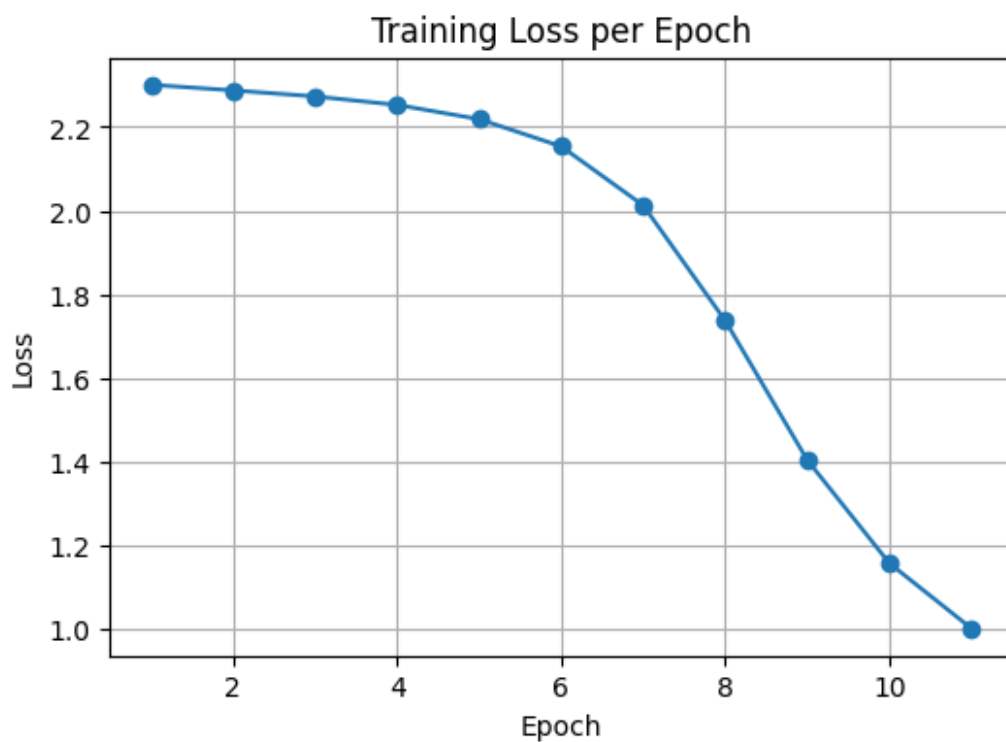
Epoch 9/15, Loss: 1.4044



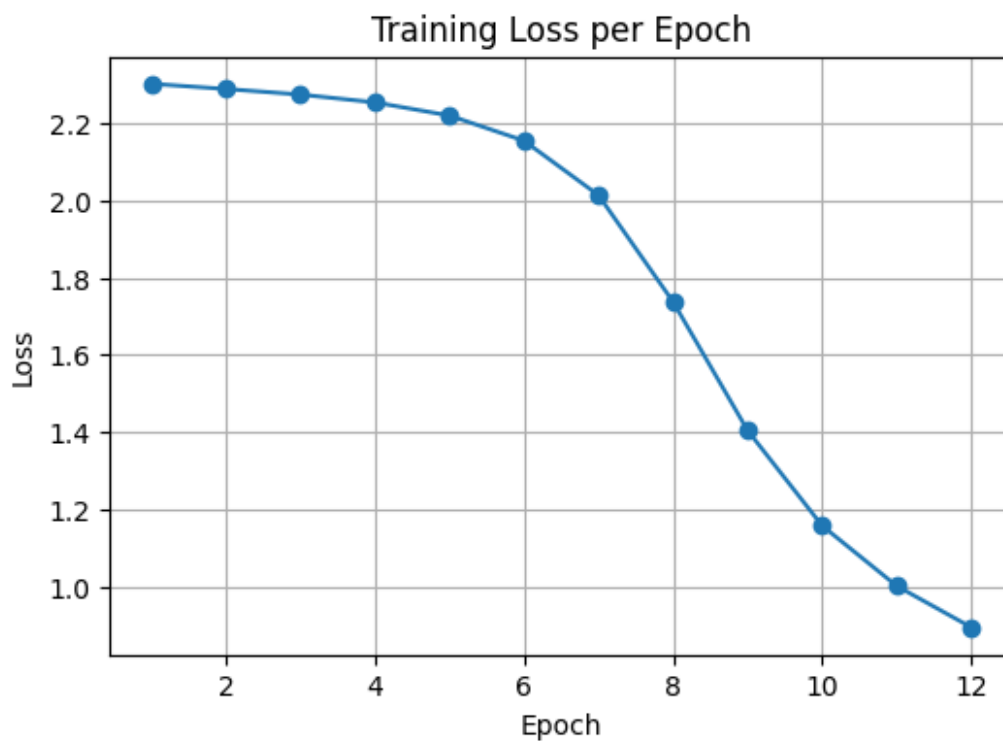
Epoch 10/15, Loss: 1.1582



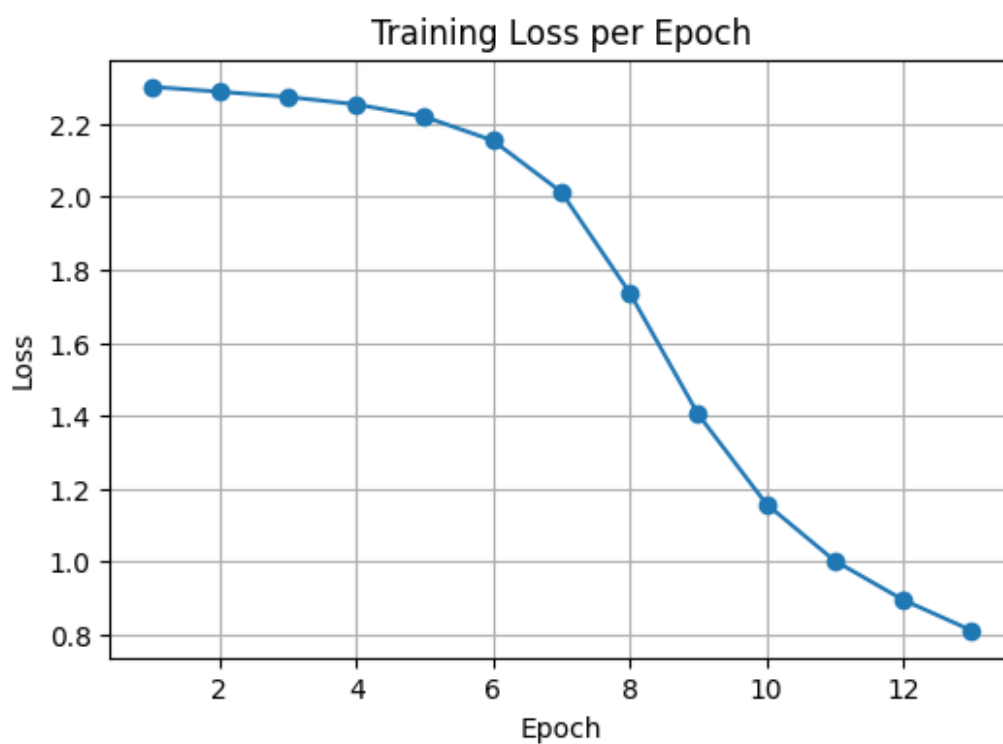
Epoch 11/15, Loss: 1.0025



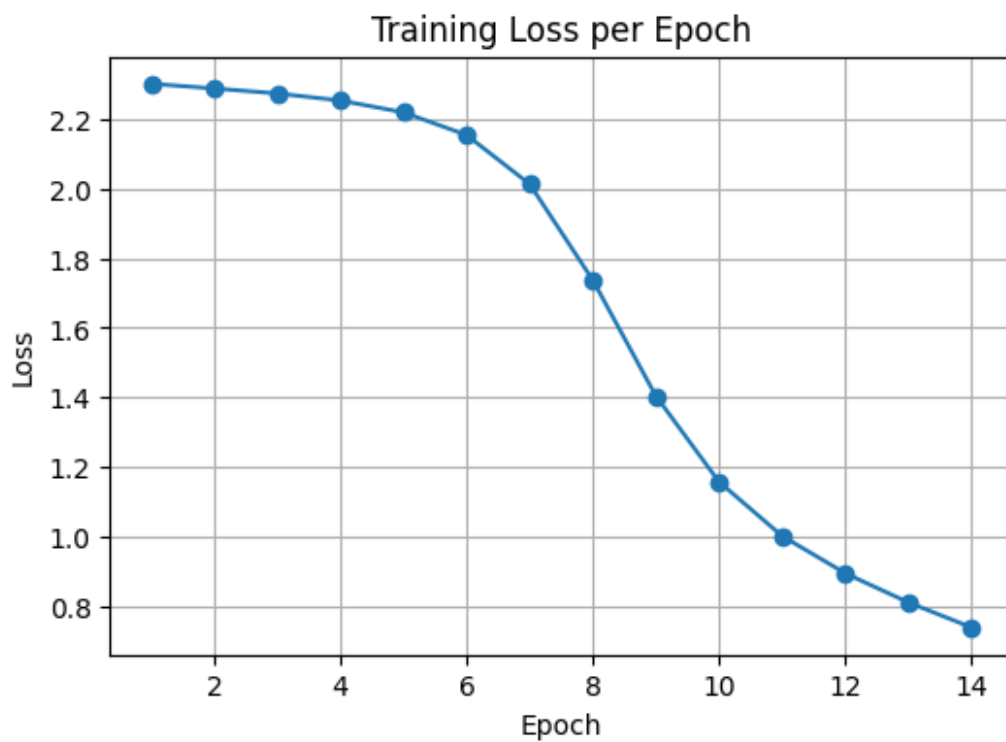
Epoch 12/15, Loss: 0.8956



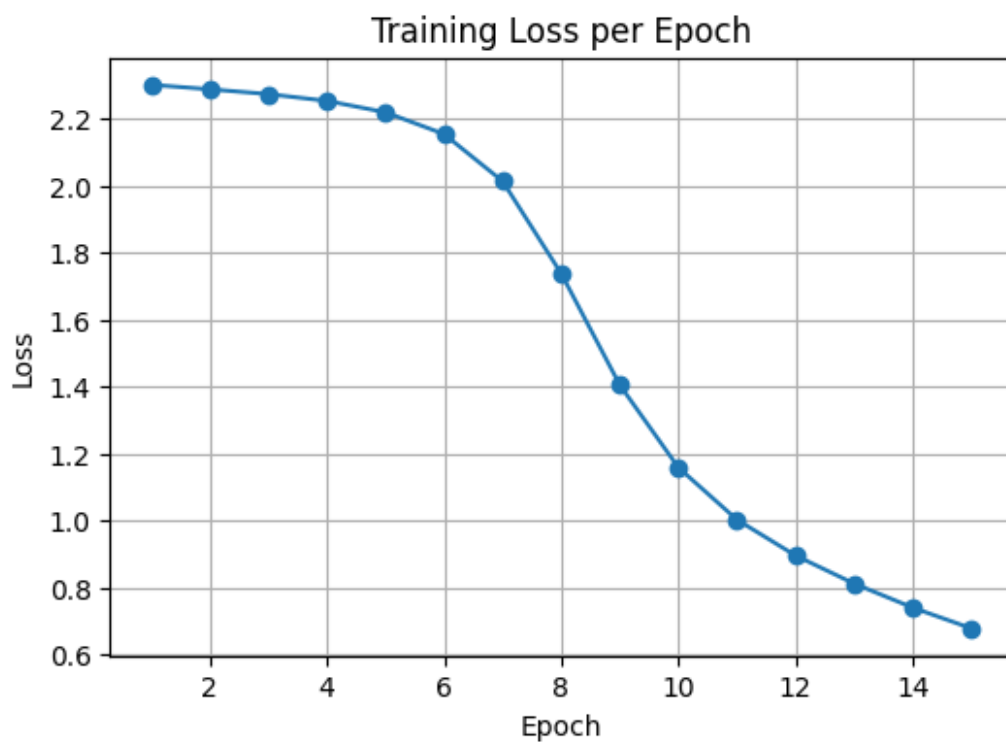
Epoch 13/15, Loss: 0.8119

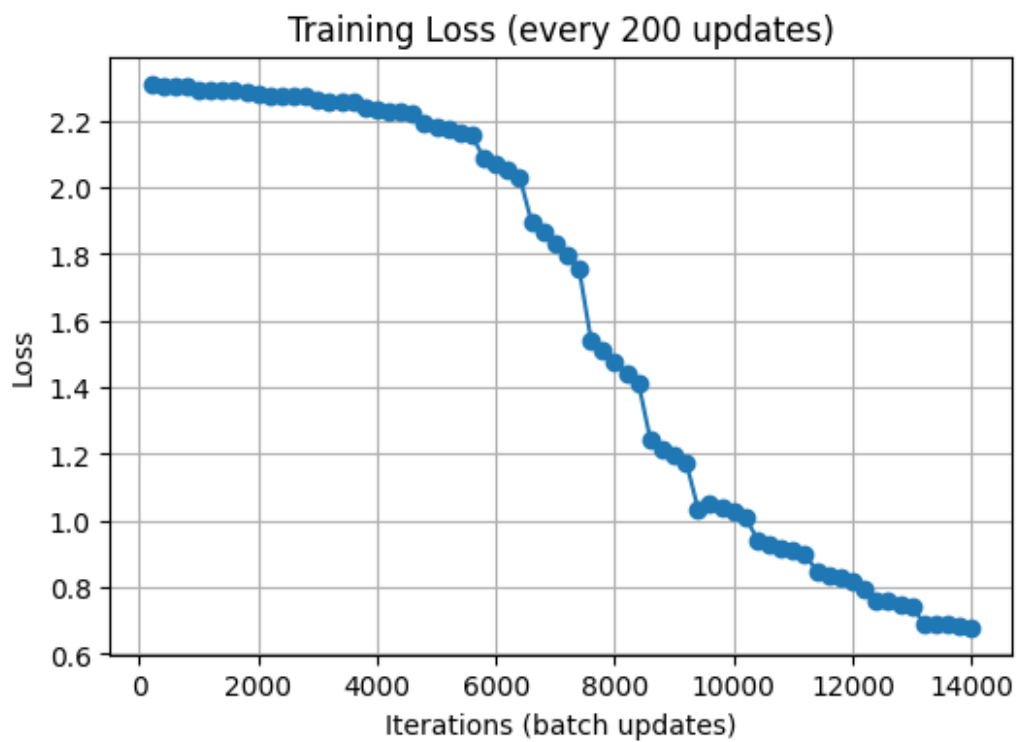


Epoch 14/15, Loss: 0.7401

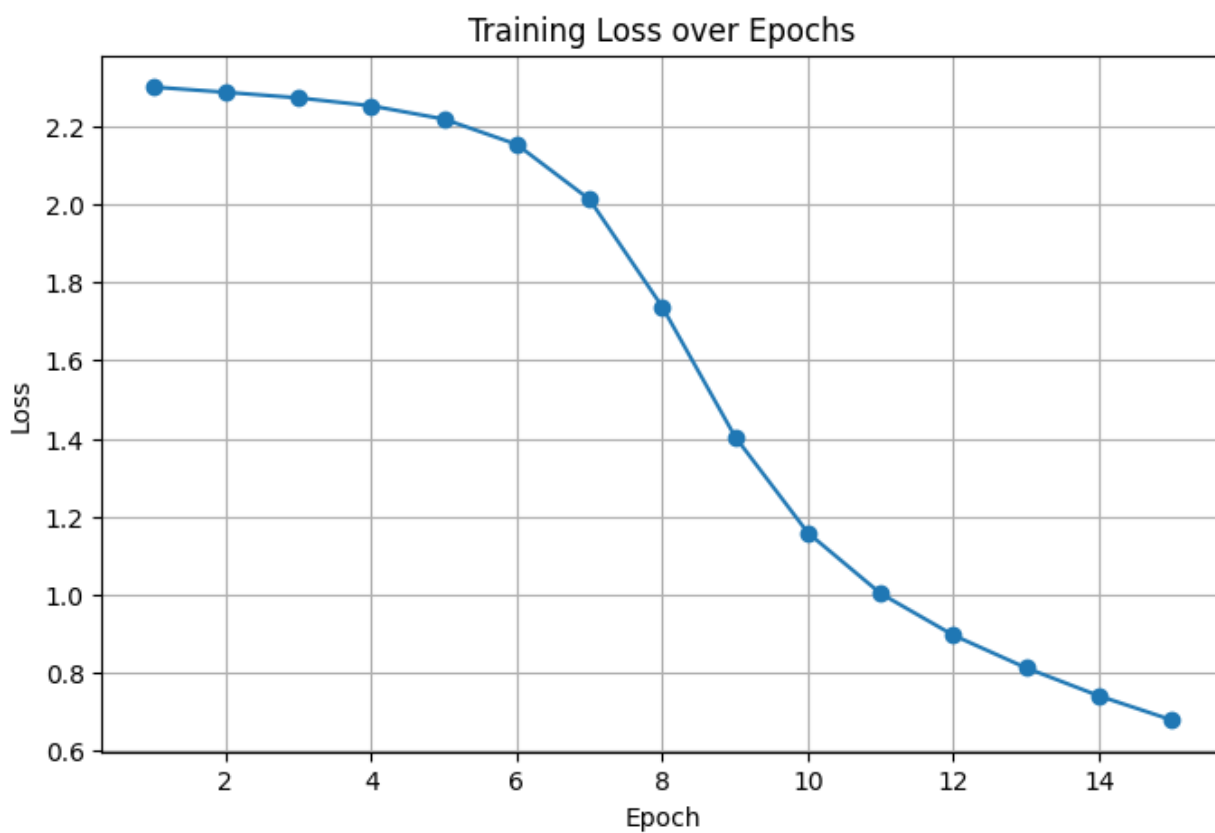


Epoch 15/15, Loss: 0.6783





```
plot_training_curves(loss_history)
```



```
# Prepare test set in numpy format
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("---- Test Set Evaluation ----")
evaluate(model, X_test_np, y_test_np)
```

```
---- Test Set Evaluation ----
```

```
Accuracy: 82.3200%
```

```
Confusion Matrix:
```

```
[ [ 944    0    2    1    3   19    9    1    1    0]
  [    0 1114    3    3    0    0    3    1   10    1]
  [   15   28  840   25   27    8   52   12   25    0]
  [    5    7   47  810    2   64    0   21   50    4]
  [    2   12    1    0  832    0   30    1    5   99]
  [   27    6   24   92   35  593   32   11   41   31]
  [   29    4   15    0   11   18  879    0    2    0]
  [    6   45   12    2   12    1    0  916    7   27]
  [   12   29   37   66   25   86   21   13  651   34]
  [   14    7    2    8  207   17    3   94    4  653]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.90	0.96	0.93	980
1	0.89	0.98	0.93	1135
2	0.85	0.81	0.83	1032
3	0.80	0.80	0.80	1010
4	0.72	0.85	0.78	982
5	0.74	0.66	0.70	892
6	0.85	0.92	0.88	958
7	0.86	0.89	0.87	1028
8	0.82	0.67	0.74	974
9	0.77	0.65	0.70	1009
accuracy			0.82	10000
macro avg	0.82	0.82	0.82	10000
weighted avg	0.82	0.82	0.82	10000

```
np.float64(0.8232)
```

Results for MLP using Sigmoid activation function

- Classification accuracy over test data = **82.30%**
- Training loss after 15 epochs = **0.6783**

#Question 1 Part B

#Implementing Multi Layer Perceptron Algorithm using ReLu activation function

Model architecture:

- Input layer (784 neurons)
- Hidden layers with 500, 250, and 100 neurons using ReLu as activation function
- Output layer with 10 neurons (linear, raw logits)
- Use of softmax for probability assignment
- Glorot uniform weight initialization and zeros biases per assignment

```
class MLP_ReLu(nn.Module):
    def __init__(self, input_dim=784, hidden_dims=[500, 250, 100],
output_dim=10 ):

        layer_sizes = [input_dim] + hidden_dims + [output_dim]

        # Initialize weights and biases (Glorot initialization, zero
biases)
        self.weights = []
        self.biases = []
        for i in range(len(layer_sizes) - 1):
            ip__layer_size = layer_sizes[i]
            op__layer_size = layer_sizes[i + 1]
            limit = np.sqrt(6 / (ip__layer_size + op__layer_size))
            W = np.random.uniform(-limit, limit, size=(ip__layer_size,
op__layer_size))
            b = np.zeros((1, op__layer_size))
            self.weights.append(W)
            self.biases.append(b)

        def relu(self, x):
            return np.maximum(0, x)

        def relu_derivative(self, x):
            return (x > 0).astype(float)

        def softmax(self, x):
            exp_x = np.exp(x - np.max(x))
            return exp_x / exp_x.sum(axis=1, keepdims=True)

        def forward(self, X):
            self.inputs = []
            self.z_values = []

            # Input layer to first hidden layer
            z1 = X @ self.weights[0] + self.biases[0]
            a1 = self.relu(z1)

            # First hidden layer to second hidden layer
```



```

z2 = a1 @ self.weights[1] + self.biases[1]
a2 = self.relu(z2)

# Second hidden layer to third one
z3 = a2 @ self.weights[2] + self.biases[2]
a3 = self.relu(z3)

# Third hidden layer to output using Softmax
z4 = a3 @ self.weights[3] + self.biases[3]
output = self.softmax(z4)

self.inputs = [X, a1, a2, a3]
self.z_values = [z1, z2, z3, z4]
return output

def compute_loss(self, y_true, y_pred):
    m = y_true.shape[0] #number of samples in the batch
    eps = 1e-9 #value to avoid log(0)
    log_preds = np.log(y_pred + eps) #log of prediction
    loss = -np.sum(y_true * log_preds) / m
    return loss

def backward(self, y_true, y_pred, learning_rate):
    m = y_true.shape[0]
    delta4 = (y_pred - y_true) / m

    dW4 = self.inputs[3].T @ delta4
    db4 = np.sum(delta4, axis=0, keepdims=True)

    delta3 = delta4 @ self.weights[3].T *
self.relu_derivative(self.inputs[3])
    dW3 = self.inputs[2].T @ delta3
    db3 = np.sum(delta3, axis=0, keepdims=True)

    delta2 = delta3 @ self.weights[2].T *
self.relu_derivative(self.inputs[2])
    dW2 = self.inputs[1].T @ delta2
    db2 = np.sum(delta2, axis=0, keepdims=True)

    delta1 = delta2 @ self.weights[1].T *
self.relu_derivative(self.inputs[1])
    dW1 = self.inputs[0].T @ delta1
    db1 = np.sum(delta1, axis=0, keepdims=True)

    # Update weights and biases
    self.weights[3] -= learning_rate * dW4
    self.biases[3] -= learning_rate * db4

    self.weights[2] -= learning_rate * dW3
    self.biases[2] -= learning_rate * db3

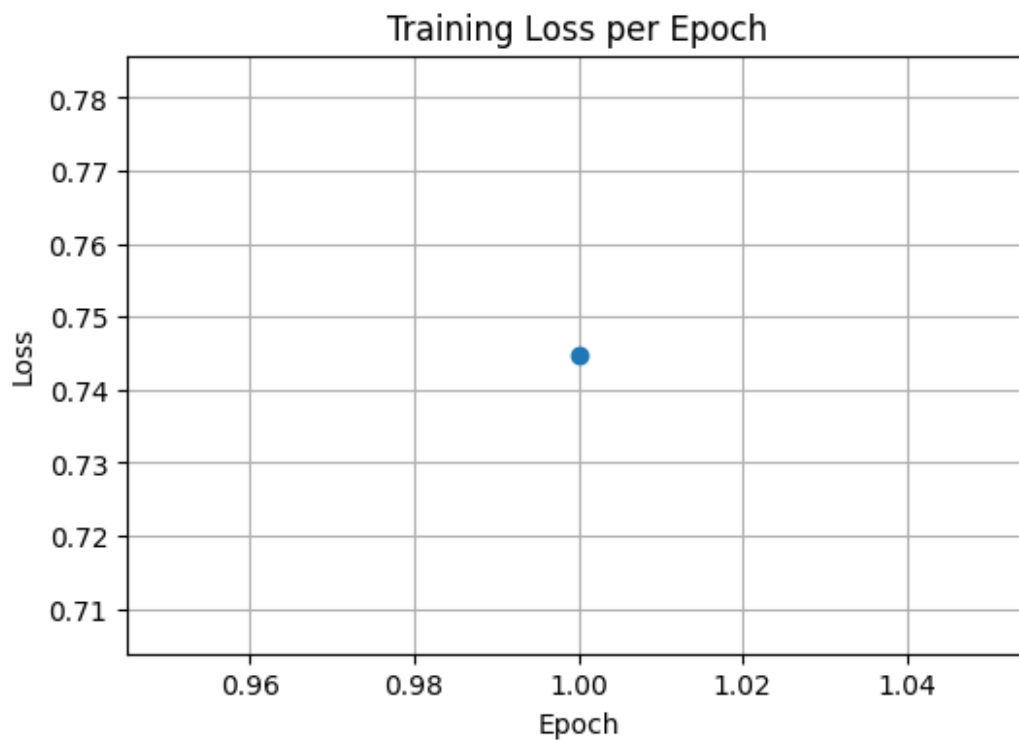
```

```
self.weights[1] -= learning_rate * dw2
self.biases[1] -= learning_rate * db2

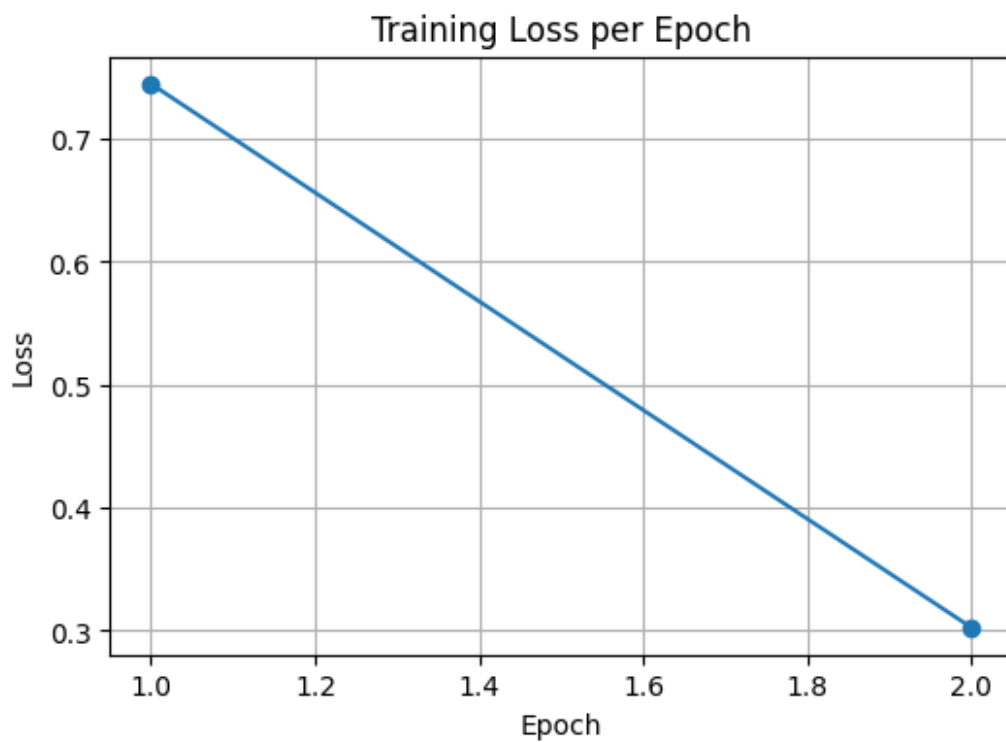
self.weights[0] -= learning_rate * dw1
self.biases[0] -= learning_rate * db1

model = MLP_ReLu()
loss_history_ReLu, iter_ReLu, iter_loss_ReLu = train(model,
train_loader)

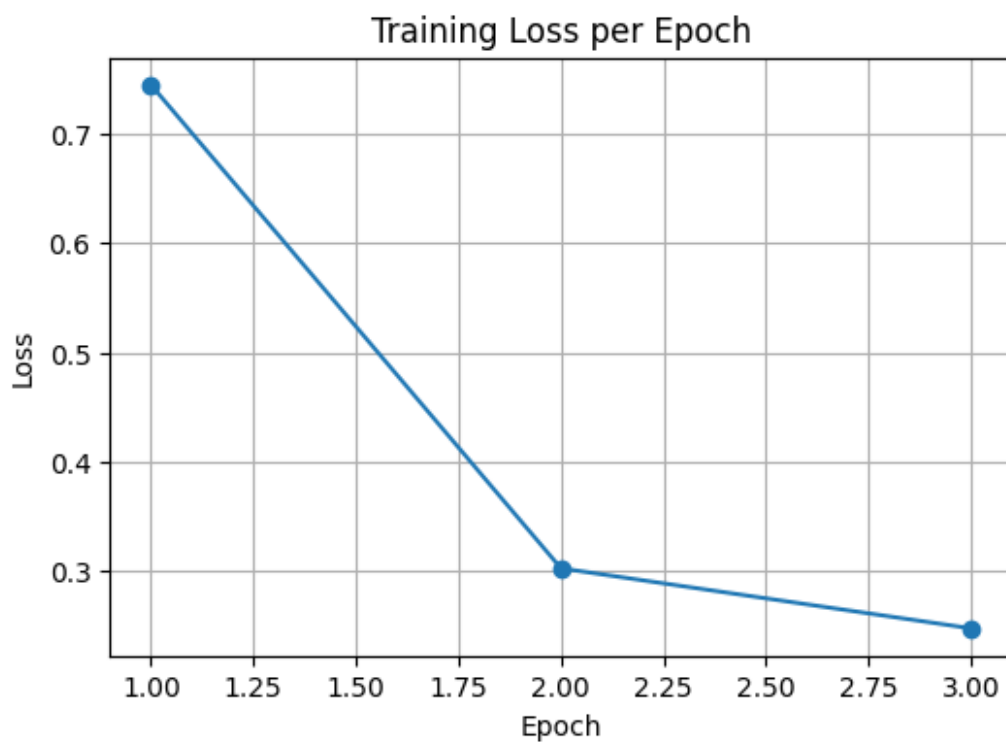
Epoch 1/15, Loss: 0.7448
```



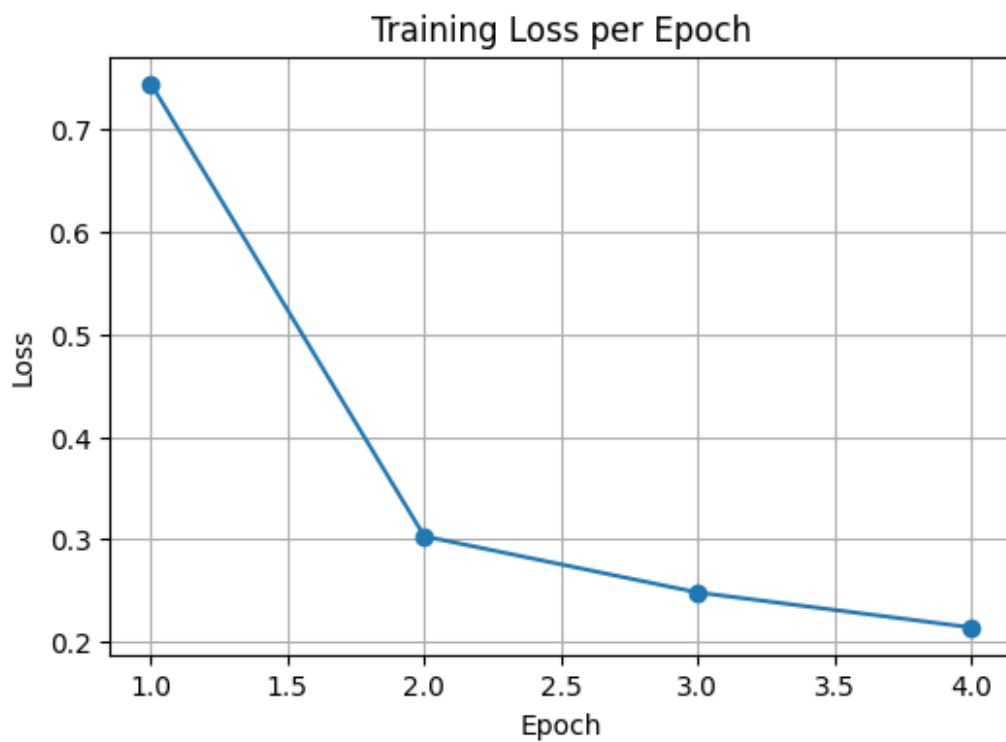
Epoch 2/15, Loss: 0.3028



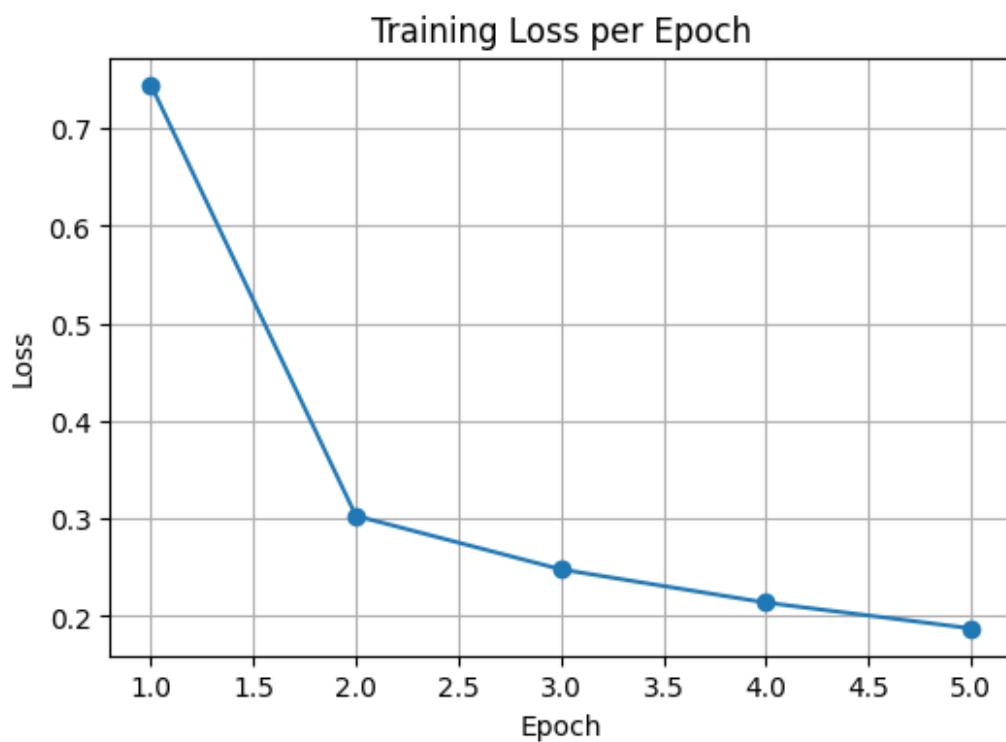
Epoch 3/15, Loss: 0.2479



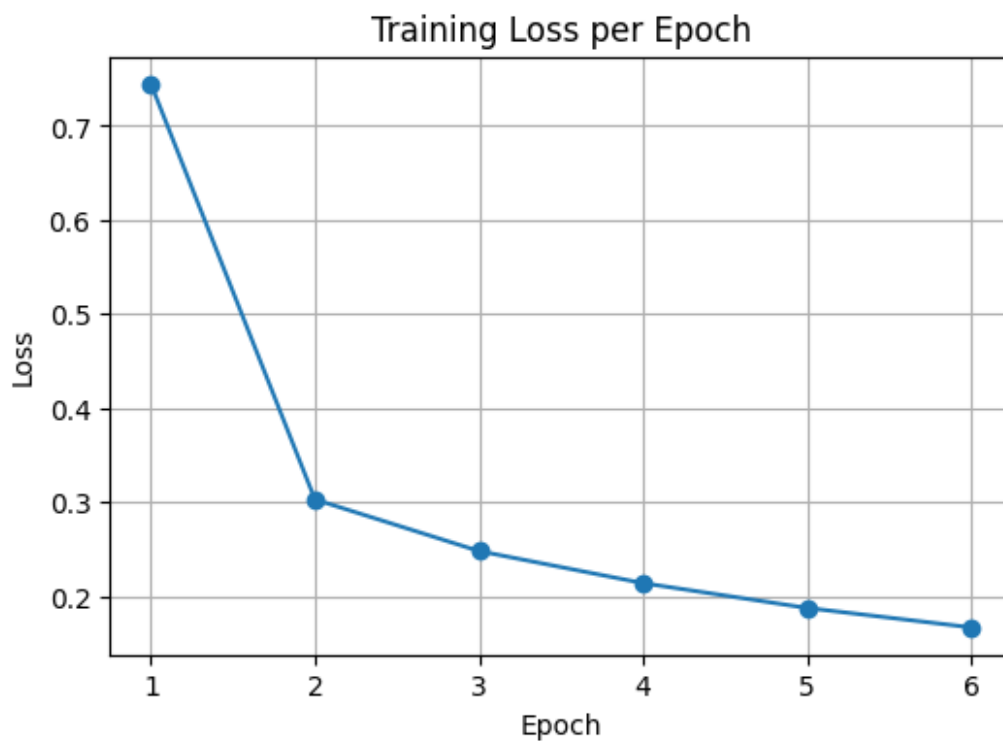
Epoch 4/15, Loss: 0.2138



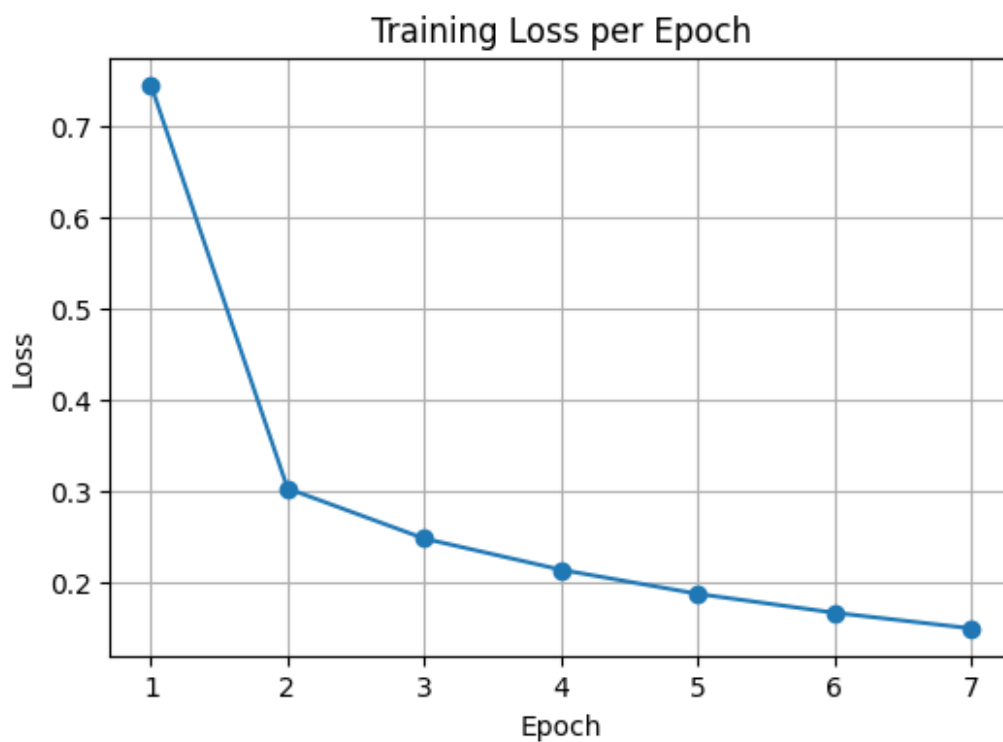
Epoch 5/15, Loss: 0.1874



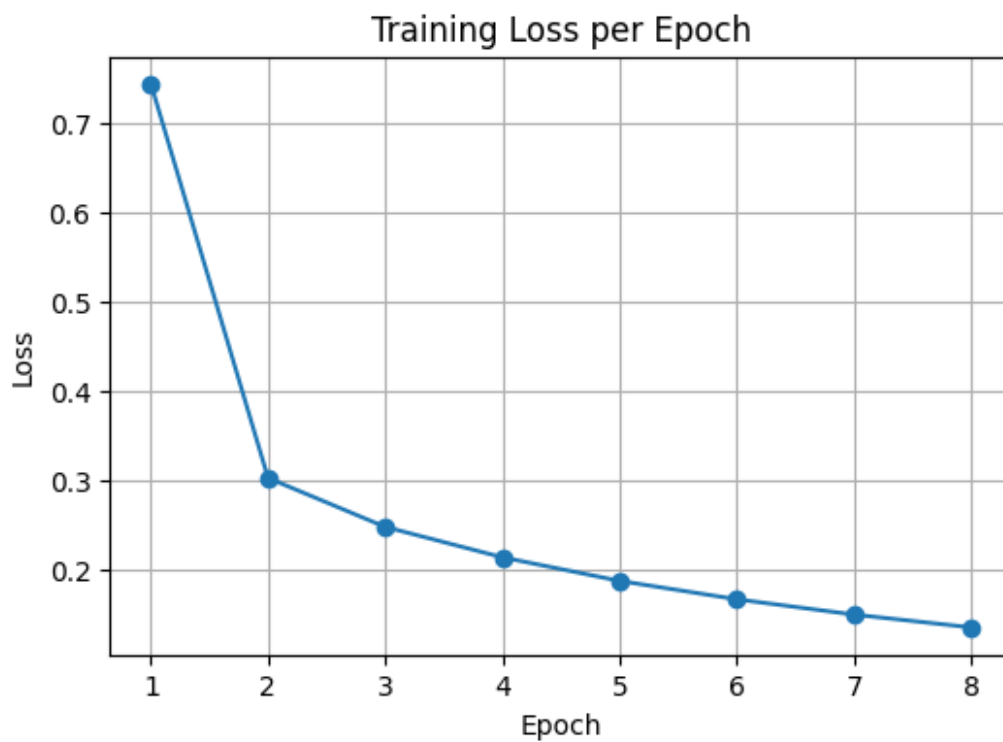
Epoch 6/15, Loss: 0.1668



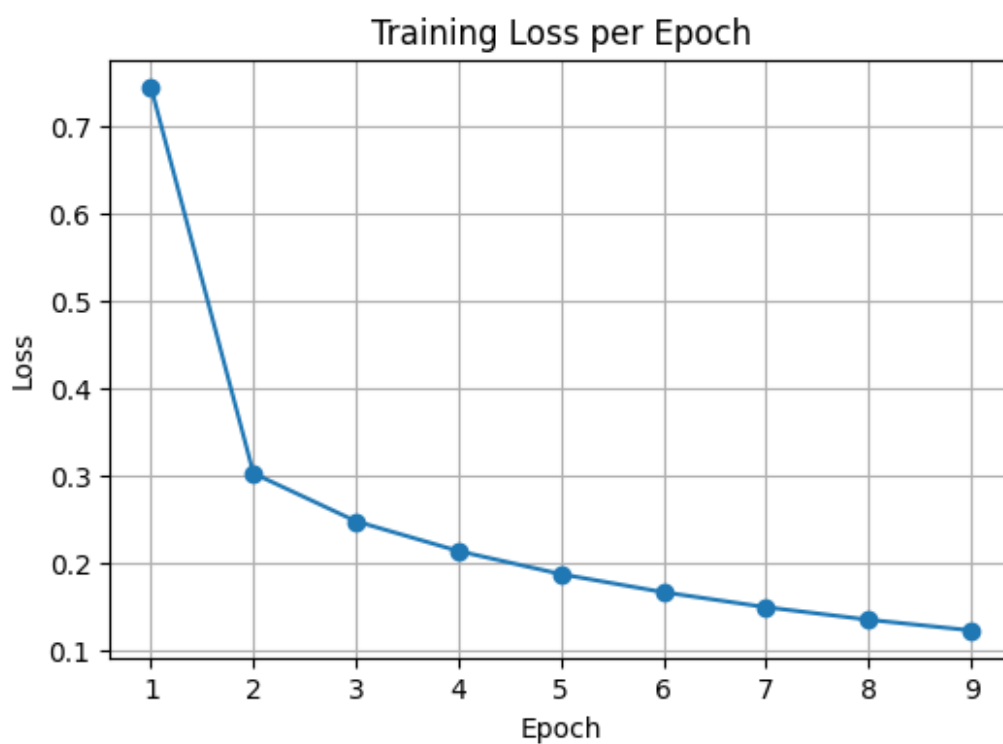
Epoch 7/15, Loss: 0.1496



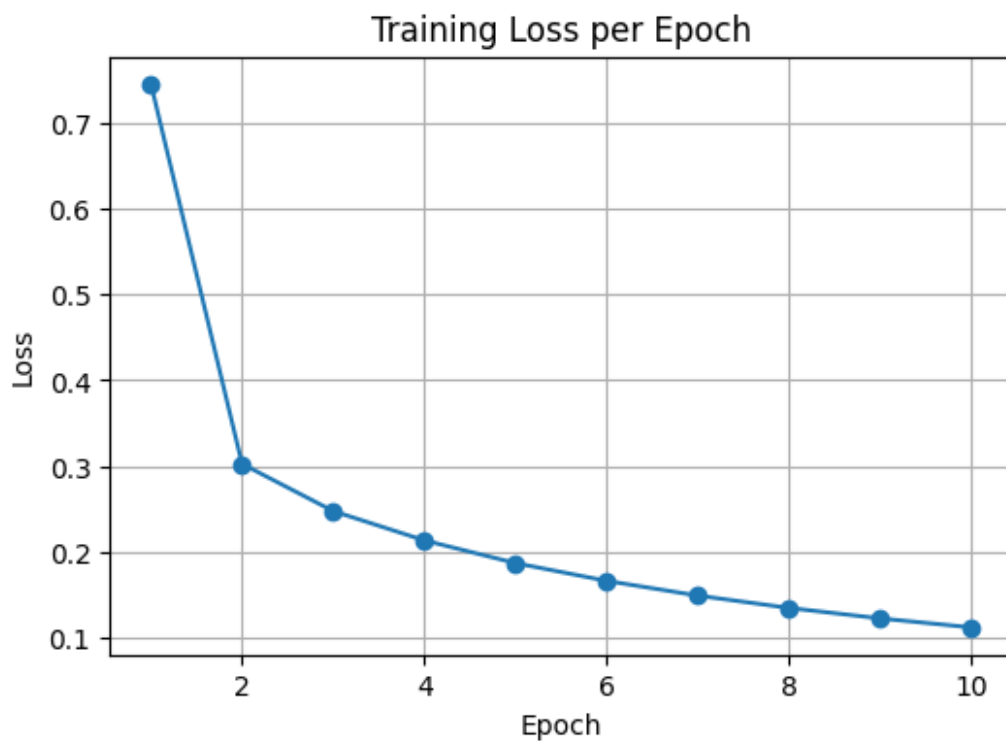
Epoch 8/15, Loss: 0.1353



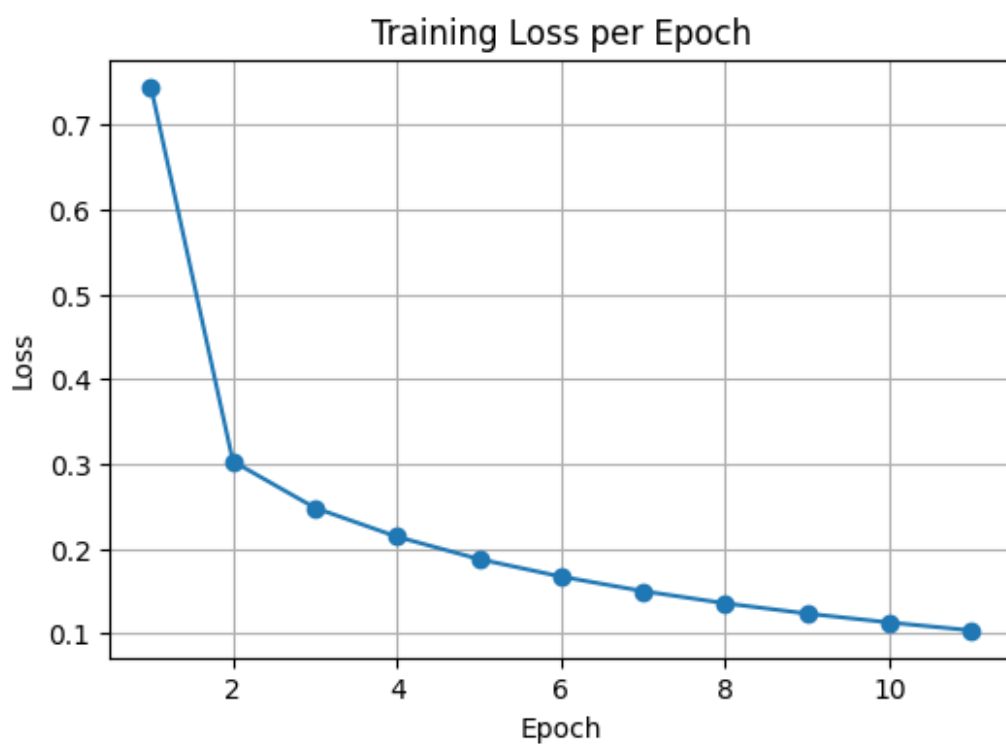
Epoch 9/15, Loss: 0.1231



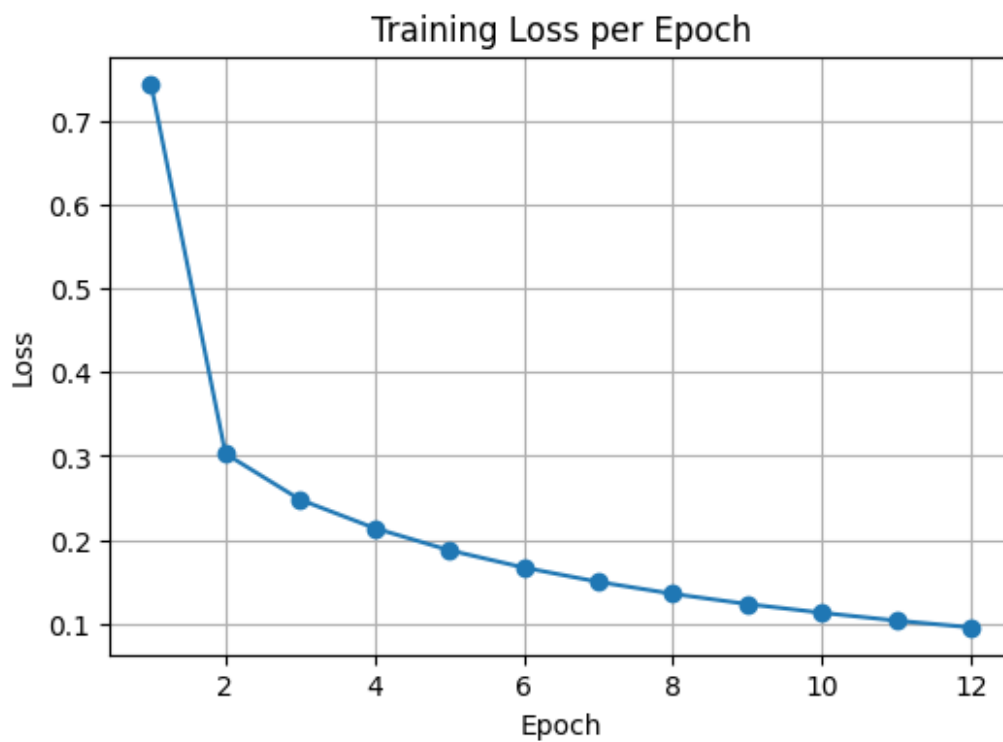
Epoch 10/15, Loss: 0.1127



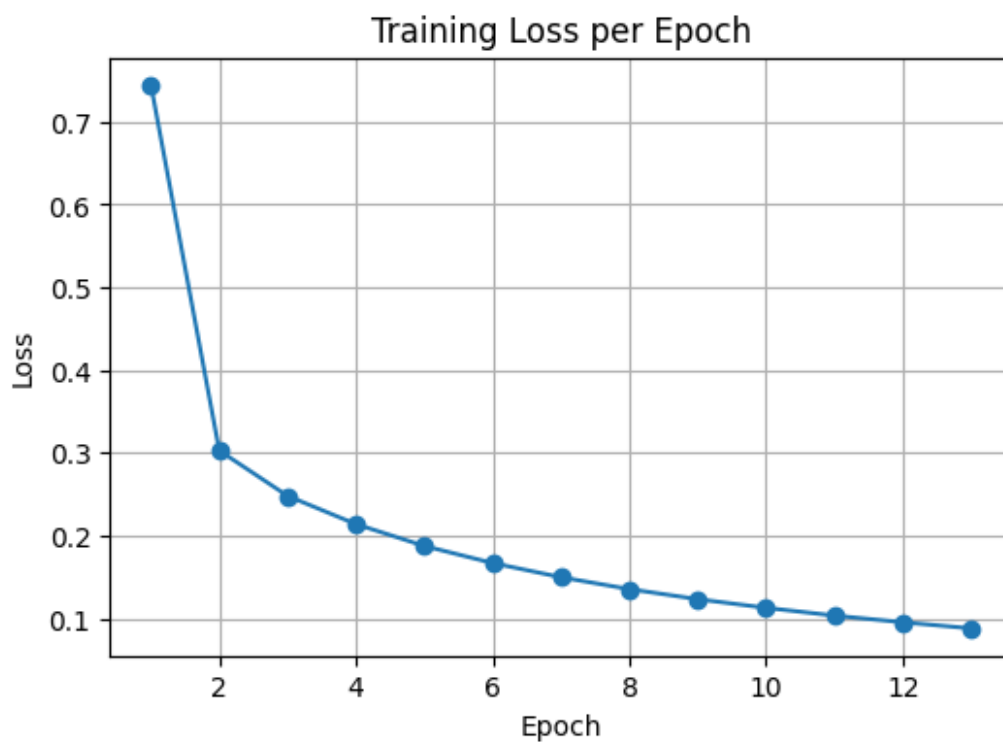
Epoch 11/15, Loss: 0.1033



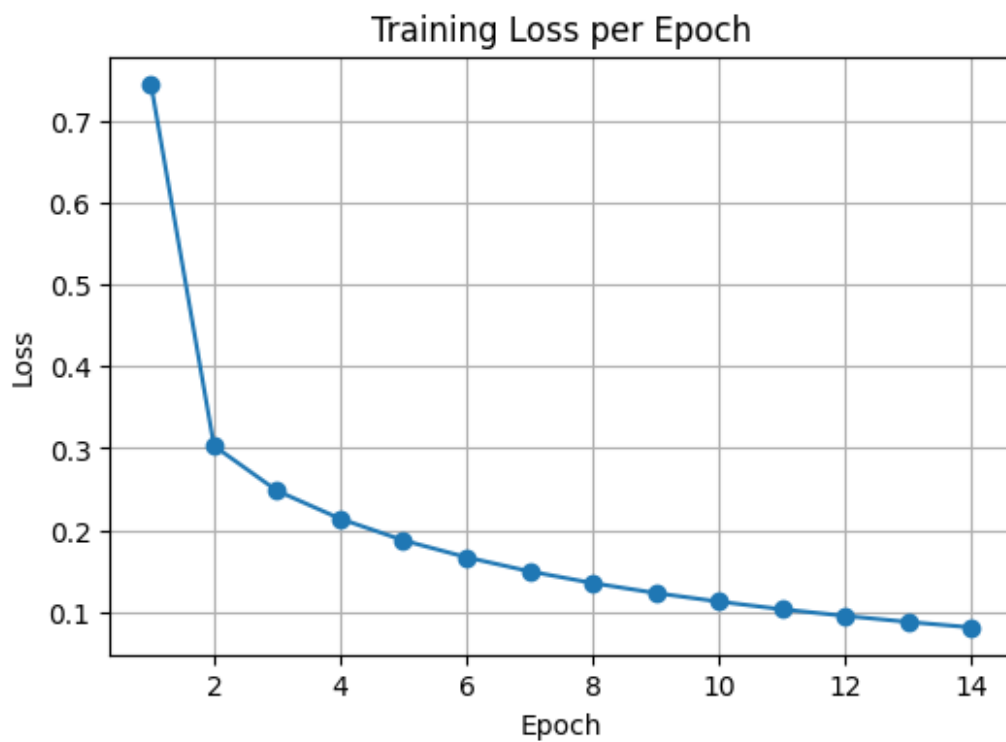
Epoch 12/15, Loss: 0.0952



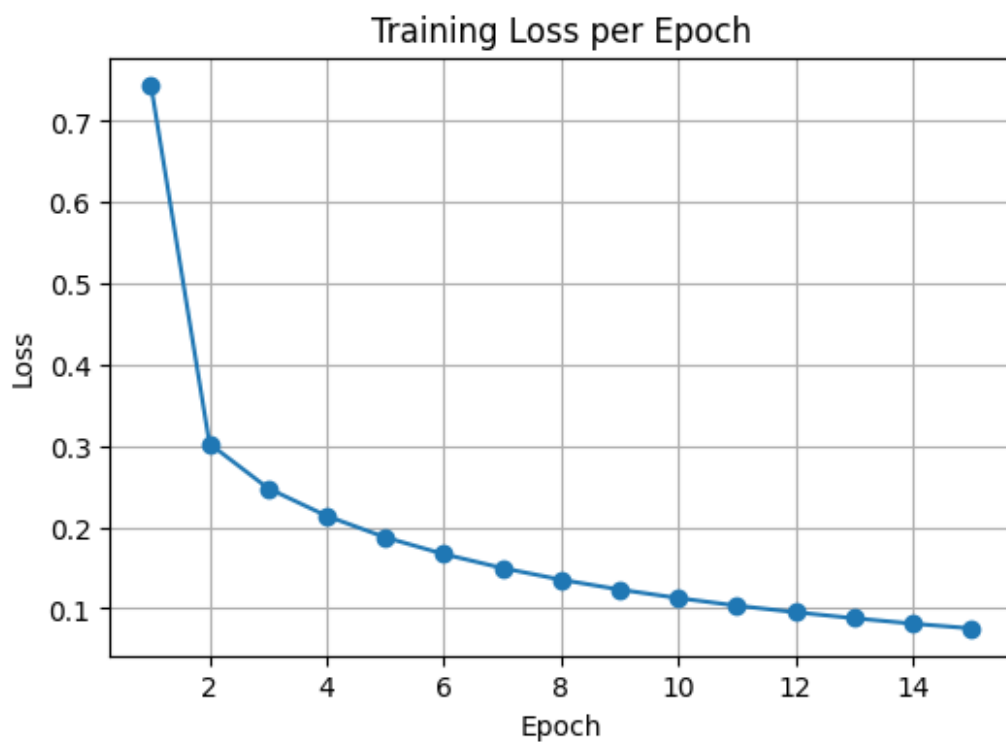
Epoch 13/15, Loss: 0.0878

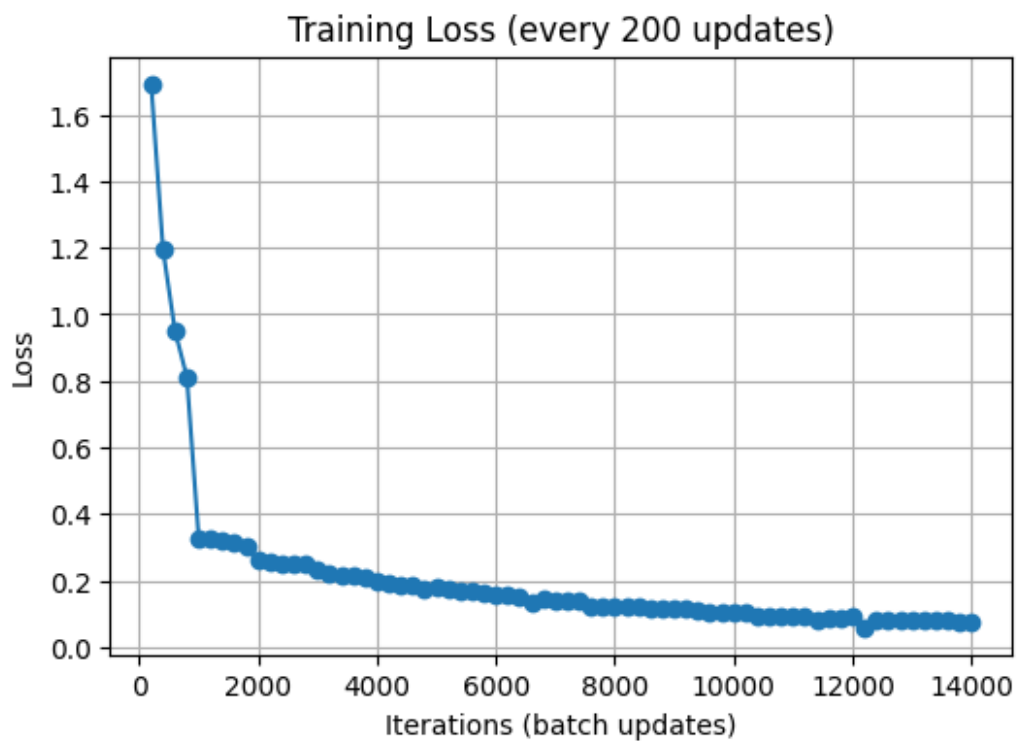


Epoch 14/15, Loss: 0.0812

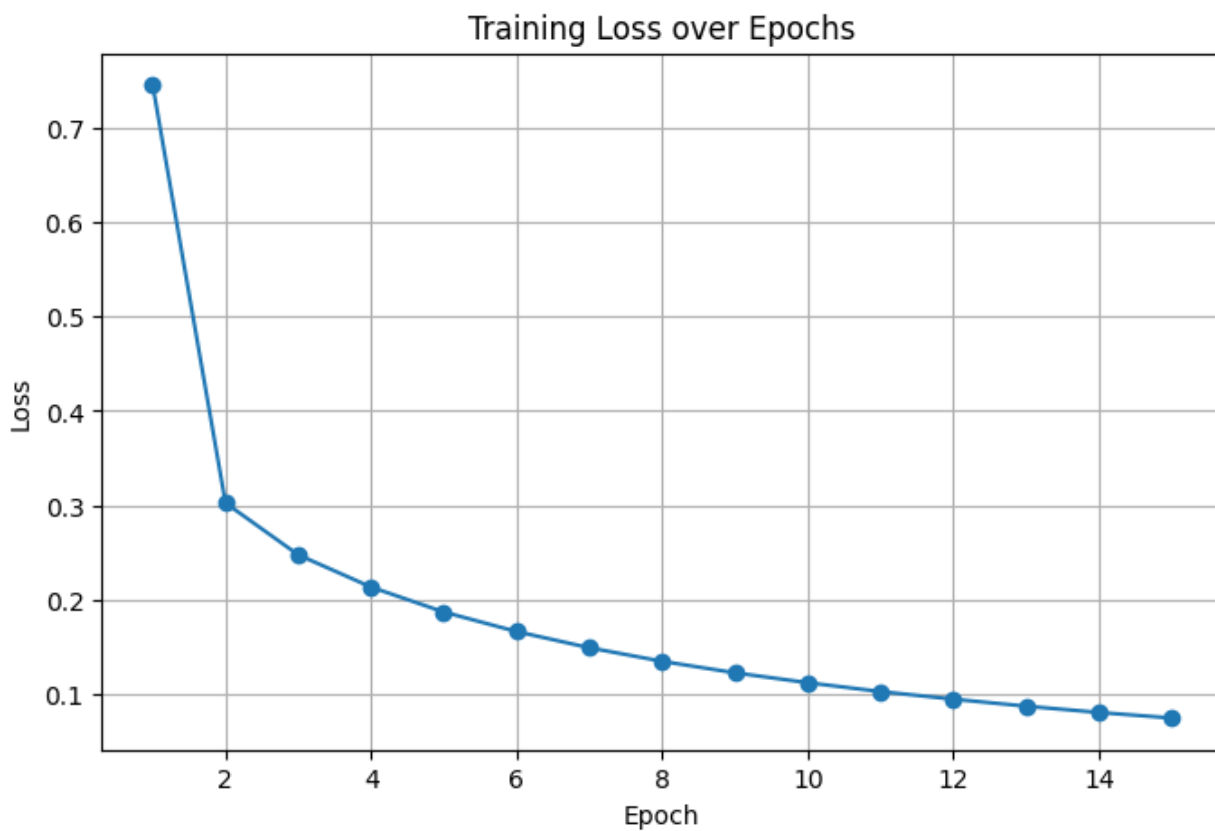


Epoch 15/15, Loss: 0.0753





```
plot_training_curves(loss_history_ReLu)
```



```
# Prepare test set in numpy format
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("---- Test Set Evaluation ----")
evaluate(model, X_test_np, y_test_np)
```

```
---- Test Set Evaluation ----
```

```
Accuracy: 97.1200%
```

```
Confusion Matrix:
```

```
[[ 963    0    1    1    0    5    4    2    2    2]
 [    0 1126    4    0    0    1    4    0    0    0]
 [    4    3 1003    4    4    0    2    6    6    0]
 [    0    0    4  983    0    8    0    7    5    3]
 [    2    0    4    1  957    0    3    2    2   11]
 [    4    0    0    6    2  866    9    0    3    2]
 [    6    3    0    0    6   10  931    0    2    0]
 [    1    9   10    4    2    1    0  990    2    9]
 [    3    0    3    9    3    6    7    6  933    4]
 [    4    5    1    8   17    3    1    7    3  960]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	980
1	0.98	0.99	0.99	1135
2	0.97	0.97	0.97	1032
3	0.97	0.97	0.97	1010
4	0.97	0.97	0.97	982
5	0.96	0.97	0.97	892
6	0.97	0.97	0.97	958
7	0.97	0.96	0.97	1028
8	0.97	0.96	0.97	974
9	0.97	0.95	0.96	1009
accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000

```
np.float64(0.9712)
```

Results for MLP using ReLU activation function

- Classification accuracy over test data = **97.12%**
- Training loss after 15 epochs = **0.0753**

#Question 1 Part C

#Implementing Multi Layer Perceptron Algorithm using Tan Hyperbolic activation function

Model architecture:

- Input layer (784 neurons)
- Hidden layers with 500, 250, and 100 neurons using Tanh as activation
- Output layer with 10 neurons (linear, raw logits)
- Use of softmax for probability assignment
- Glorot uniform weight initialization and zeros biases per assignment

```
class MLP_tanh:
    def __init__(self, input_dim = 784, hidden_dims = [500, 250, 100],
output_dim = 10 ):

        layer_sizes = [input_dim] + hidden_dims + [output_dim]

        # Initialize weights and biases (Glorot initialization, zero
biases)
        self.weights = []
        self.biases = []
        for i in range(len(layer_sizes) - 1):
            ip__layer_size = layer_sizes[i]
            op__layer_size = layer_sizes[i + 1]
            limit = np.sqrt(6 / (ip__layer_size + op__layer_size))
            W = np.random.uniform(-limit, limit, size=(ip__layer_size,
op__layer_size))
            b = np.zeros((1, op__layer_size))
            self.weights.append(W)
            self.biases.append(b)

        def tanh(self, x):
            return np.tanh(x)

        def tanh_derivative(self, x):
            return 1.0 - np.tanh(x)**2

        def softmax(self, x):
            exp_x = np.exp(x - np.max(x, axis=1, keepdims=True))
            return exp_x / exp_x.sum(axis=1, keepdims=True)

        def forward(self, X):
            self.inputs = []
            self.z_values = []

            # Input layer to first hidden layer
            z1 = X @ self.weights[0] + self.biases[0]
            a1 = self.tanh(z1)

            # First hidden layer to second hidden layer
```

```

z2 = a1 @ self.weights[1] + self.biases[1]
a2 = self.tanh(z2)

# Second hidden layer to third one
z3 = a2 @ self.weights[2] + self.biases[2]
a3 = self.tanh(z3)

# Third hidden layer to output using Softmax
z4 = a3 @ self.weights[3] + self.biases[3]
output = self.softmax(z4)

self.inputs = [X, a1, a2, a3]
self.z_values = [z1, z2, z3, z4]
return output

def compute_loss(self, y_true, y_pred):
    m = y_true.shape[0] #number of samples in the batch
    eps = 1e-9 #value to avoid log(0)
    log_preds = np.log(y_pred + eps) #log of prediction
    loss = -np.sum(y_true * log_preds) / m
    return loss

def backward(self, y_true, y_pred, learning_rate):
    m = y_true.shape[0]
    delta4 = (y_pred - y_true) / m

    dW4 = self.inputs[3].T @ delta4
    db4 = np.sum(delta4, axis=0, keepdims=True)

    delta3 = delta4 @ self.weights[3].T *
self.tanh_derivative(self.inputs[3])
    dW3 = self.inputs[2].T @ delta3
    db3 = np.sum(delta3, axis=0, keepdims=True)

    delta2 = delta3 @ self.weights[2].T *
self.tanh_derivative(self.inputs[2])
    dW2 = self.inputs[1].T @ delta2
    db2 = np.sum(delta2, axis=0, keepdims=True)

    delta1 = delta2 @ self.weights[1].T *
self.tanh_derivative(self.inputs[1])
    dW1 = self.inputs[0].T @ delta1
    db1 = np.sum(delta1, axis=0, keepdims=True)

    # Update weights and biases
    self.weights[3] -= learning_rate * dW4
    self.biases[3] -= learning_rate * db4

    self.weights[2] -= learning_rate * dW3
    self.biases[2] -= learning_rate * db3

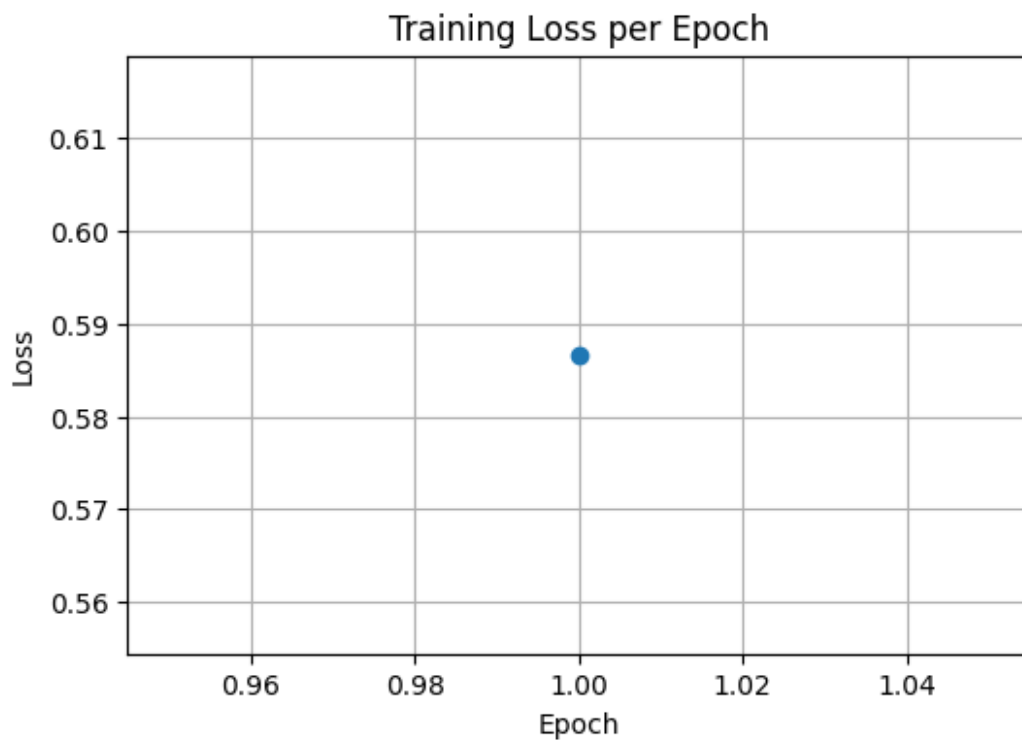
```

```
self.weights[1] -= learning_rate * dW2
self.biases[1] -= learning_rate * db2

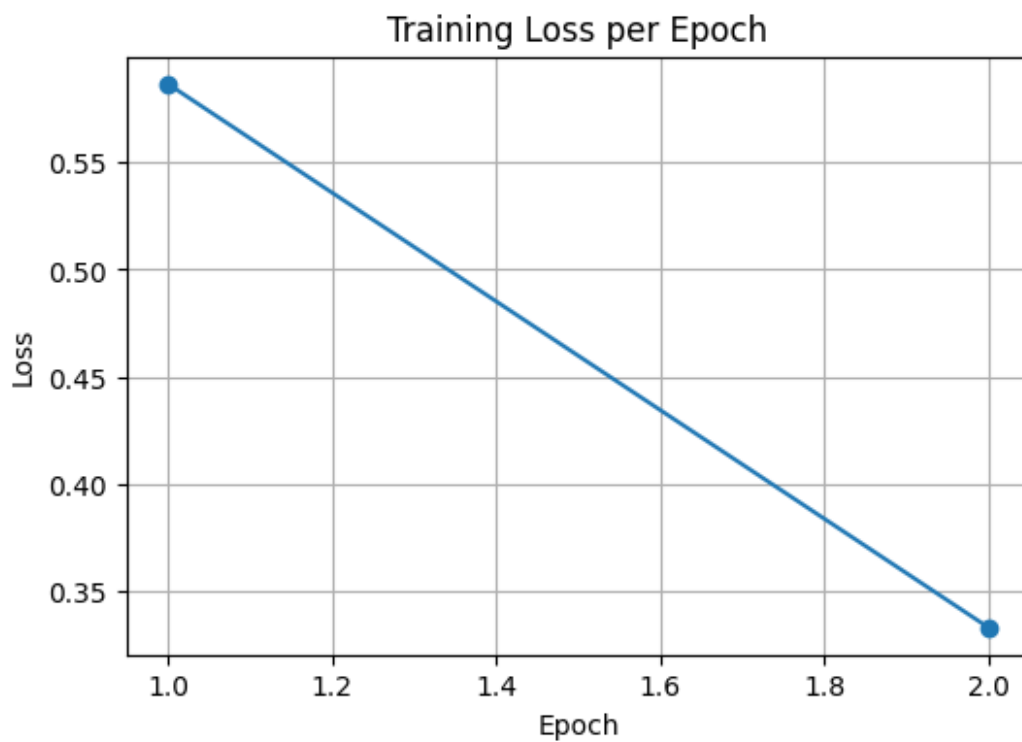
self.weights[0] -= learning_rate * dW1
self.biases[0] -= learning_rate * db1

model = MLP_tanh()
loss_history_tanh, iter_tanh, iter_loss_tanh = train(model,
train_loader)

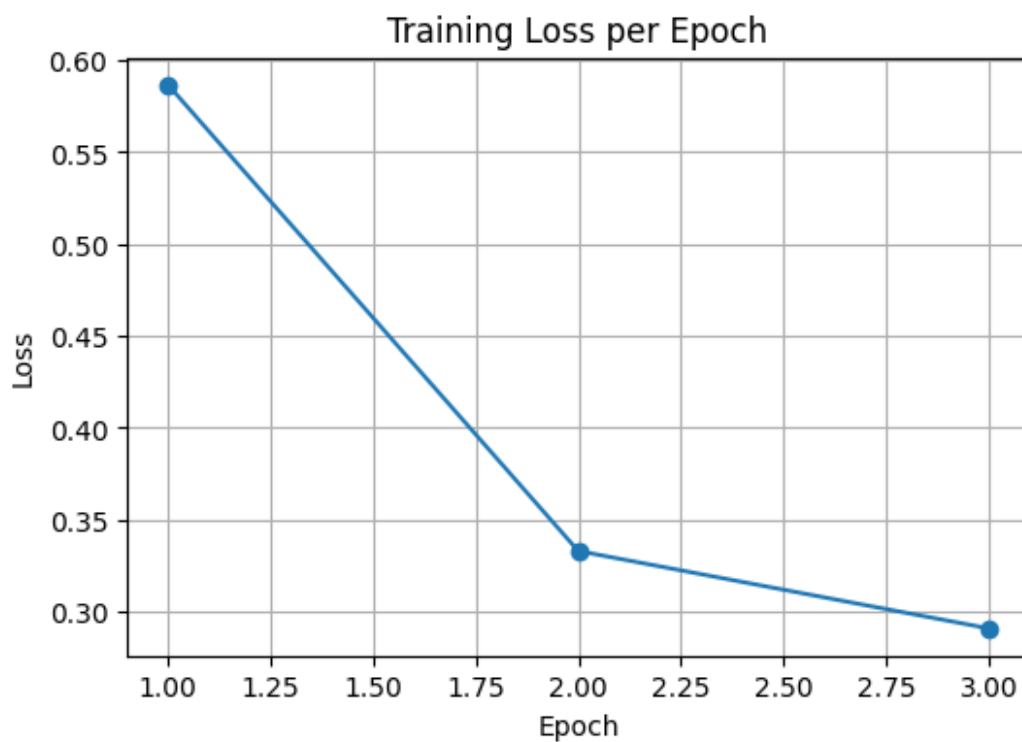
Epoch 1/15, Loss: 0.5867
```



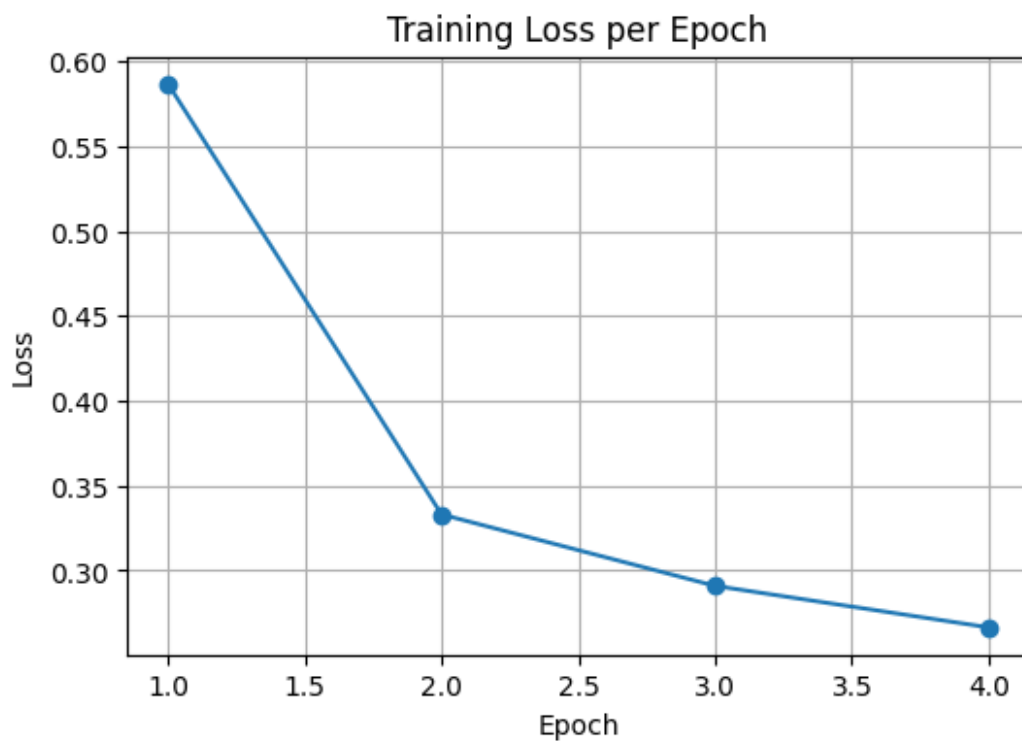
Epoch 2/15, Loss: 0.3330



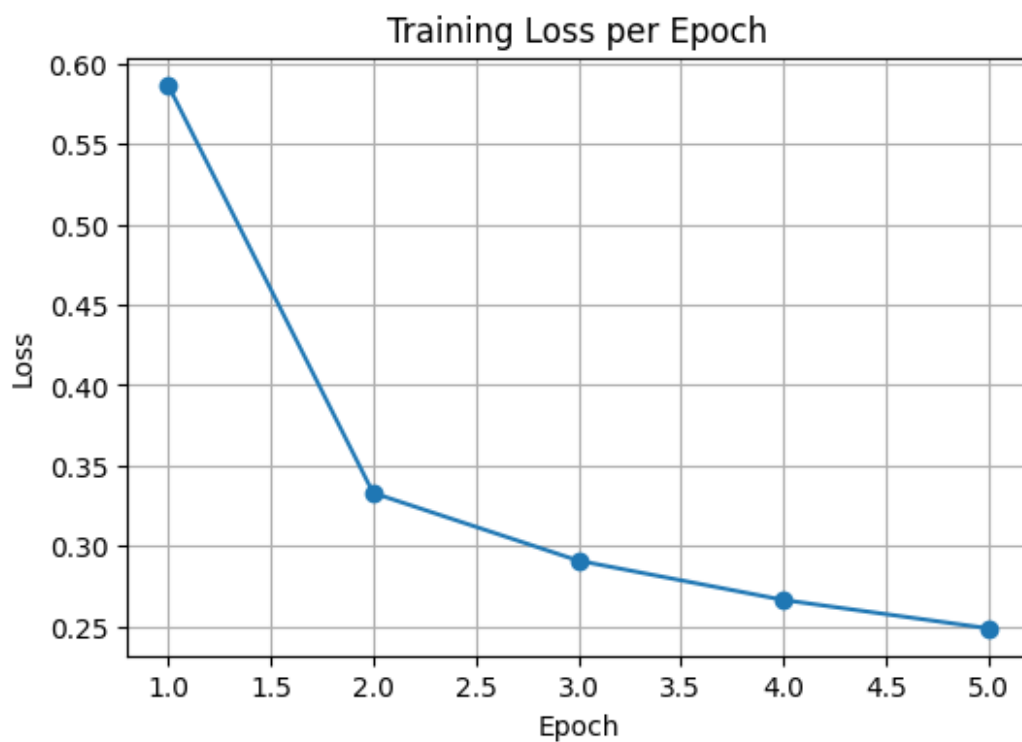
Epoch 3/15, Loss: 0.2909



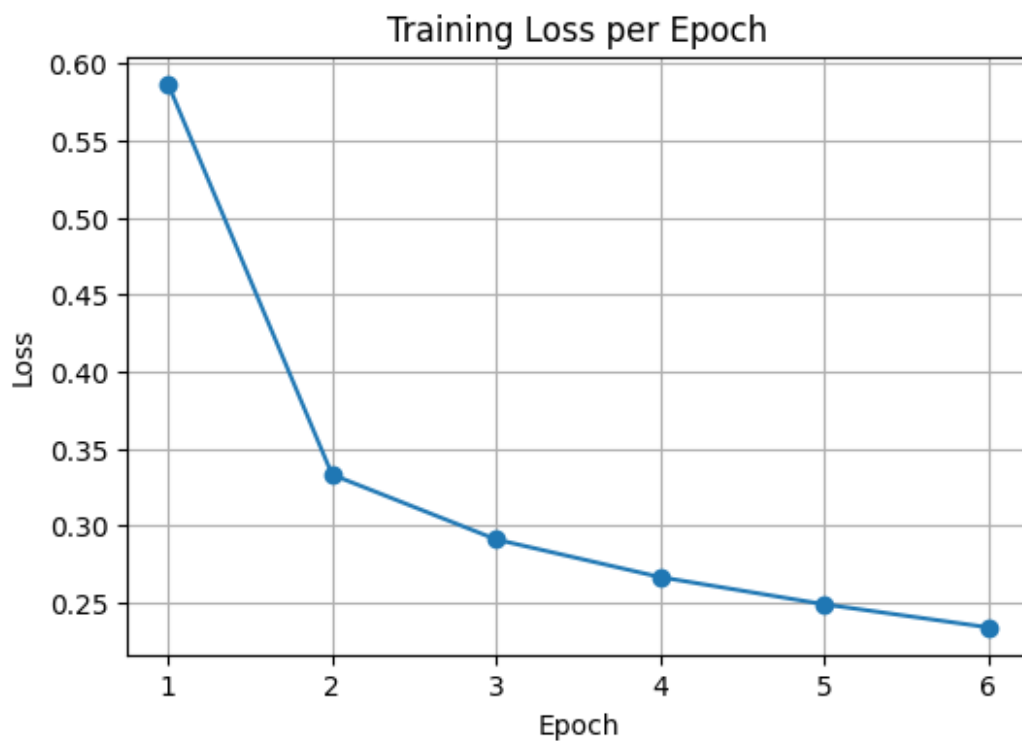
Epoch 4/15, Loss: 0.2665



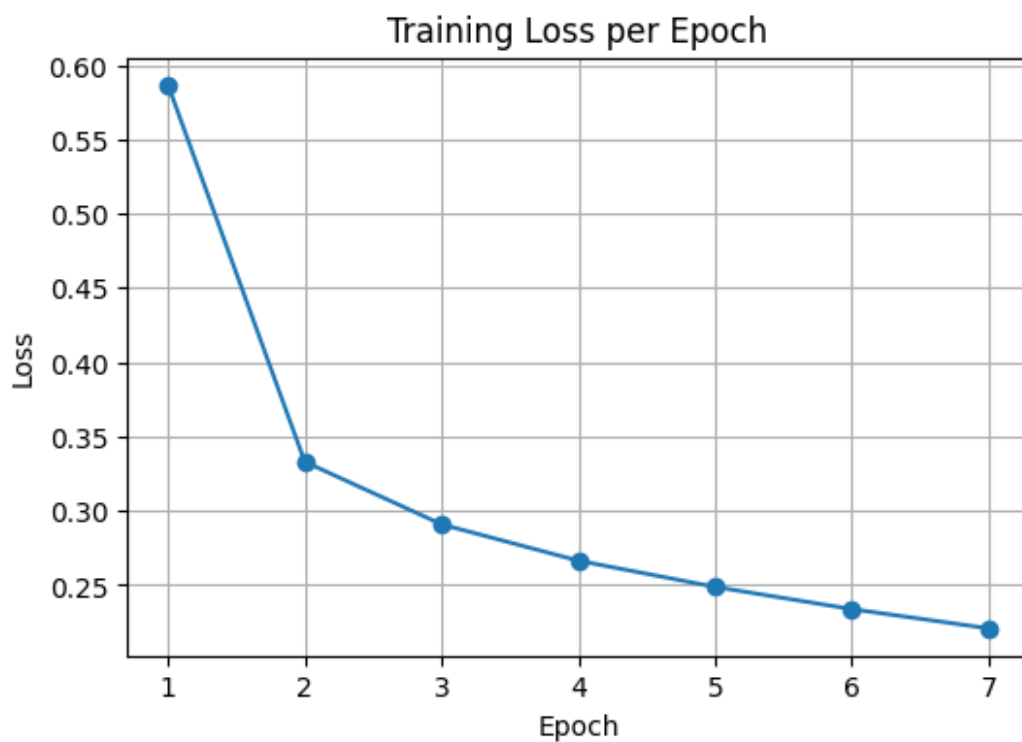
Epoch 5/15, Loss: 0.2488



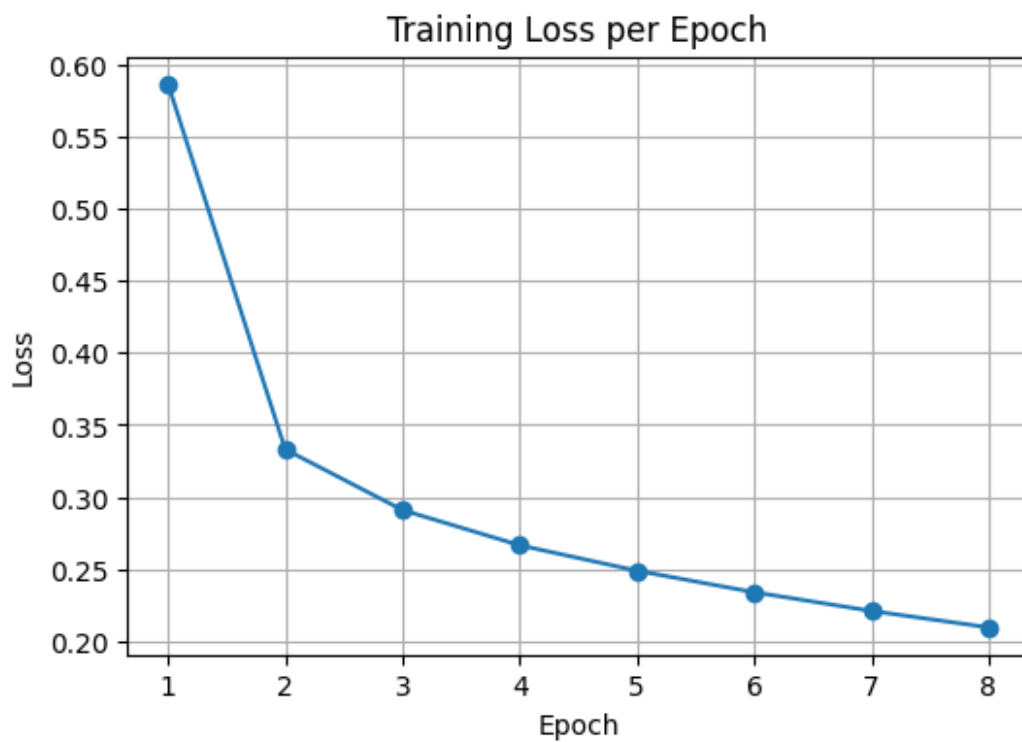
Epoch 6/15, Loss: 0.2339



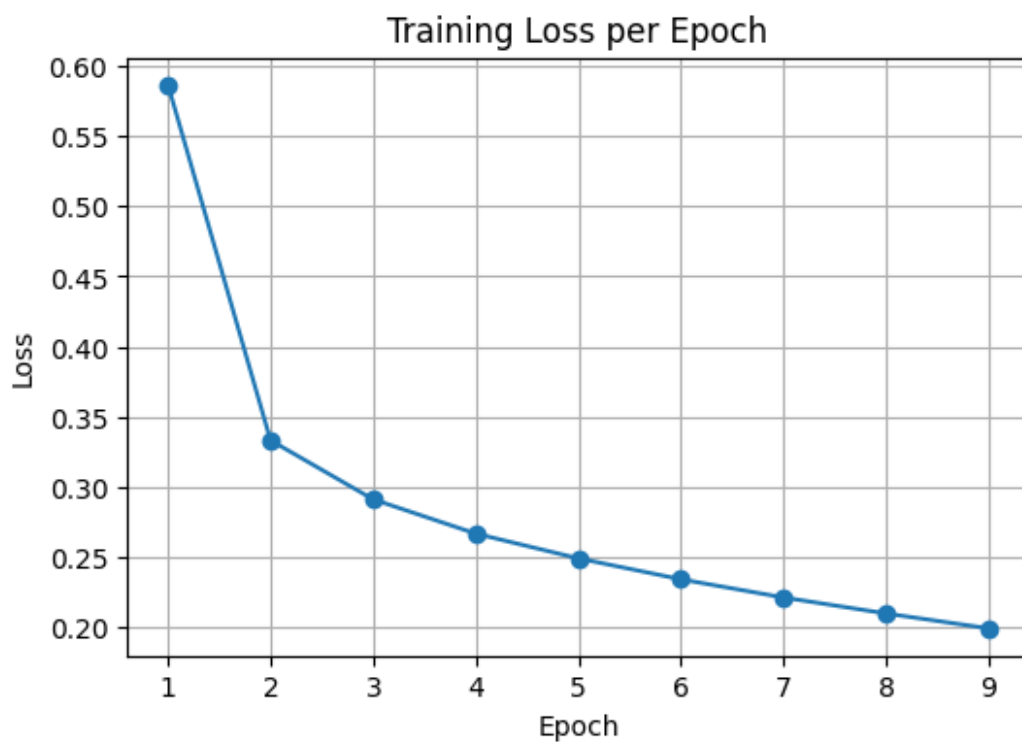
Epoch 7/15, Loss: 0.2209



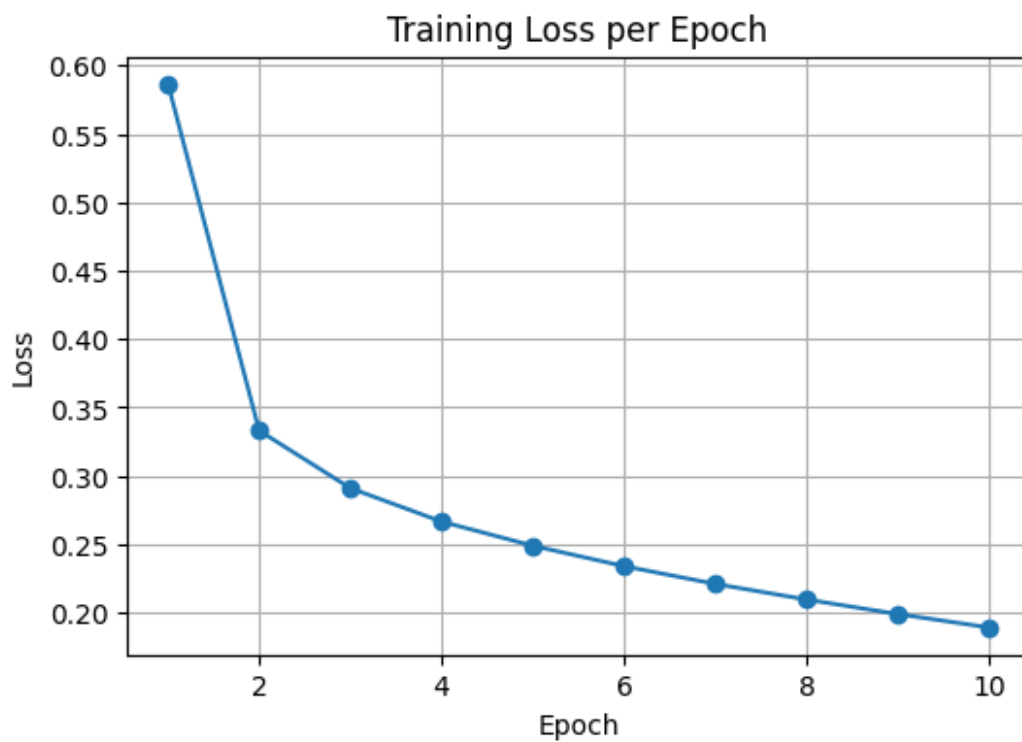
Epoch 8/15, Loss: 0.2095



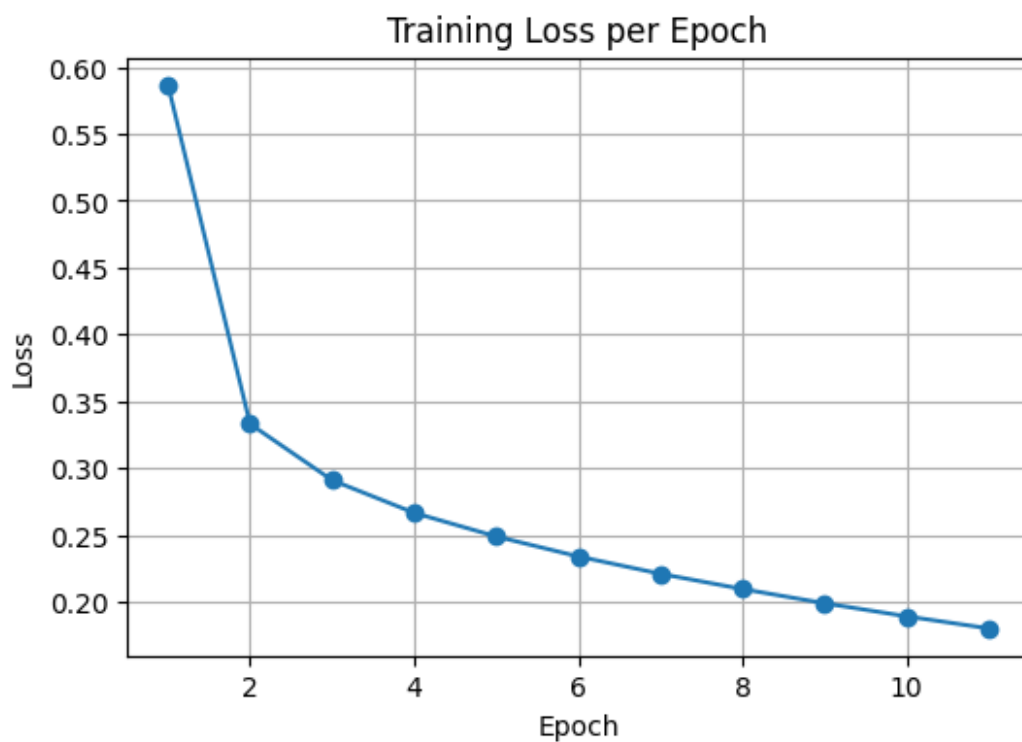
Epoch 9/15, Loss: 0.1989



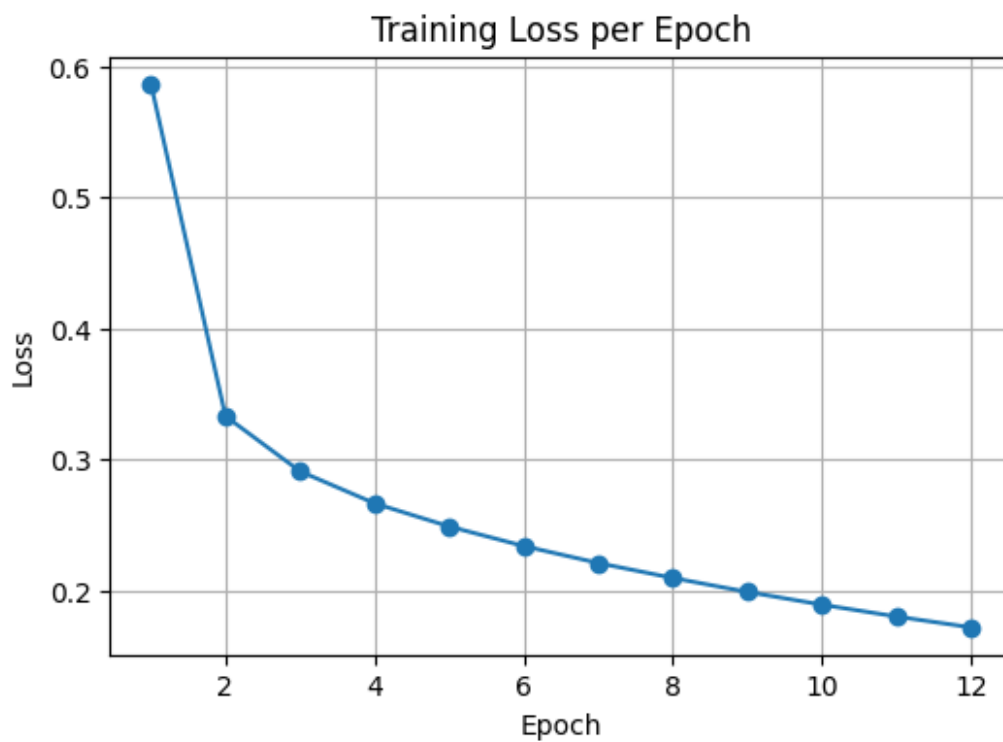
Epoch 10/15, Loss: 0.1891



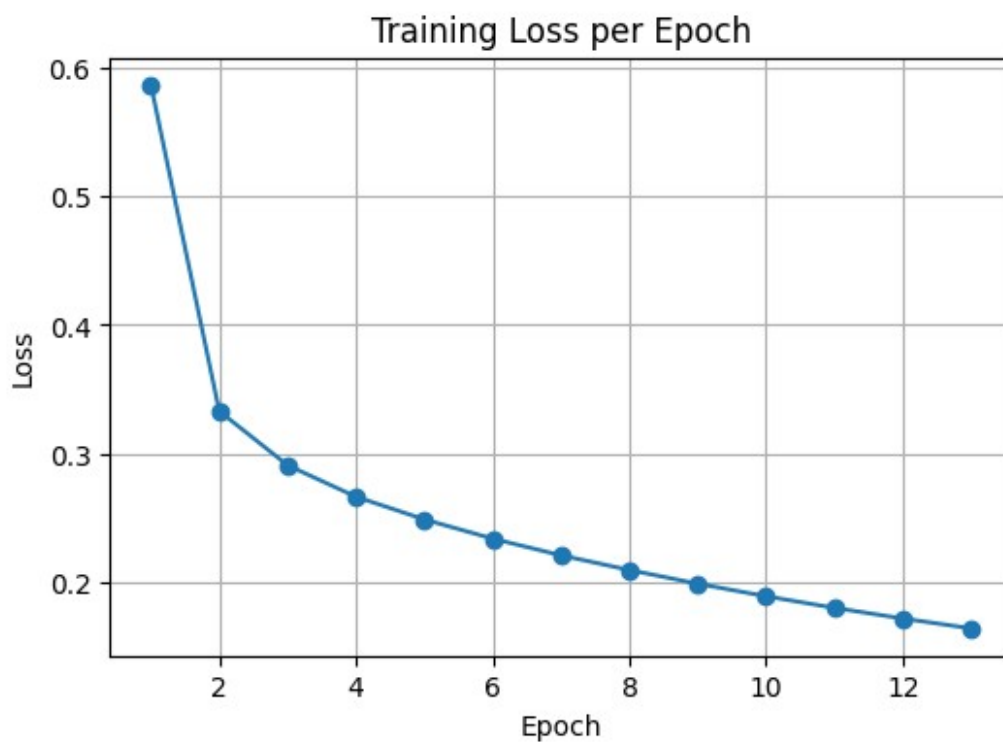
Epoch 11/15, Loss: 0.1801



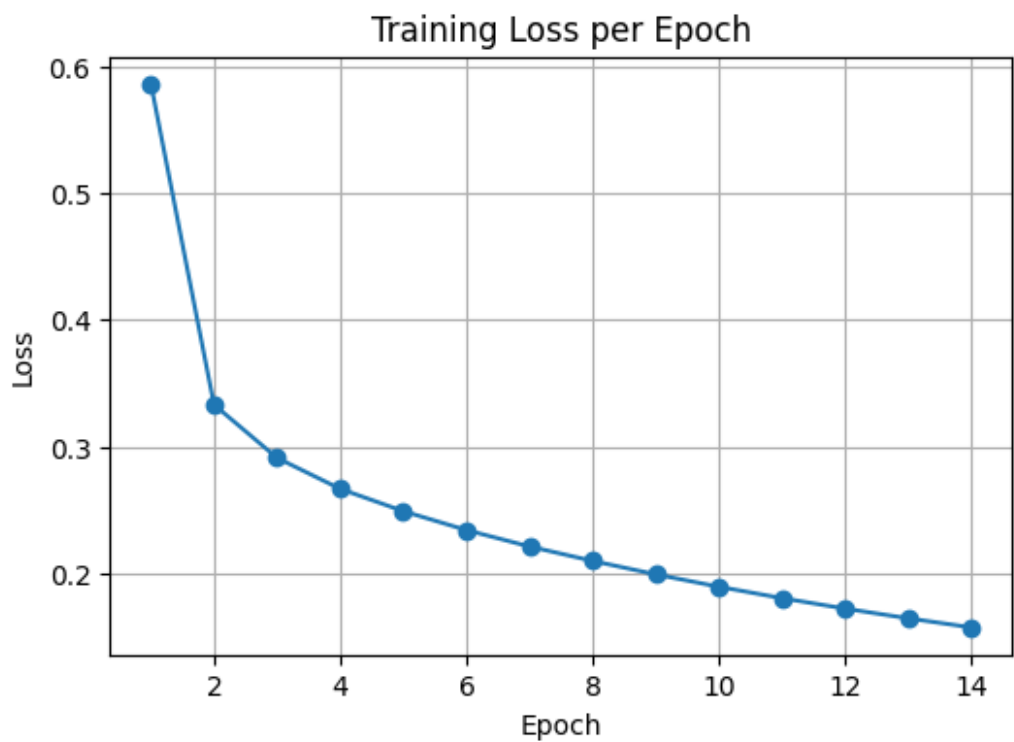
Epoch 12/15, Loss: 0.1718



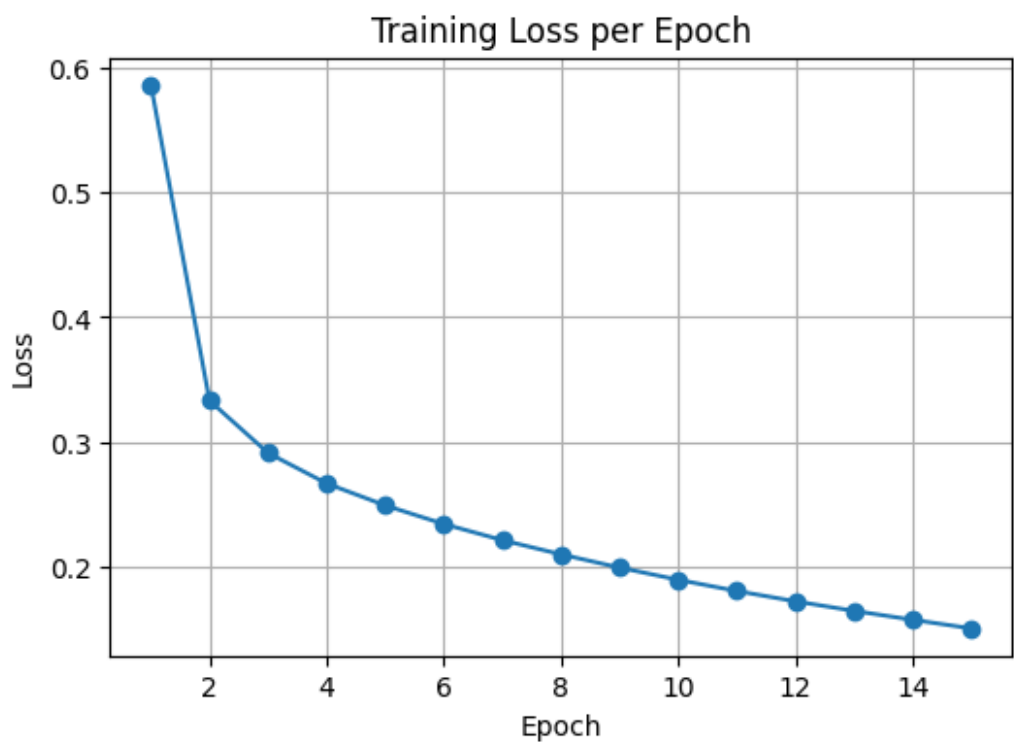
Epoch 13/15, Loss: 0.1642

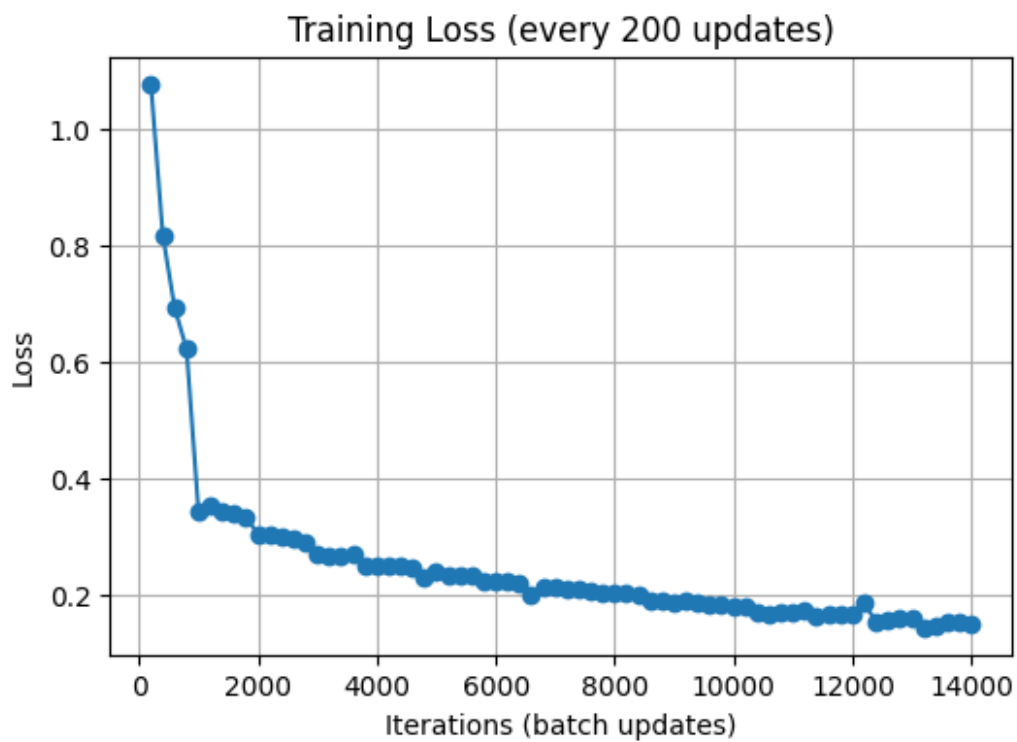


Epoch 14/15, Loss: 0.1570

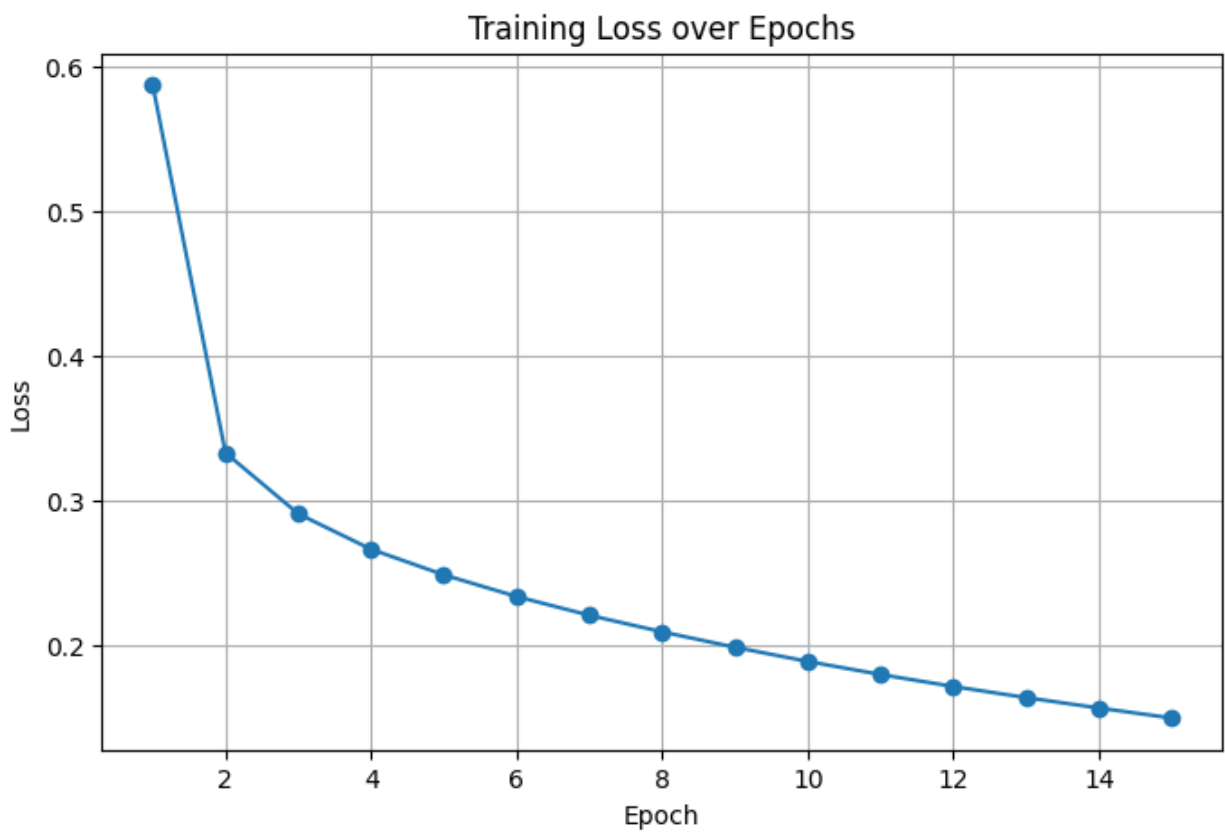


Epoch 15/15, Loss: 0.1501





```
plot_training_curves(loss_history_tanh)
```



```
# Prepare test set in numpy format
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("---- Test Set Evaluation ----")
evaluate(model, X_test_np, y_test_np)
```

```
---- Test Set Evaluation ----
```

```
Accuracy: 95.4700%
```

```
Confusion Matrix:
```

```
[[ 968    0    1    3    0    3    3    1    1    0]
 [   0 1120    3    2    0    1    3    2    4    0]
 [   6    1  988    5    4    1    9    6   11    1]
 [   0    0   12  972    0    7    0    9    9    1]
 [   1    0    9    0  931    0    6    3    5   27]
 [   8    2    0   24    2  827   12    2   11    4]
 [   9    3    6    2    3    8  923    1    3    0]
 [   2    6   20    6    2    2    0  969    2   19]
 [   3    3    8   20    6    9    7    9  909    0]
 [   9    6    2   13   18    4    1   12    4  940]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	980
1	0.98	0.99	0.98	1135
2	0.94	0.96	0.95	1032
3	0.93	0.96	0.95	1010
4	0.96	0.95	0.96	982
5	0.96	0.93	0.94	892
6	0.96	0.96	0.96	958
7	0.96	0.94	0.95	1028
8	0.95	0.93	0.94	974
9	0.95	0.93	0.94	1009
accuracy			0.95	10000
macro avg	0.95	0.95	0.95	10000
weighted avg	0.95	0.95	0.95	10000

```
np.float64(0.9547)
```

Results for MLP using Tan Hyperbolic activation function

- Classification accuracy over test data = **95.47%**
- Training loss after 15 epochs = **0.1501**

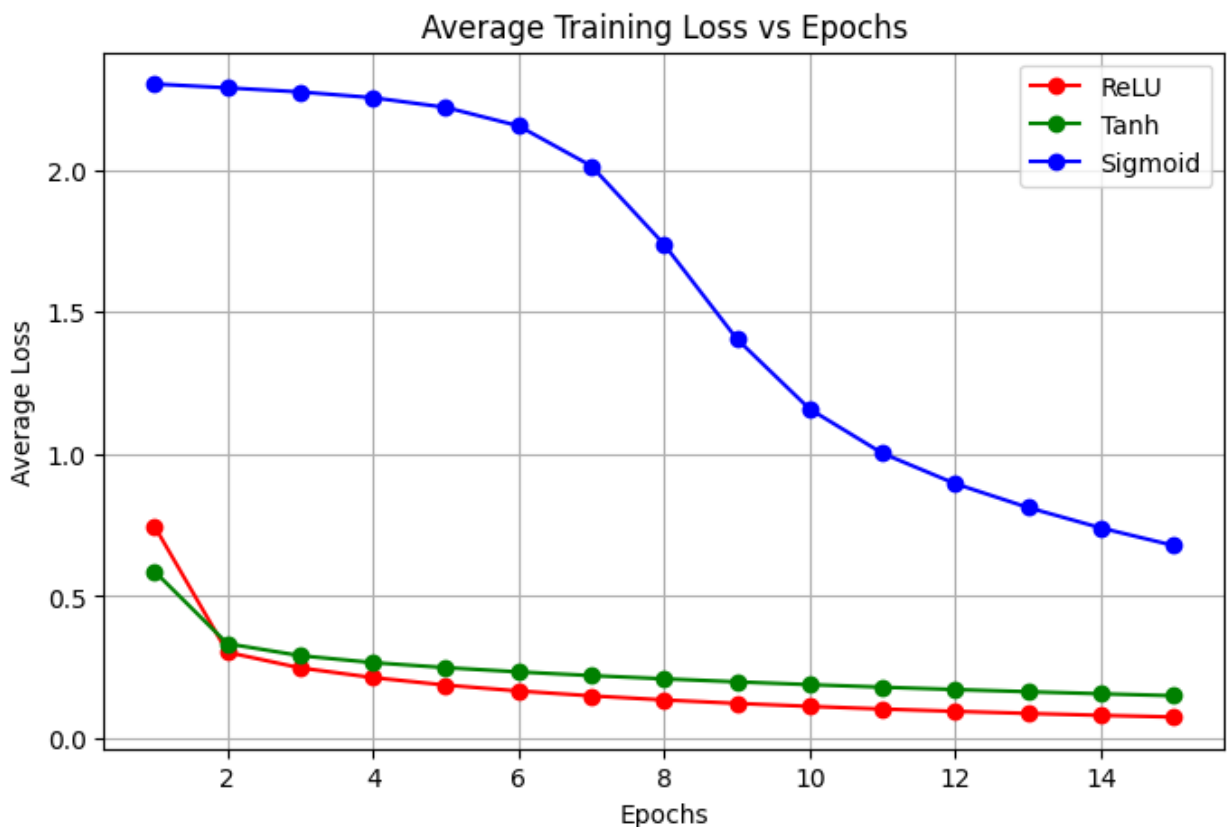
```

epochs = range(1, len(loss_history) + 1)

plt.figure(figsize=(8,5))
plt.plot(epochs, loss_history_ReLu, marker='o', label="ReLU",
color='r')
plt.plot(epochs, loss_history_tanh, marker='o', label="Tanh",
color='g')
plt.plot(epochs, loss_history, marker='o', label="Sigmoid", color='b')

plt.title("Average Training Loss vs Epochs")
plt.xlabel("Epochs")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True)
plt.show()

```



```

plt.figure(figsize=(8,5))

# Plot each activation using its own iteration indices
plt.plot(iter_tanh, iter_loss, marker='o', color='b', label="Sigmoid")
plt.plot(iter_tanh, iter_loss_tanh, marker='o', color='g',
label="Tanh")
plt.plot(iter_ReLu, iter_loss_ReLu, marker='o', color='r',

```

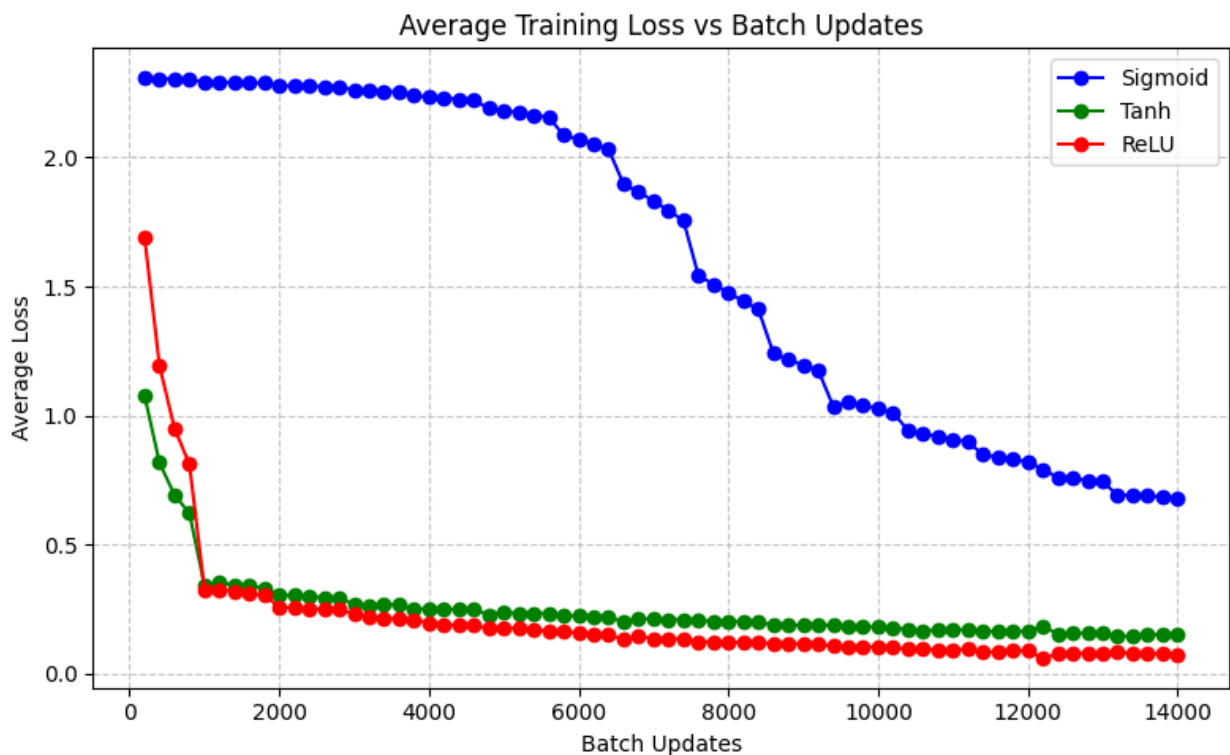


```

label="ReLU")

plt.title("Average Training Loss vs Batch Updates")
plt.xlabel("Batch Updates")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()

```



Comparison for the different activation functions

- Test accuracy for sigmoid activation function = 82.30%
- Test accuracy for ReLU activation function = 97.12%
- Test accuracy for Tan Hyperbolic activation function = 95.47%

Conclusions:-

- The Tanh and ReLU activation functions perform almost identical.
- The Sigmoid activation function struggles to deliver good training and testing processes.

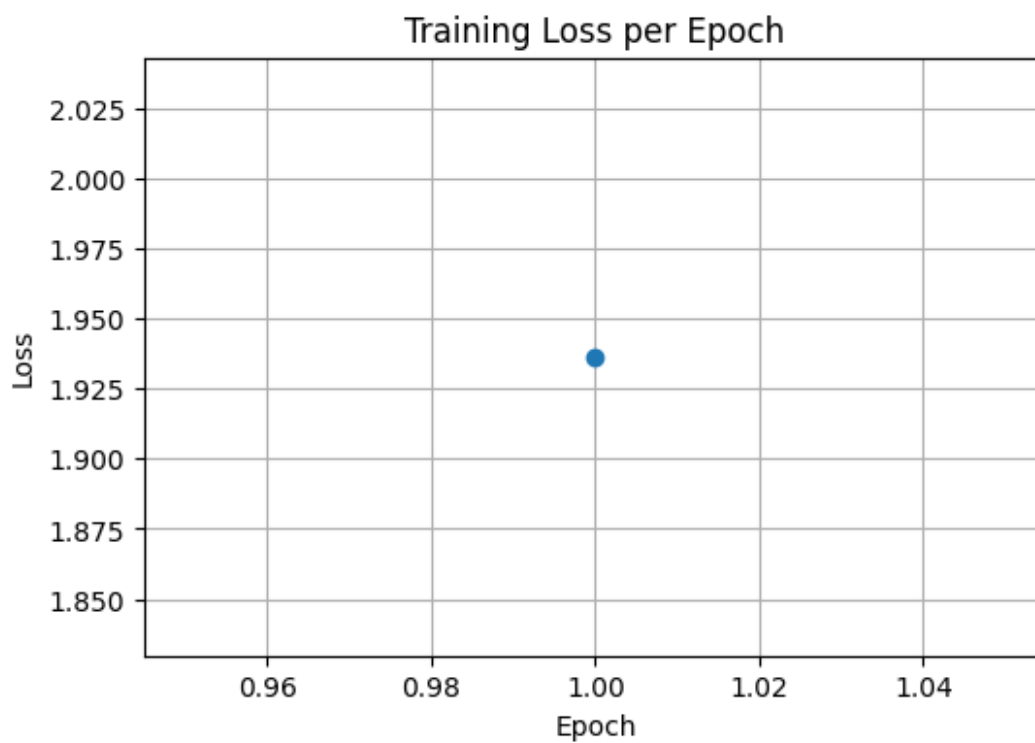
- The ReLU activation function performs relatively better.

Analyzing the MLP algorithm for different learning rates

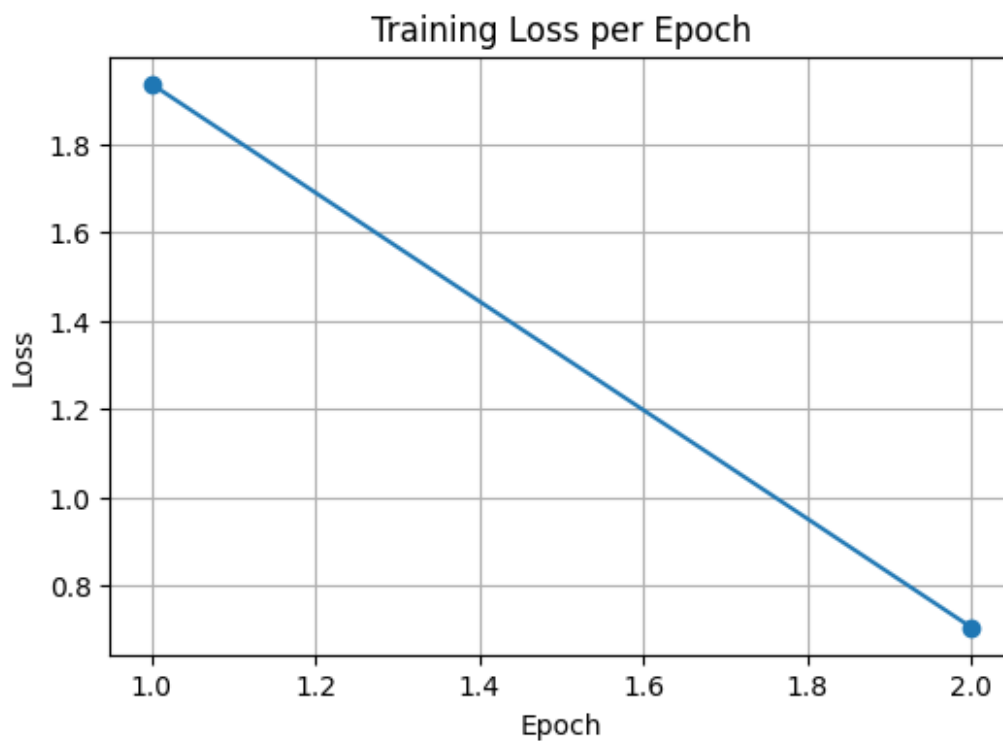
- Learning rate = 0.1

```
model_1 = MLP()  
loss_history_sigmoid_1, iter_sigmoid_1, iter_loss_sigmoid_1 =  
train(model_1, train_loader, 15, 0.1)
```

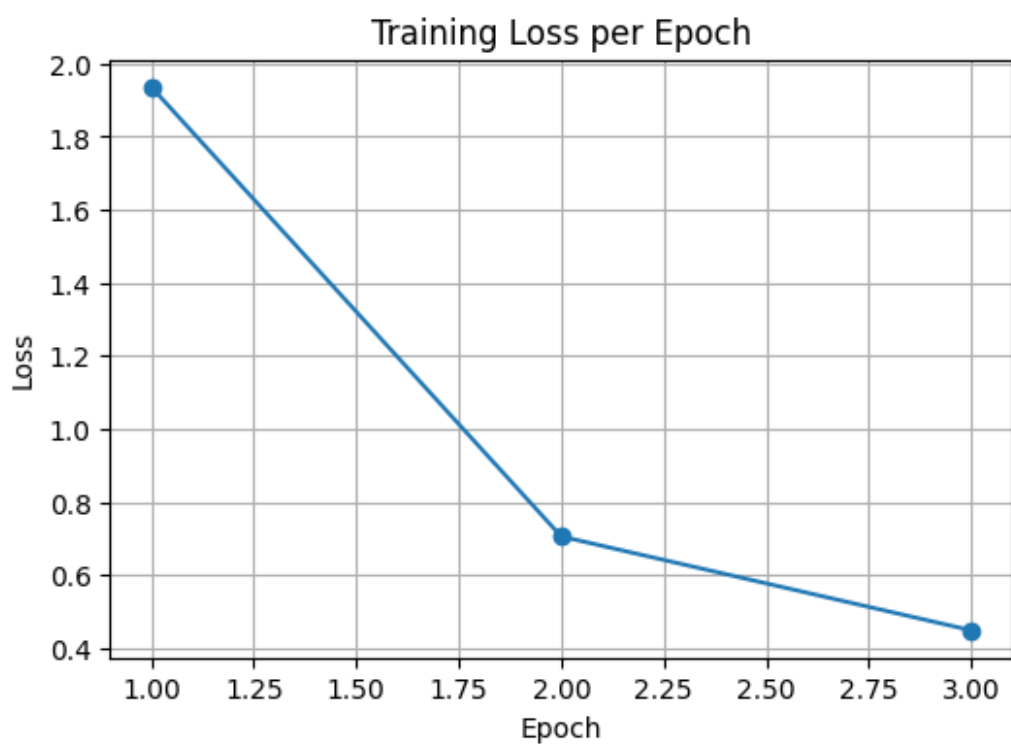
Epoch 1/15, Loss: 1.9364



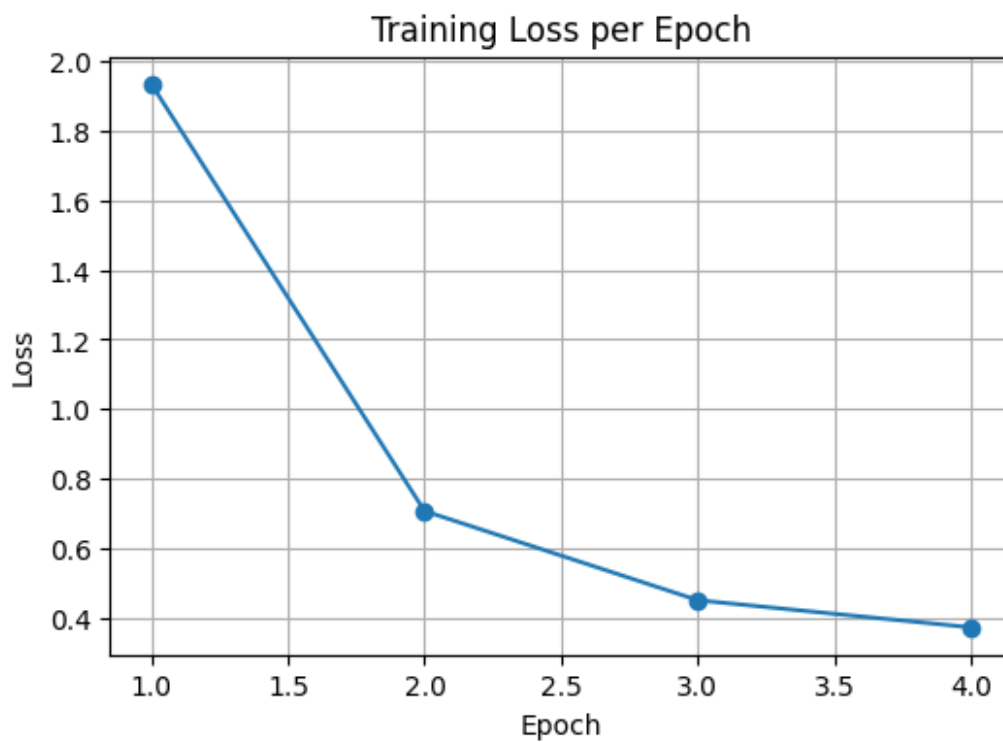
Epoch 2/15, Loss: 0.7062



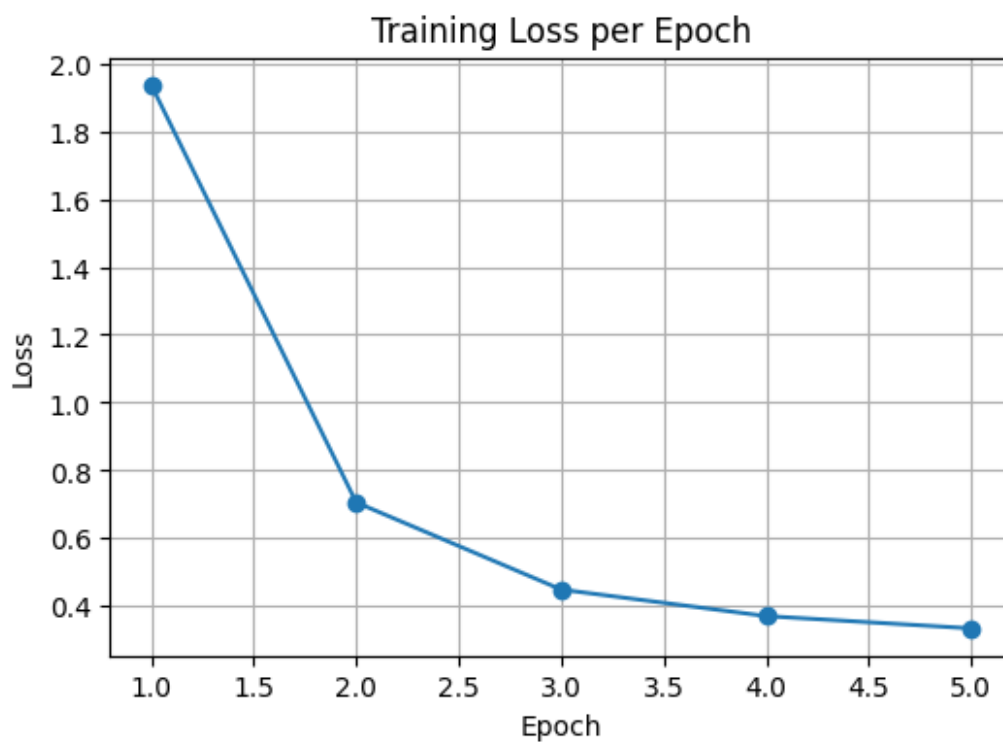
Epoch 3/15, Loss: 0.4492



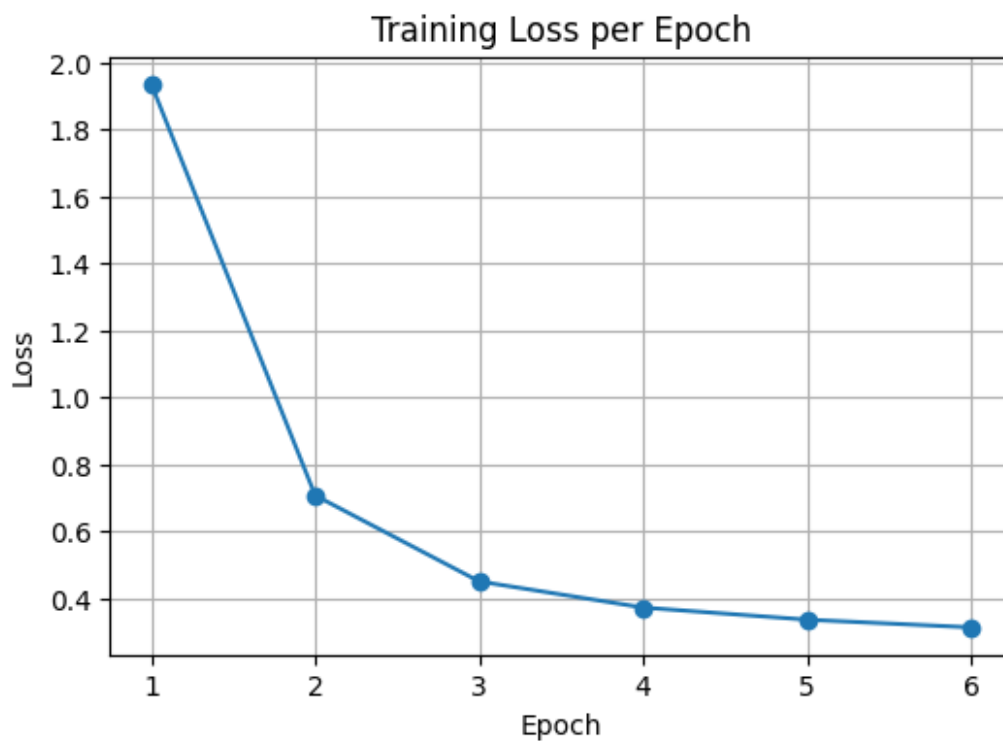
Epoch 4/15, Loss: 0.3712



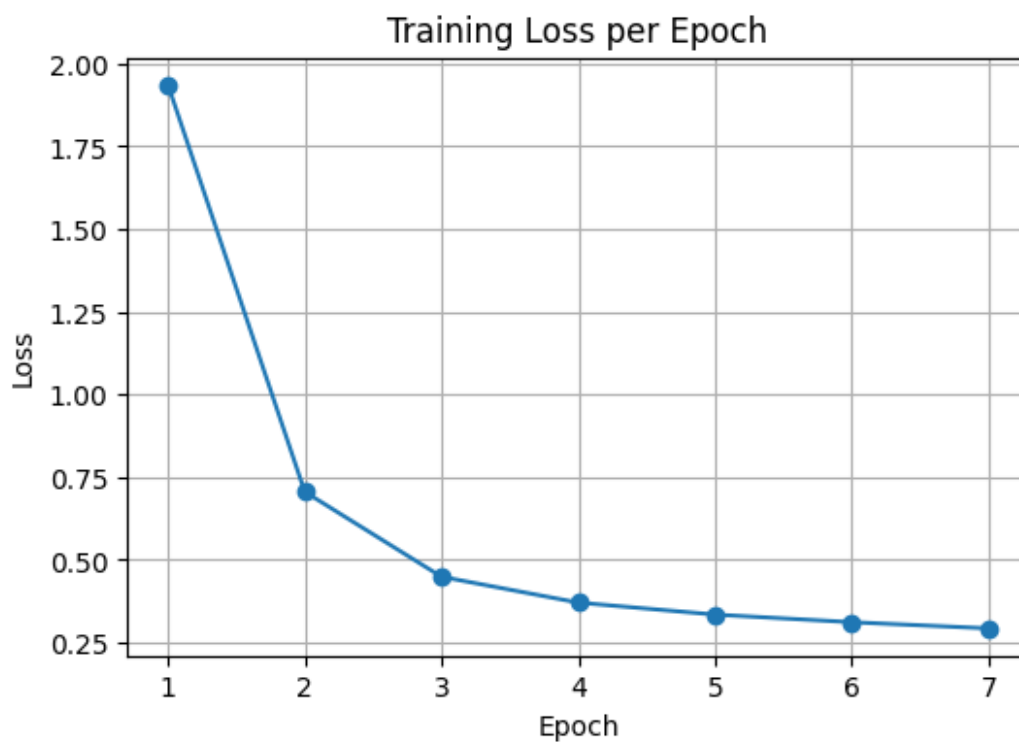
Epoch 5/15, Loss: 0.3349



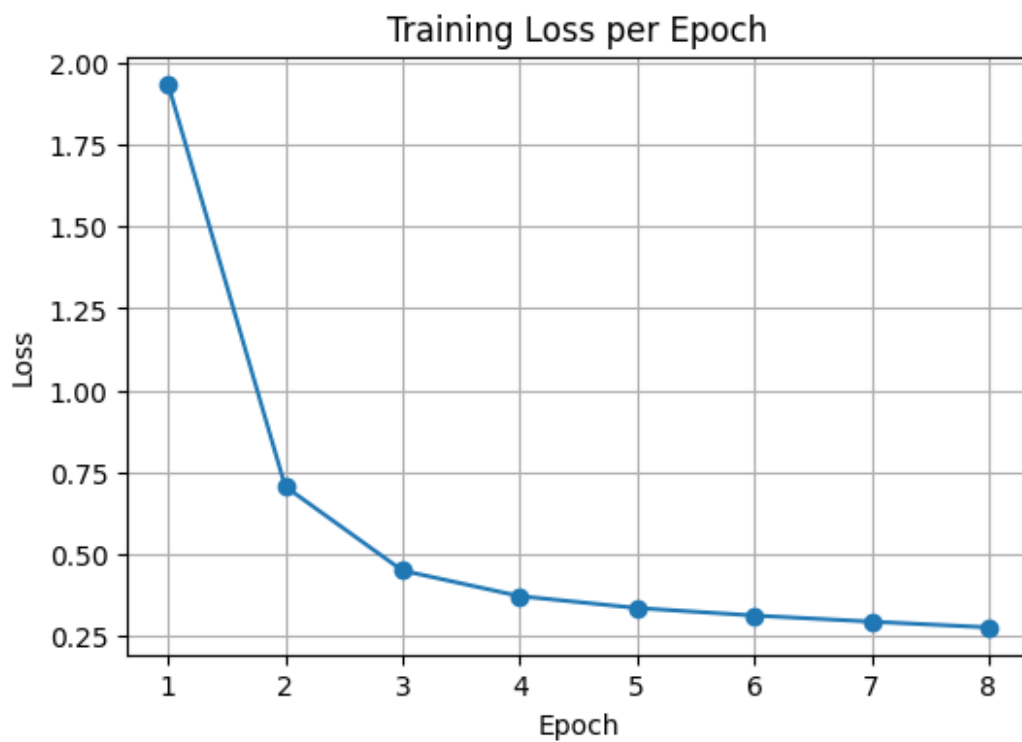
Epoch 6/15, Loss: 0.3120



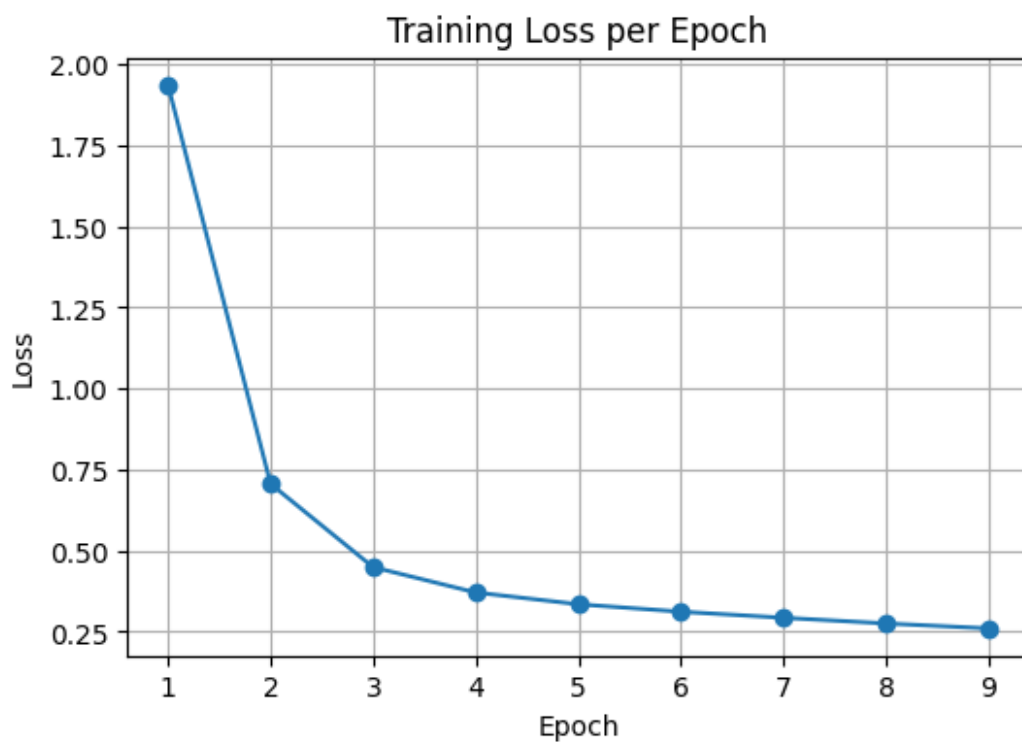
Epoch 7/15, Loss: 0.2933



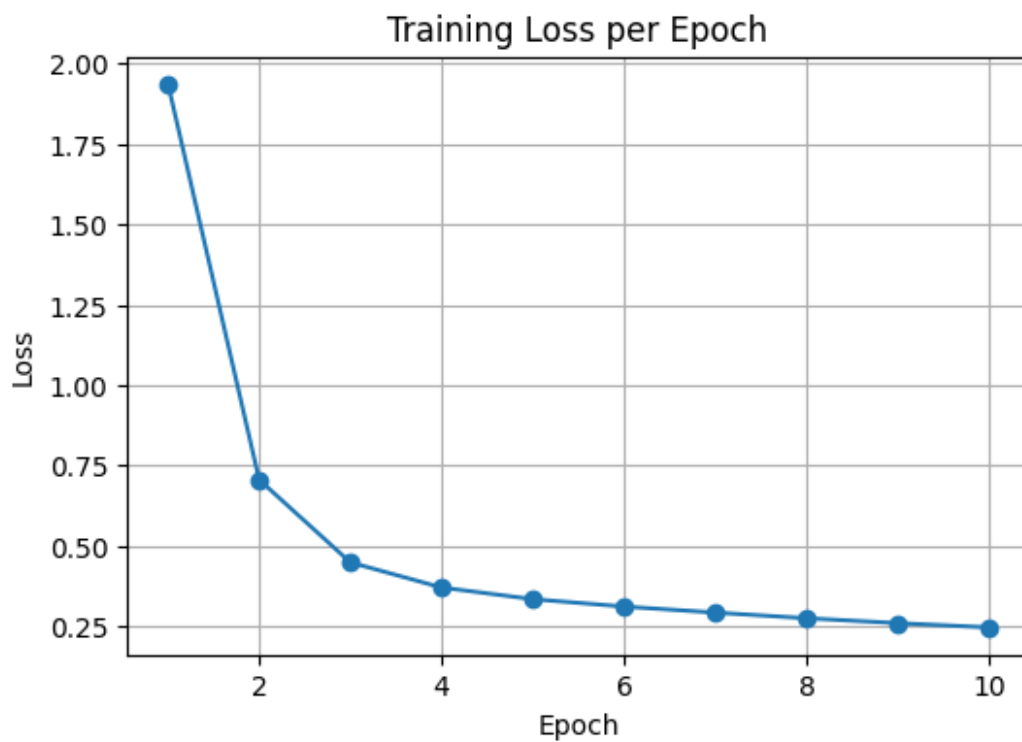
Epoch 8/15, Loss: 0.2757



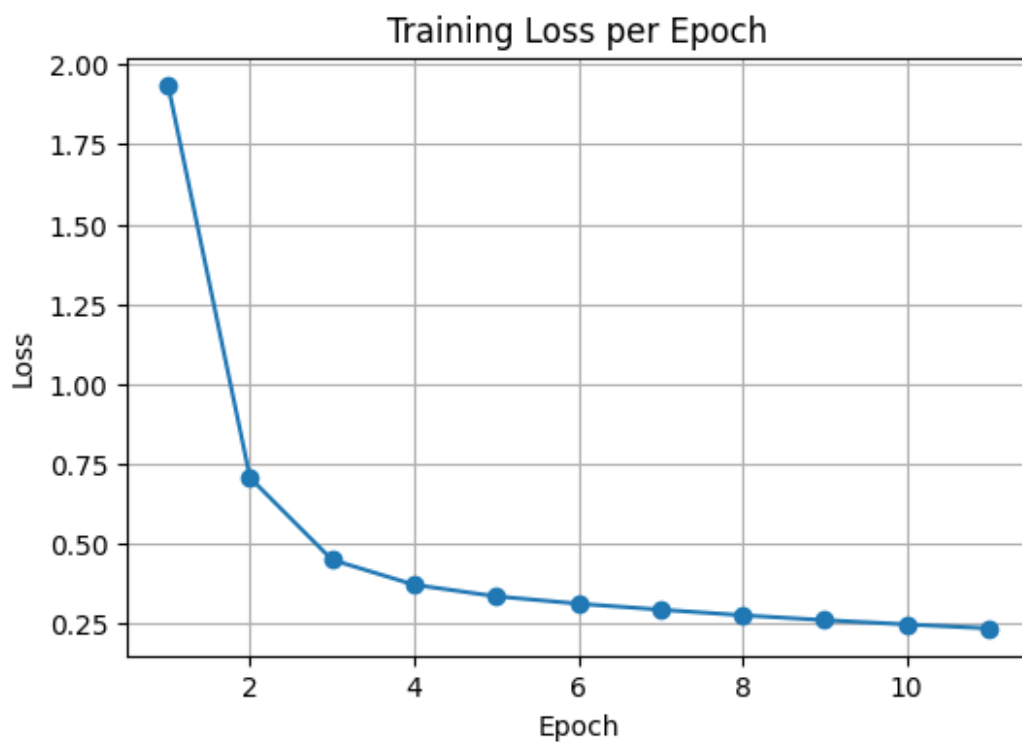
Epoch 9/15, Loss: 0.2606



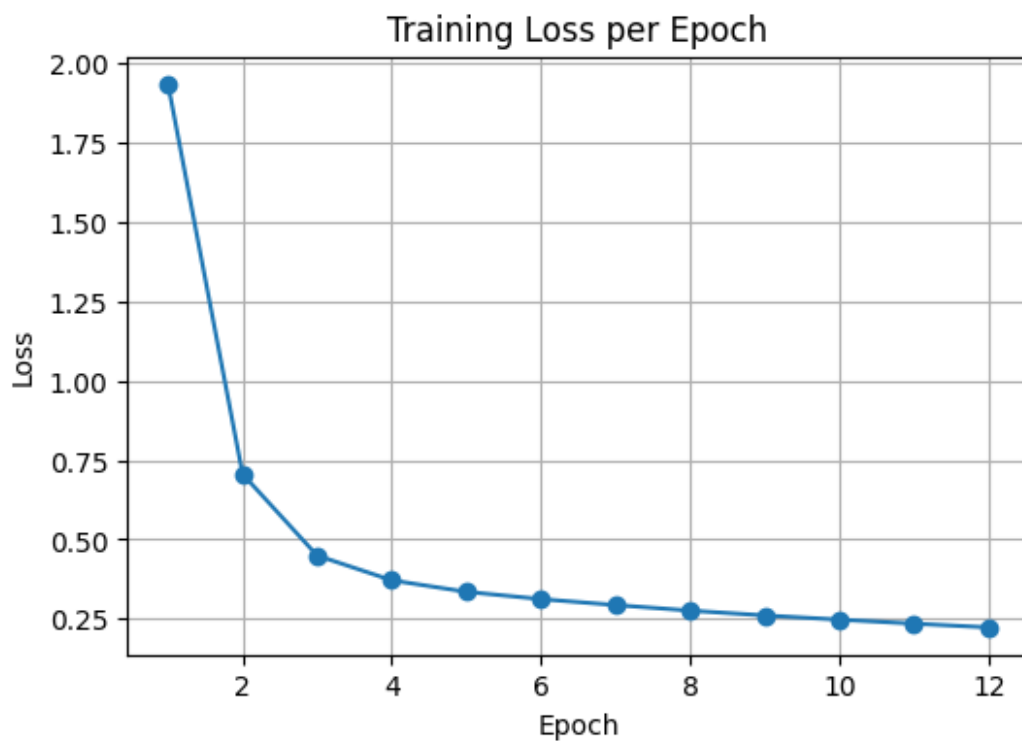
Epoch 10/15, Loss: 0.2473



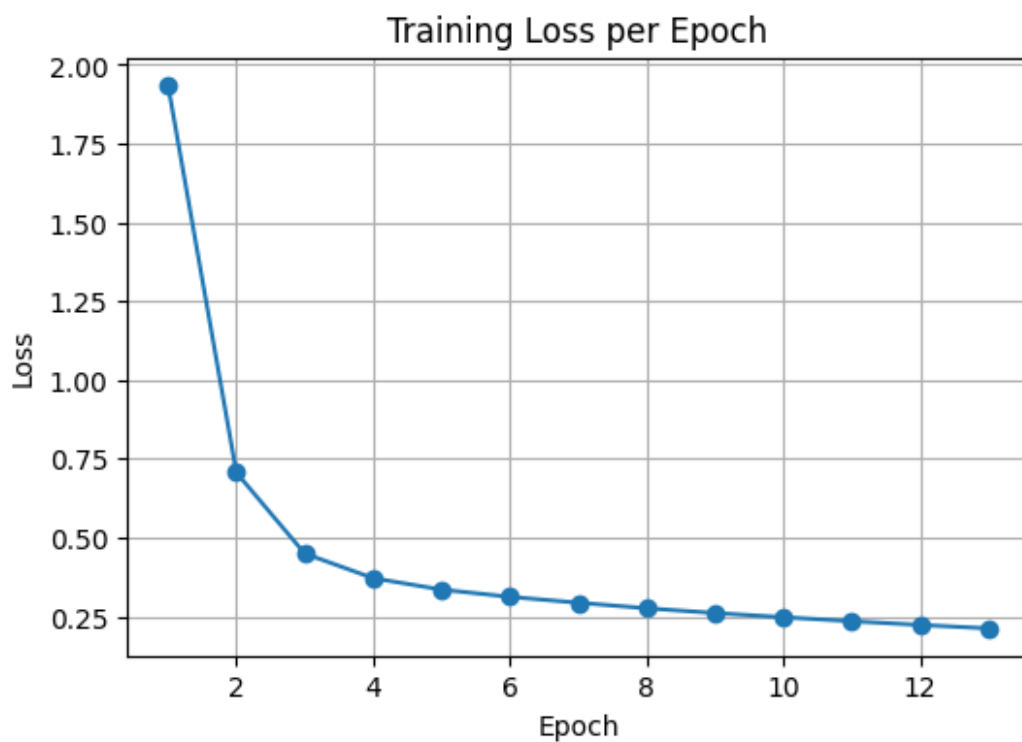
Epoch 11/15, Loss: 0.2347



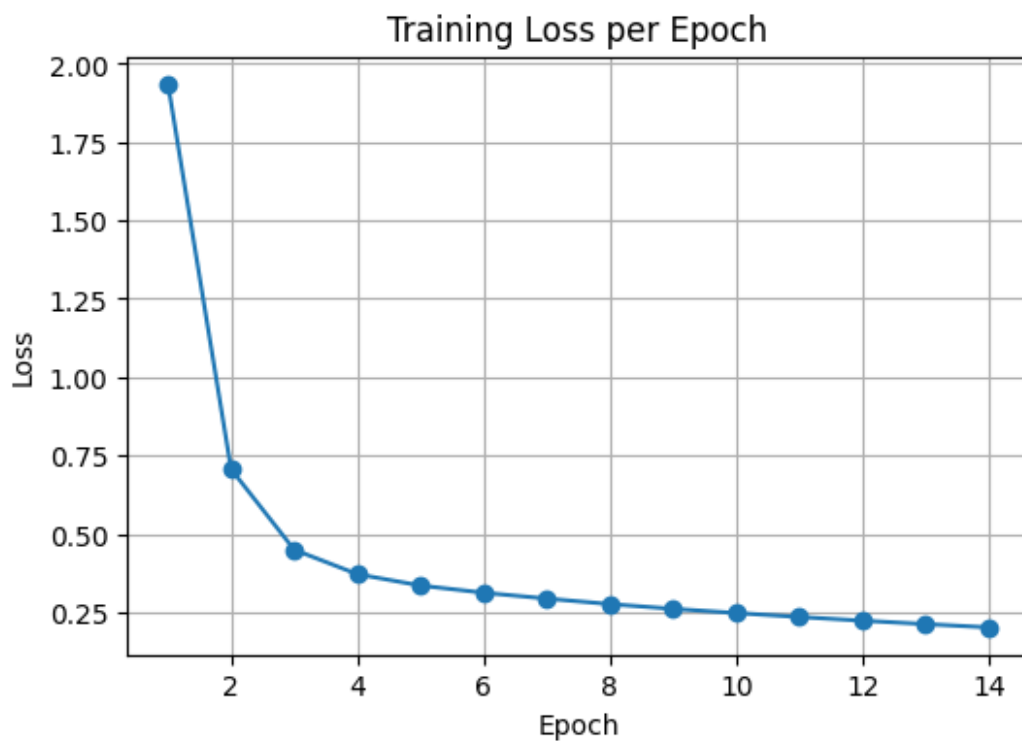
Epoch 12/15, Loss: 0.2229



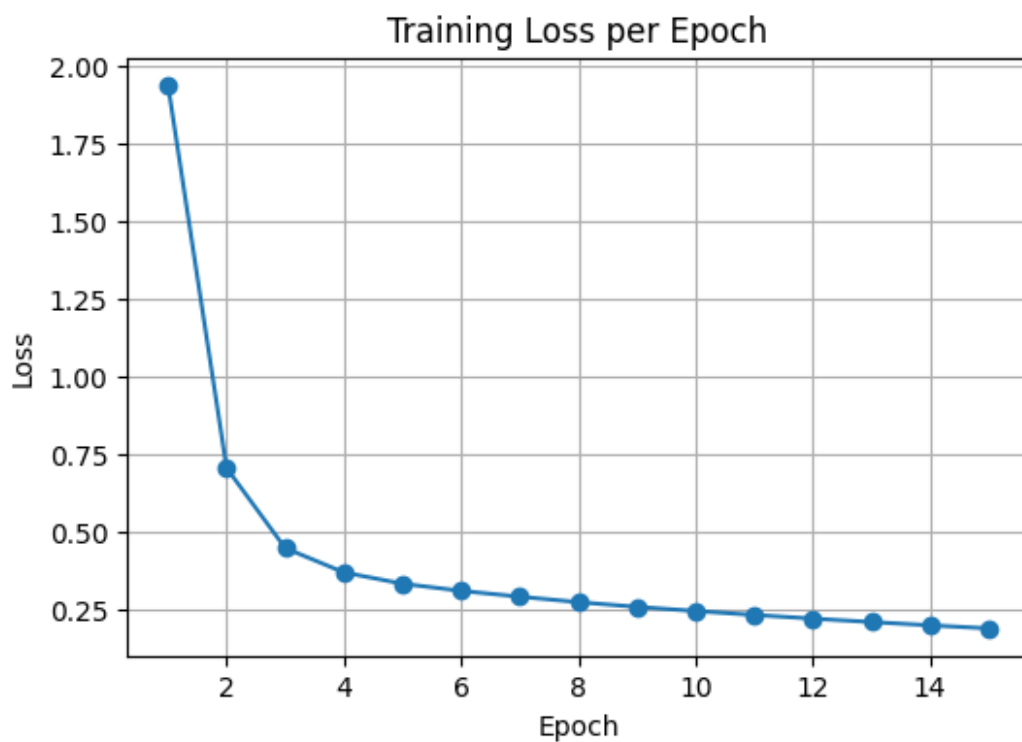
Epoch 13/15, Loss: 0.2119

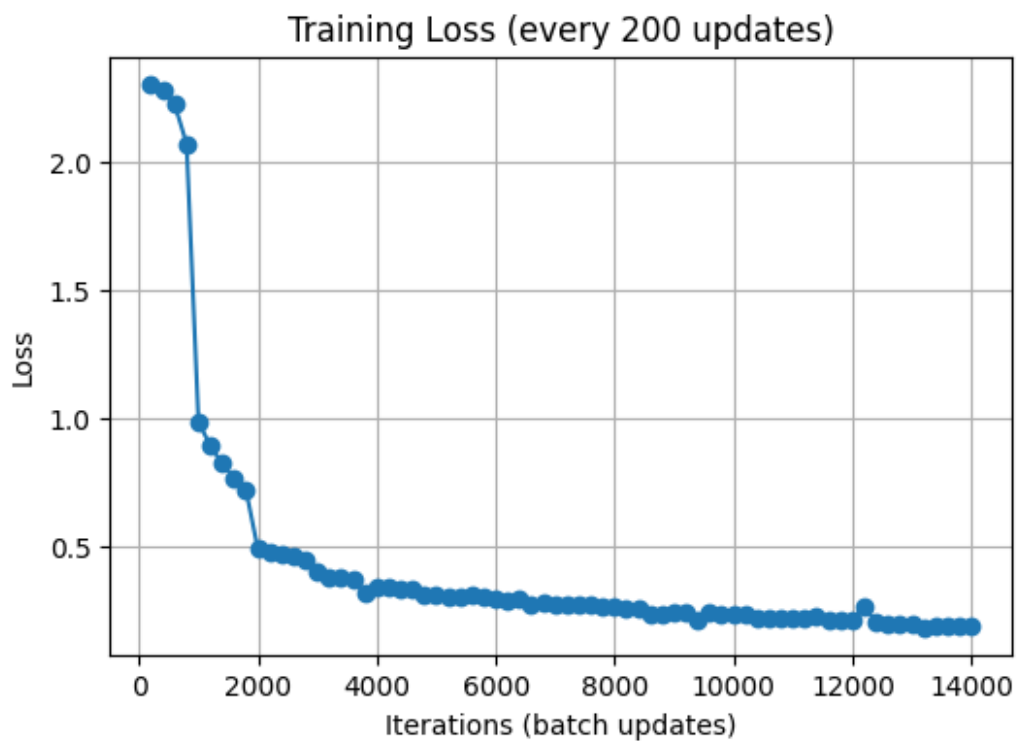


Epoch 14/15, Loss: 0.2014

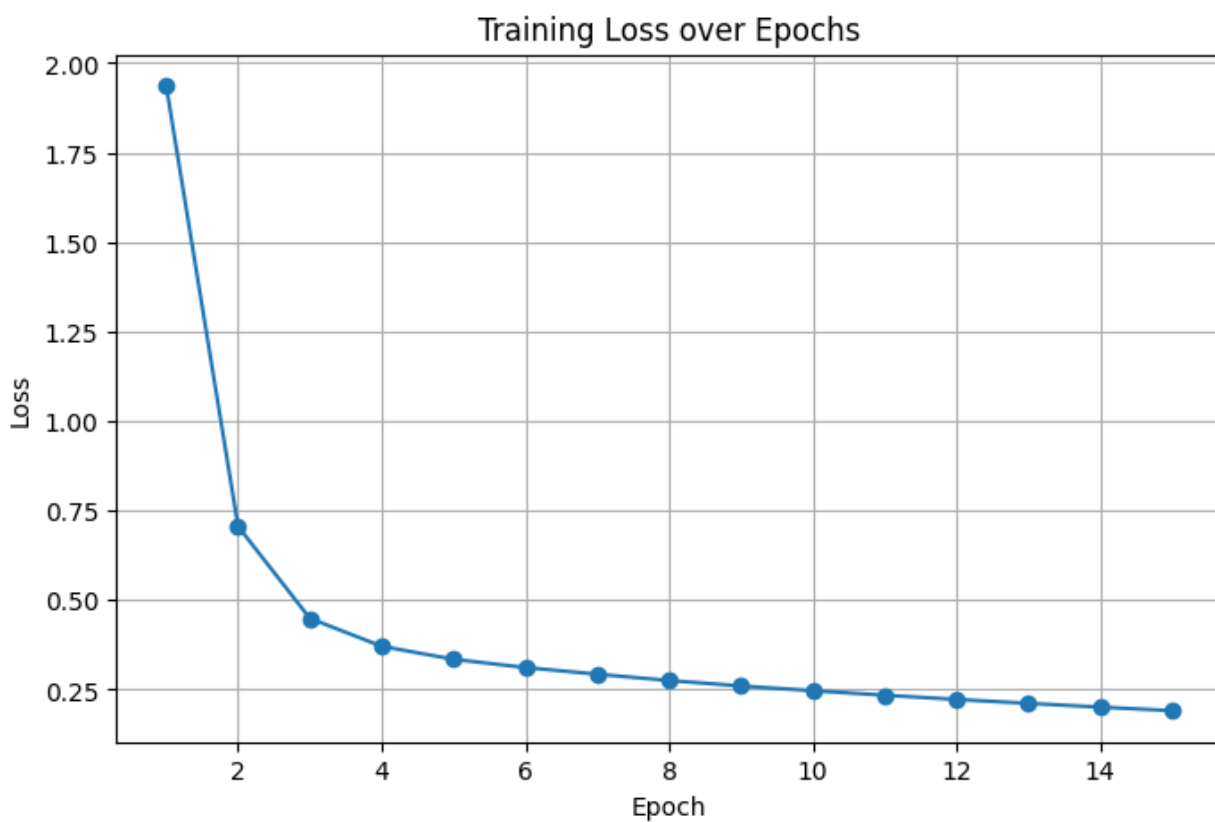


Epoch 15/15, Loss: 0.1915





```
plot_training_curves(loss_history_sigmoid_1)
```



```
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("Test Set Evaluation for Learning rate = 0.1")
evaluate(model_1, X_test_np, y_test_np)
```

Test Set Evaluation for Learning rate = 0.1

Accuracy: 94.2500%

Confusion Matrix:

```
[[ 963    0    1    1    0    8    5    2    0    0]
 [    0 1113    2    2    0    2    4    2   10    0]
 [   12    2   965    6    4    4   13   11   13    2]
 [    0    0   15   953    0   25    1    8    6    2]
 [    3    2    4    0   906    2   12    2    4   47]
 [    7    1    0   16    2  837   11    1   13    4]
 [   12    3    2    1    7   23  907    0    3    0]
 [    3    8   18    8    1    2    0  972    0   16]
 [    6    4    4   20    5   49    8   10  865    3]
 [   10    5    1    8   15    8    1   13    4  944]]
```

Classification Report:

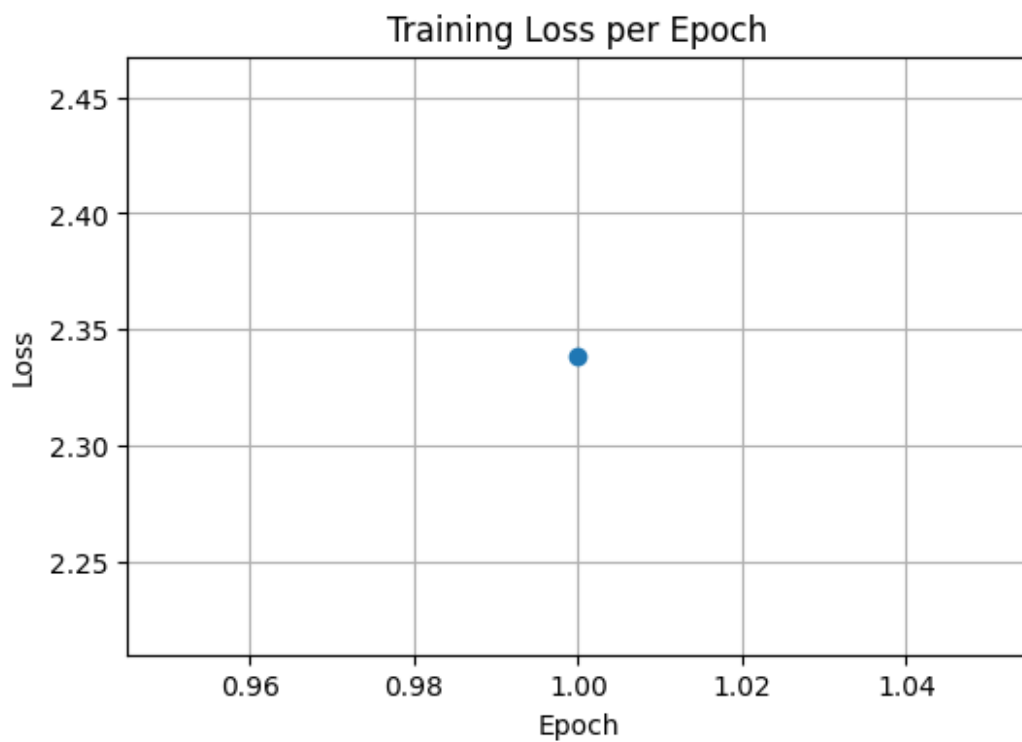
	precision	recall	f1-score	support
0	0.95	0.98	0.96	980
1	0.98	0.98	0.98	1135
2	0.95	0.94	0.94	1032
3	0.94	0.94	0.94	1010
4	0.96	0.92	0.94	982
5	0.87	0.94	0.90	892
6	0.94	0.95	0.94	958
7	0.95	0.95	0.95	1028
8	0.94	0.89	0.91	974
9	0.93	0.94	0.93	1009
accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000

```
np.float64(0.9425)
```

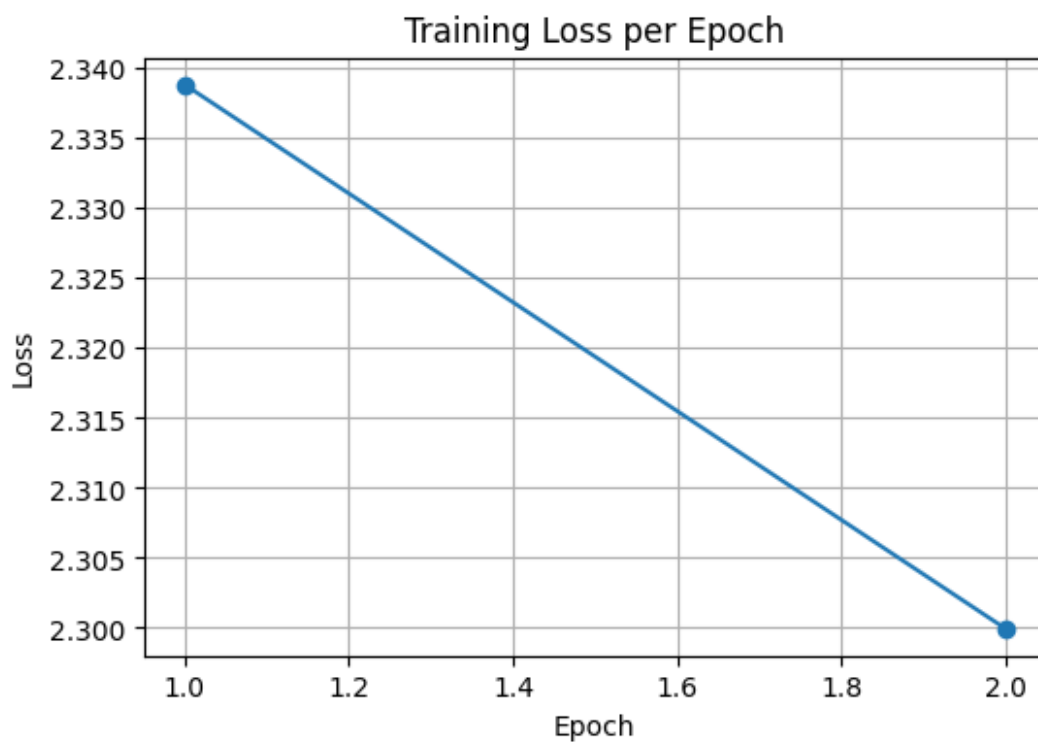
- **Learning rate = 0.001**

```
model_2 = MLP()
loss_history_sigmoid_2, iter_sigmoid_2, iter_loss_sigmoid_2 =
train(model_2, train_loader, 15, 0.001)
```

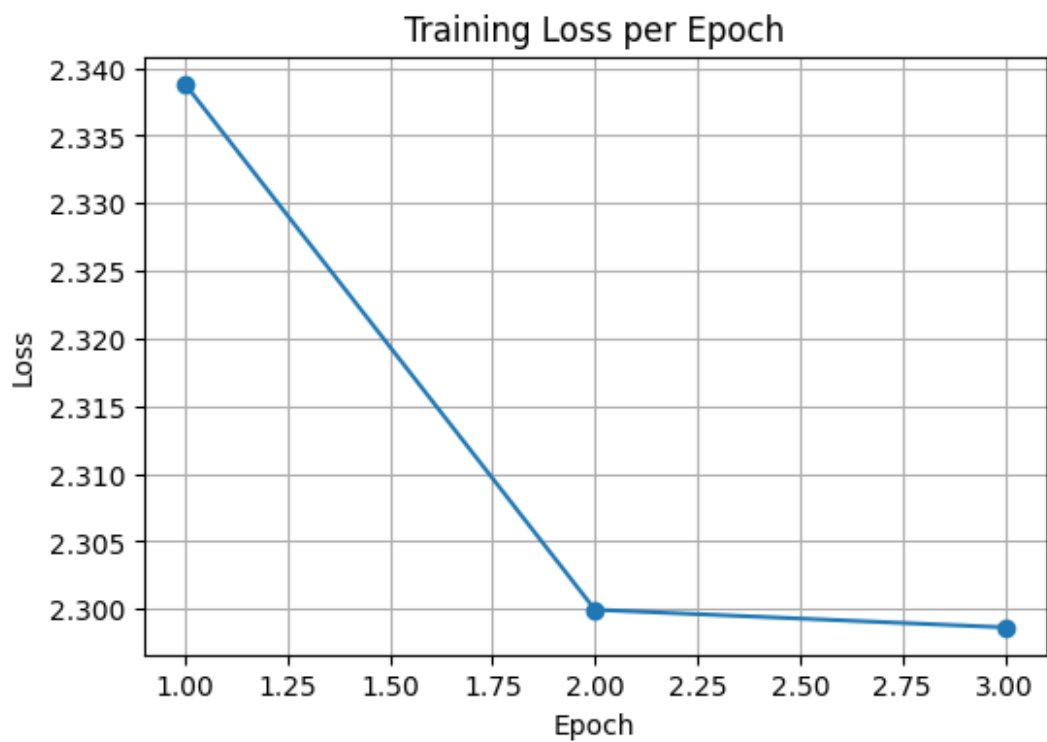
Epoch 1/15, Loss: 2.3388



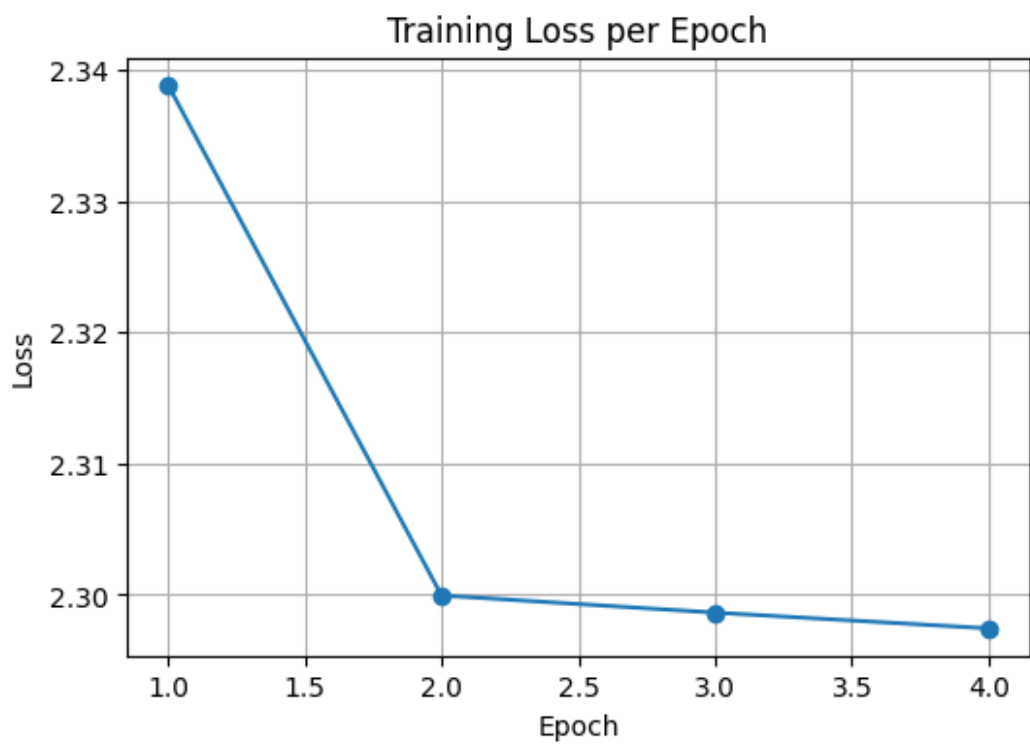
Epoch 2/15, Loss: 2.3000



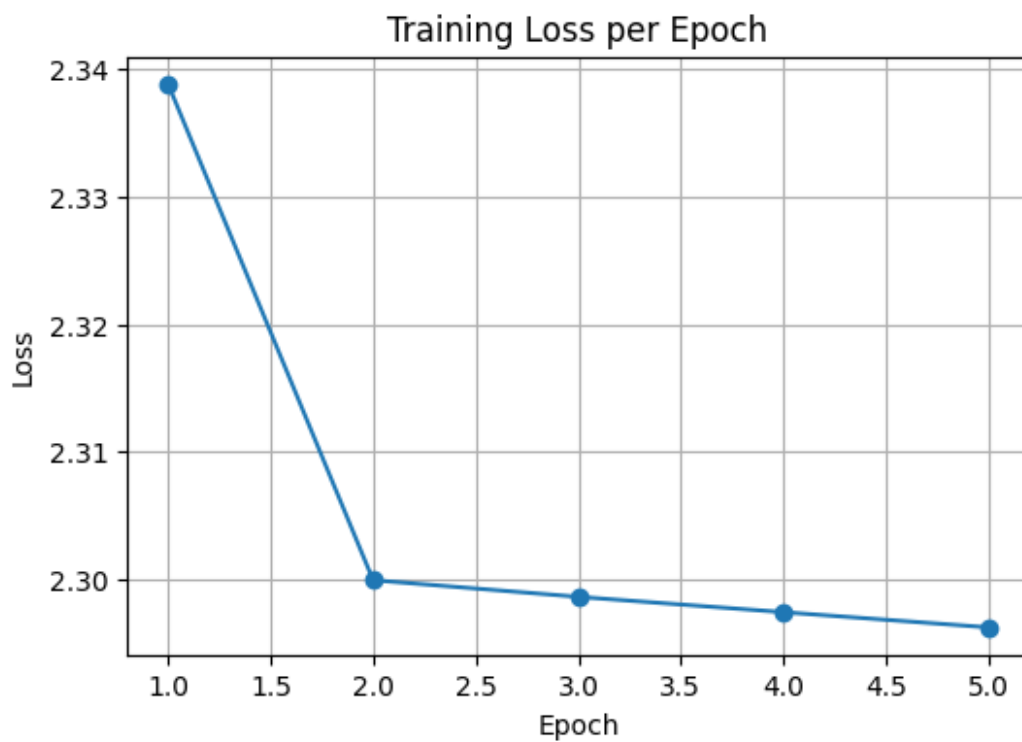
Epoch 3/15, Loss: 2.2987



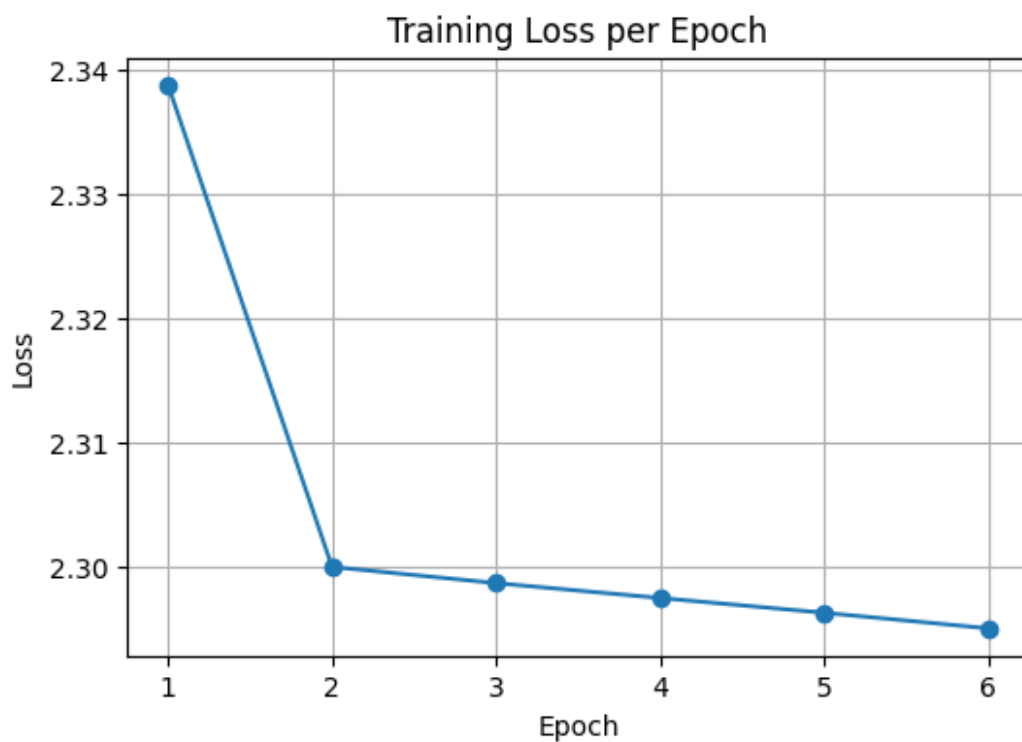
Epoch 4/15, Loss: 2.2975



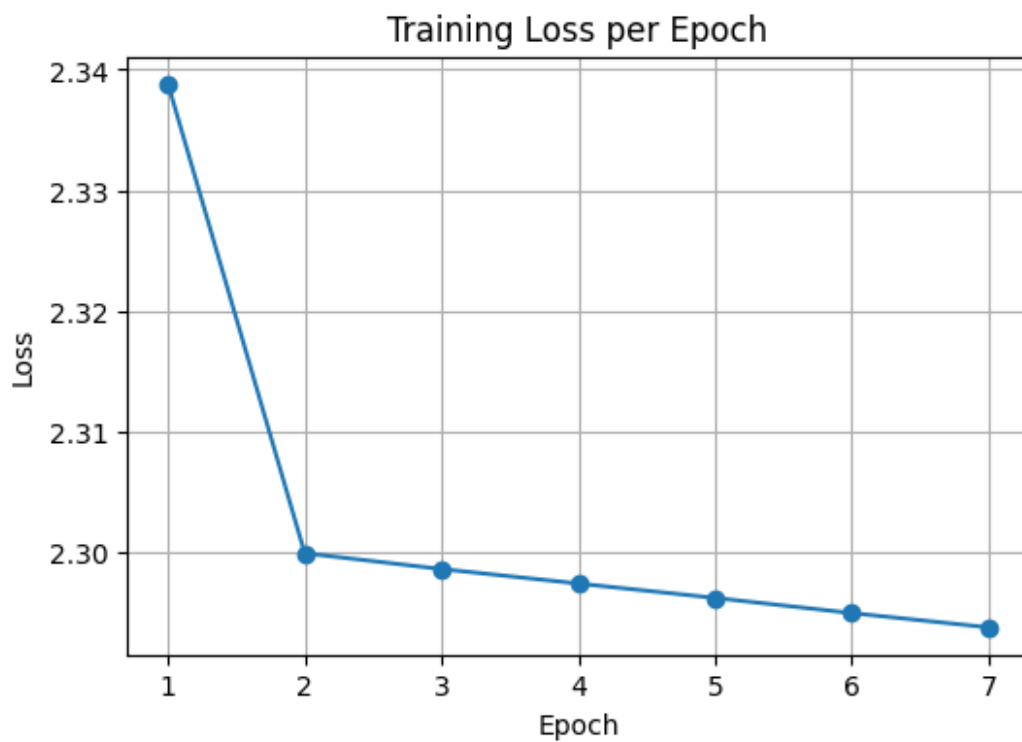
Epoch 5/15, Loss: 2.2963



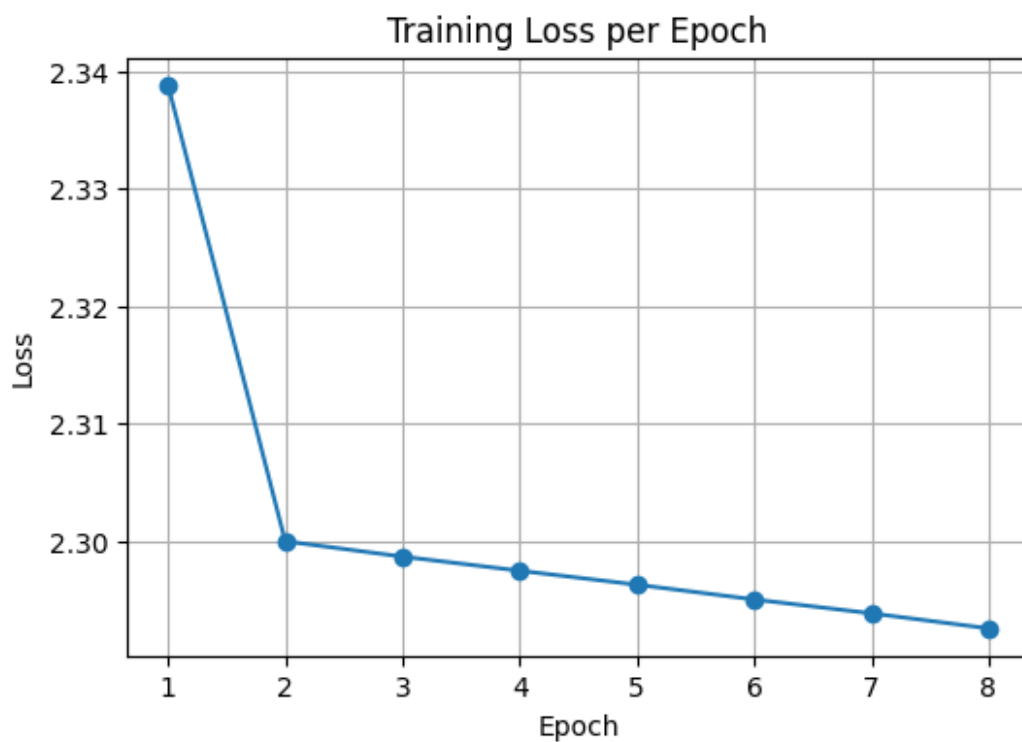
Epoch 6/15, Loss: 2.2950



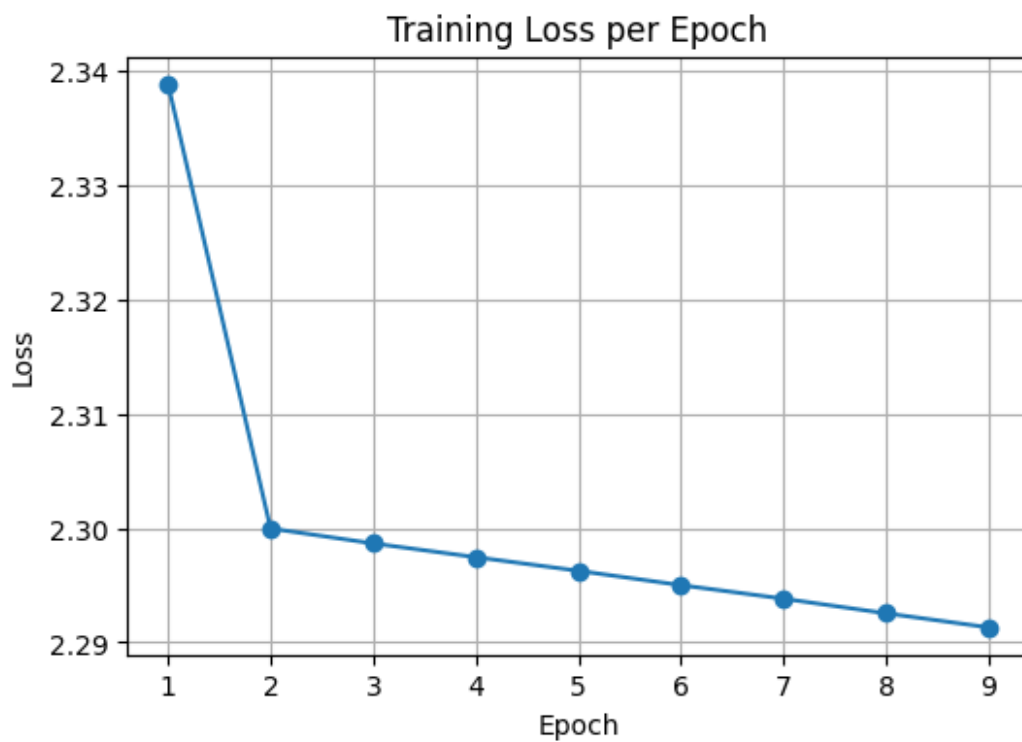
Epoch 7/15, Loss: 2.2938



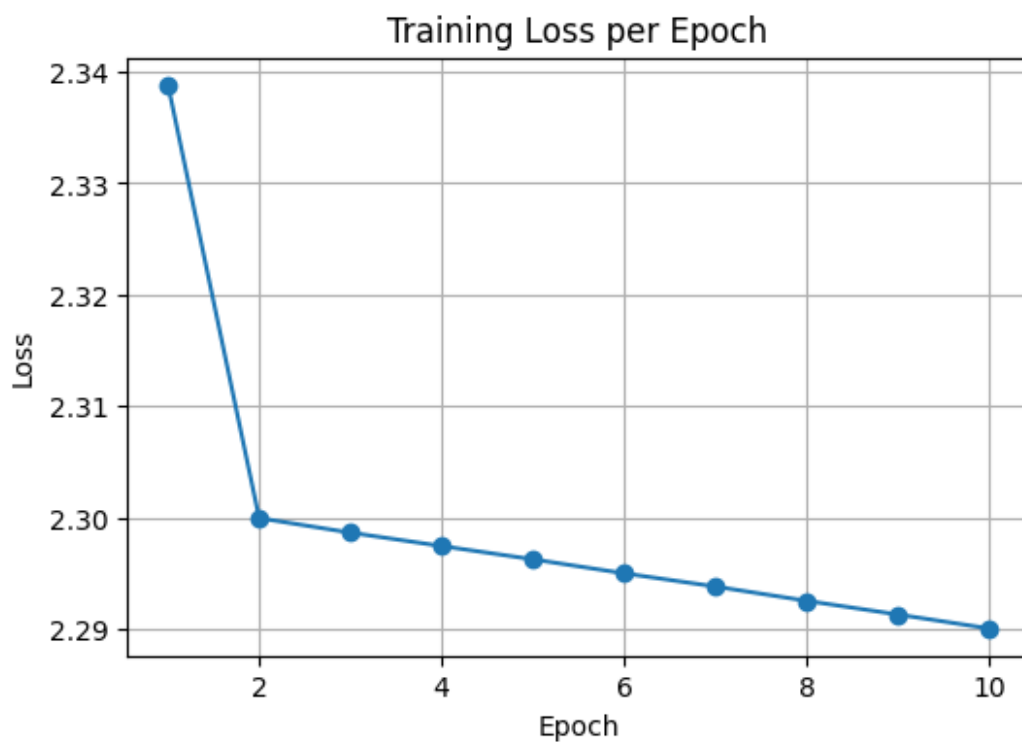
Epoch 8/15, Loss: 2.2925



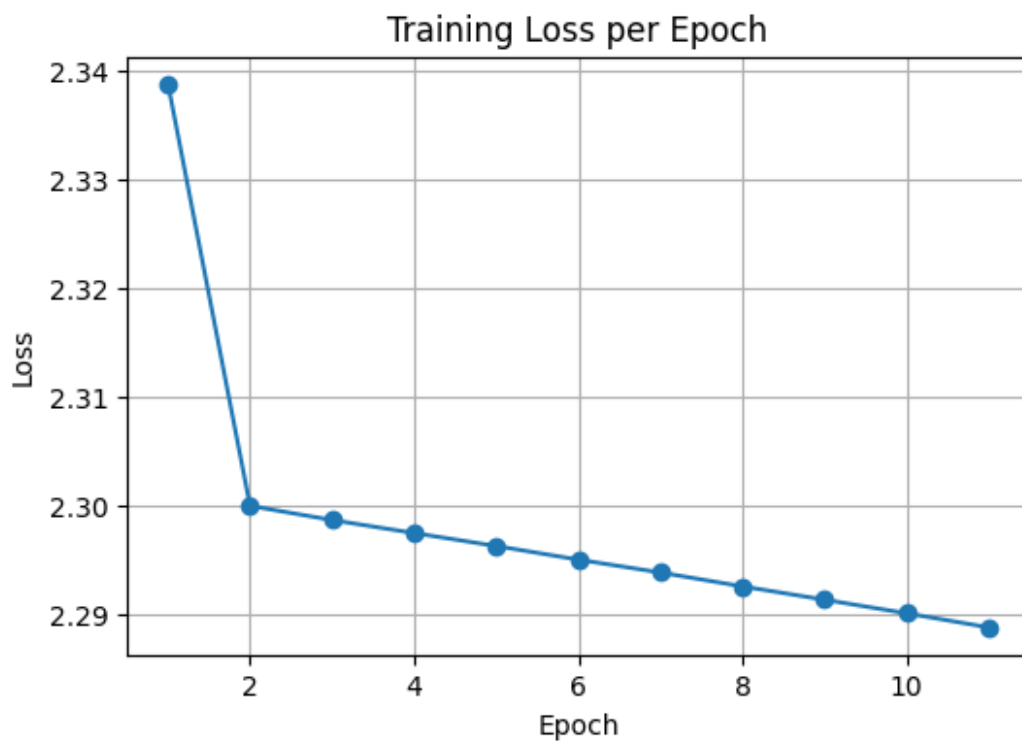
Epoch 9/15, Loss: 2.2913



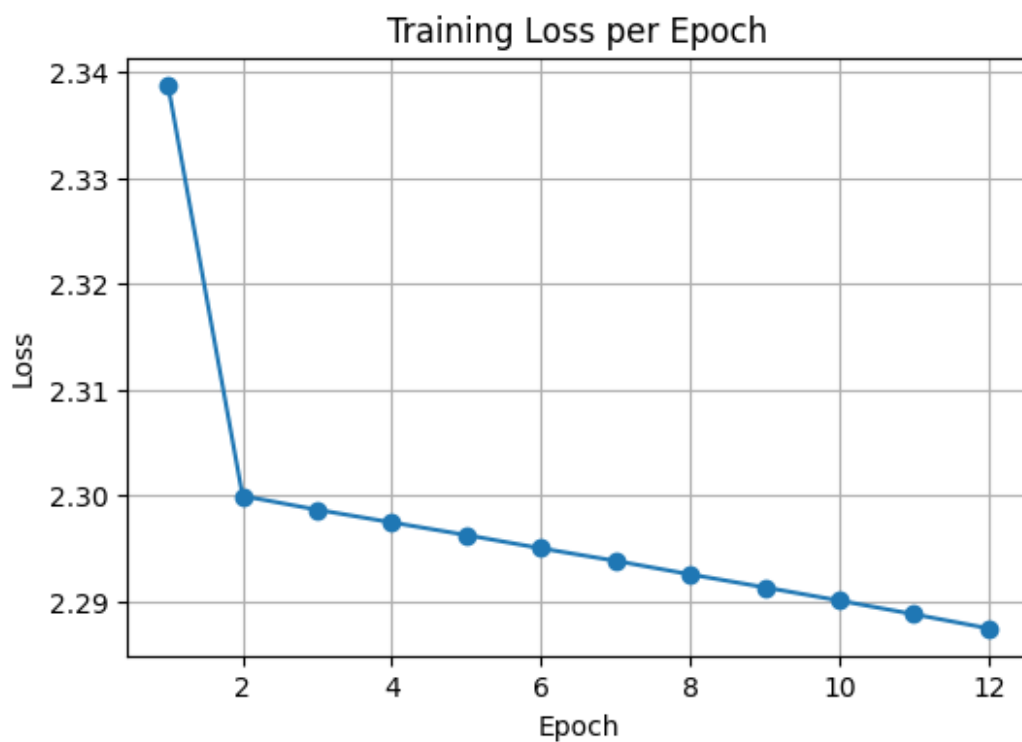
Epoch 10/15, Loss: 2.2901



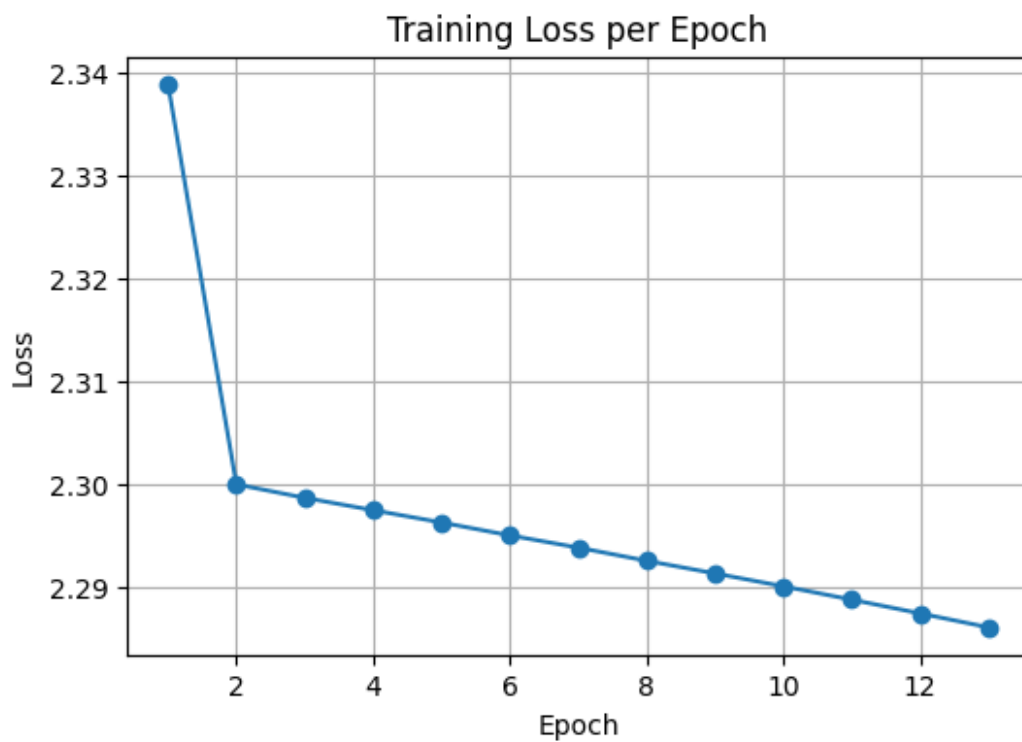
Epoch 11/15, Loss: 2.2888



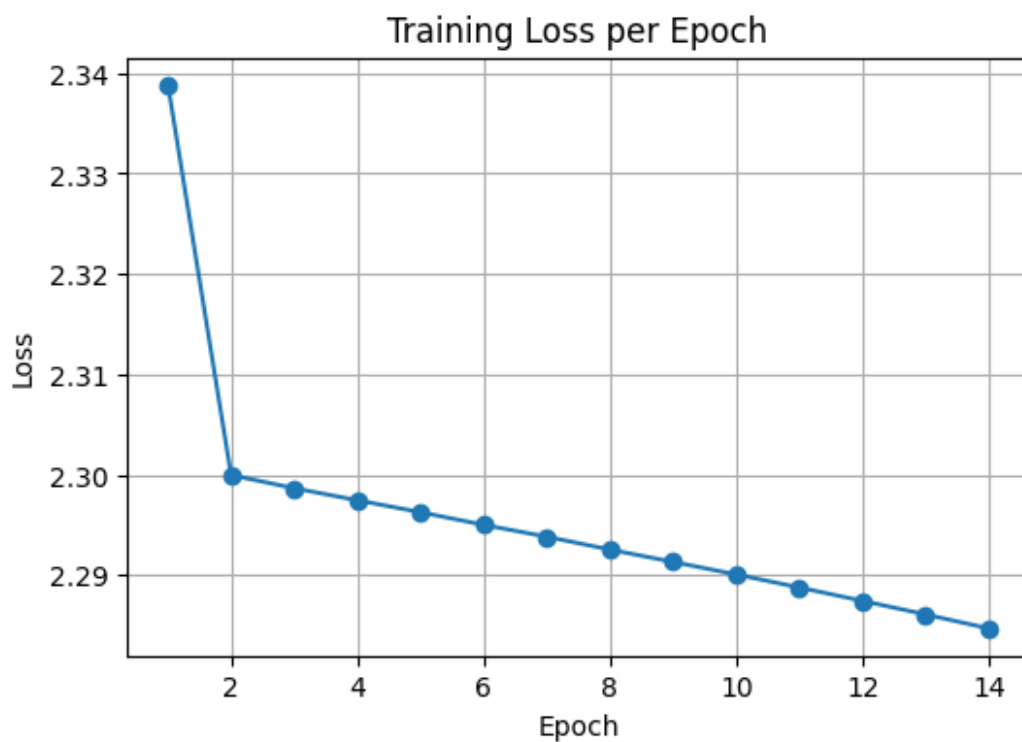
Epoch 12/15, Loss: 2.2874



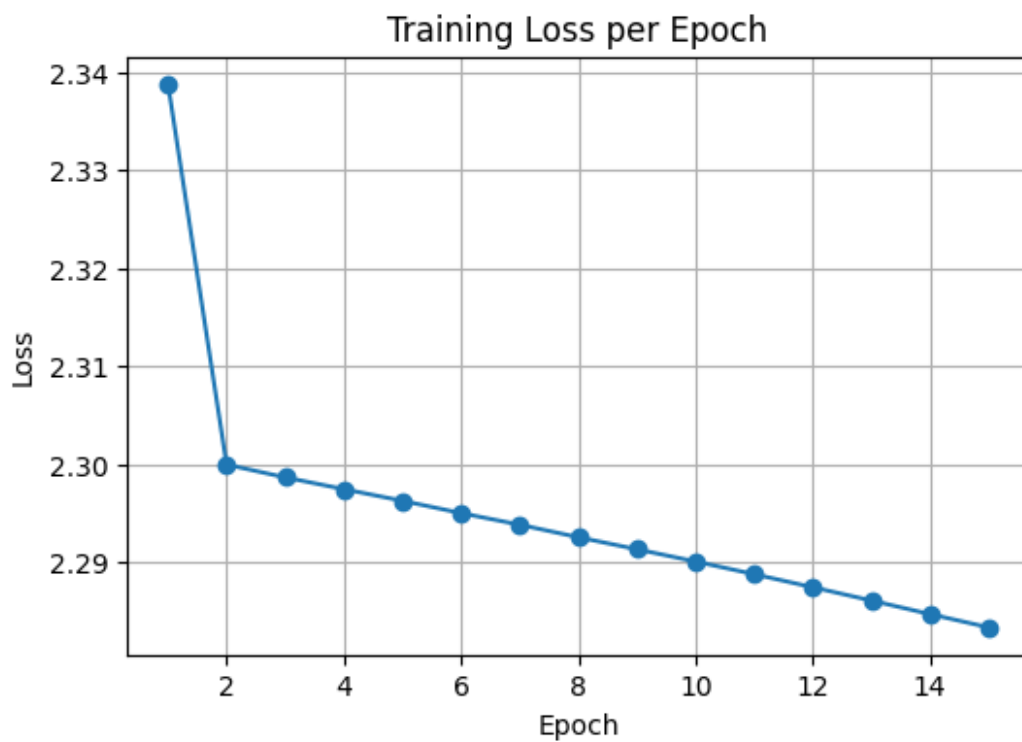
Epoch 13/15, Loss: 2.2861



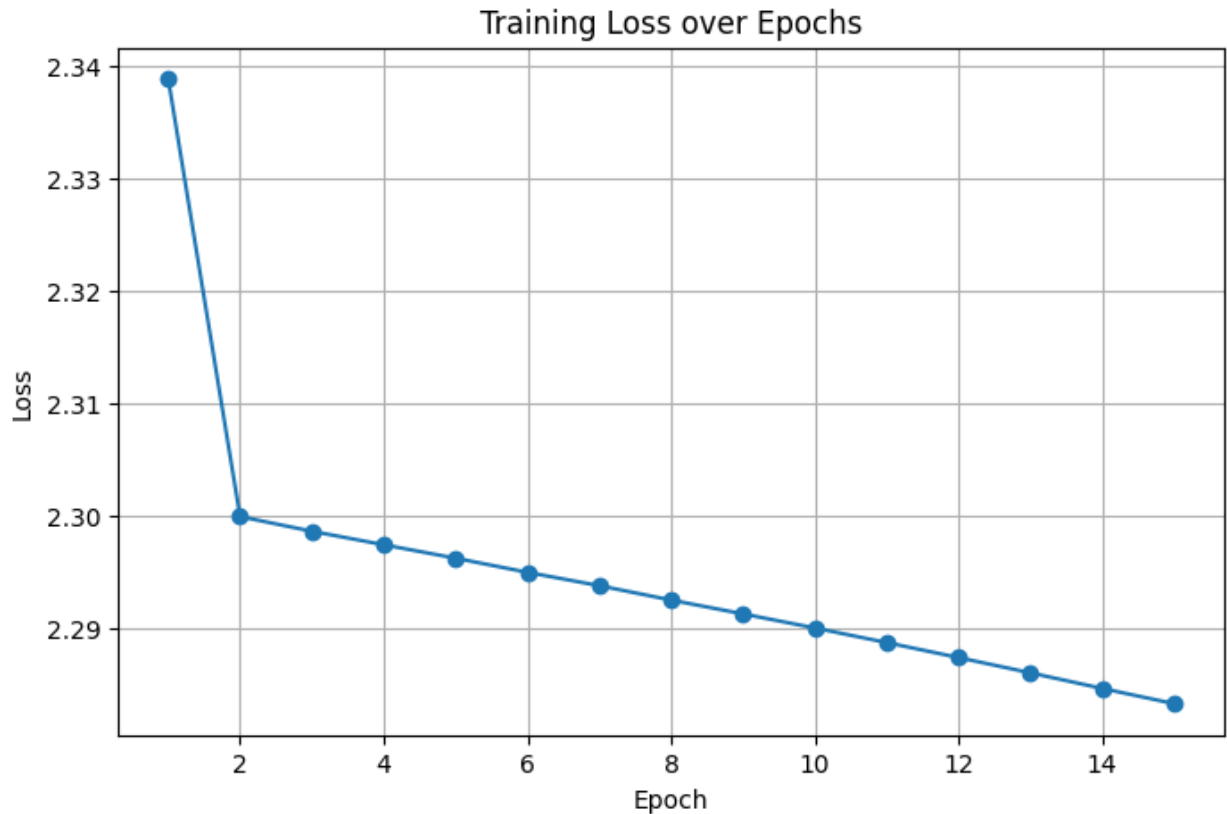
Epoch 14/15, Loss: 2.2847



Epoch 15/15, Loss: 2.2833



```
plot_training_curves(loss_history_sigmoid_2)
```



```
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("Test Set Evaluation for Learning rate = 0.001")
evaluate(model_2, X_test_np, y_test_np)
```

Test Set Evaluation for Learning rate = 0.001

Accuracy: 11.5600%

Confusion Matrix:

```
[ [ 19 961  0  0  0  0  0  0  0  0]
 [  0 1135  0  0  0  0  0  0  0  0]
 [  0 1032  0  0  0  0  0  0  0  0]
 [  0 1008  0  2  0  0  0  0  0  0]
 [  0  982  0  0  0  0  0  0  0  0]
 [  0  892  0  0  0  0  0  0  0  0]
 [  0  958  0  0  0  0  0  0  0  0]
 [  0 1028  0  0  0  0  0  0  0  0]
 [  0  974  0  0  0  0  0  0  0  0]
 [  0 1009  0  0  0  0  0  0  0  0]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.02	0.04	980
1	0.11	1.00	0.20	1135

2	0.00	0.00	0.00	1032
3	1.00	0.00	0.00	1010
4	0.00	0.00	0.00	982
5	0.00	0.00	0.00	892
6	0.00	0.00	0.00	958
7	0.00	0.00	0.00	1028
8	0.00	0.00	0.00	974
9	0.00	0.00	0.00	1009
accuracy			0.12	10000
macro avg	0.21	0.10	0.02	10000
weighted avg	0.21	0.12	0.03	10000

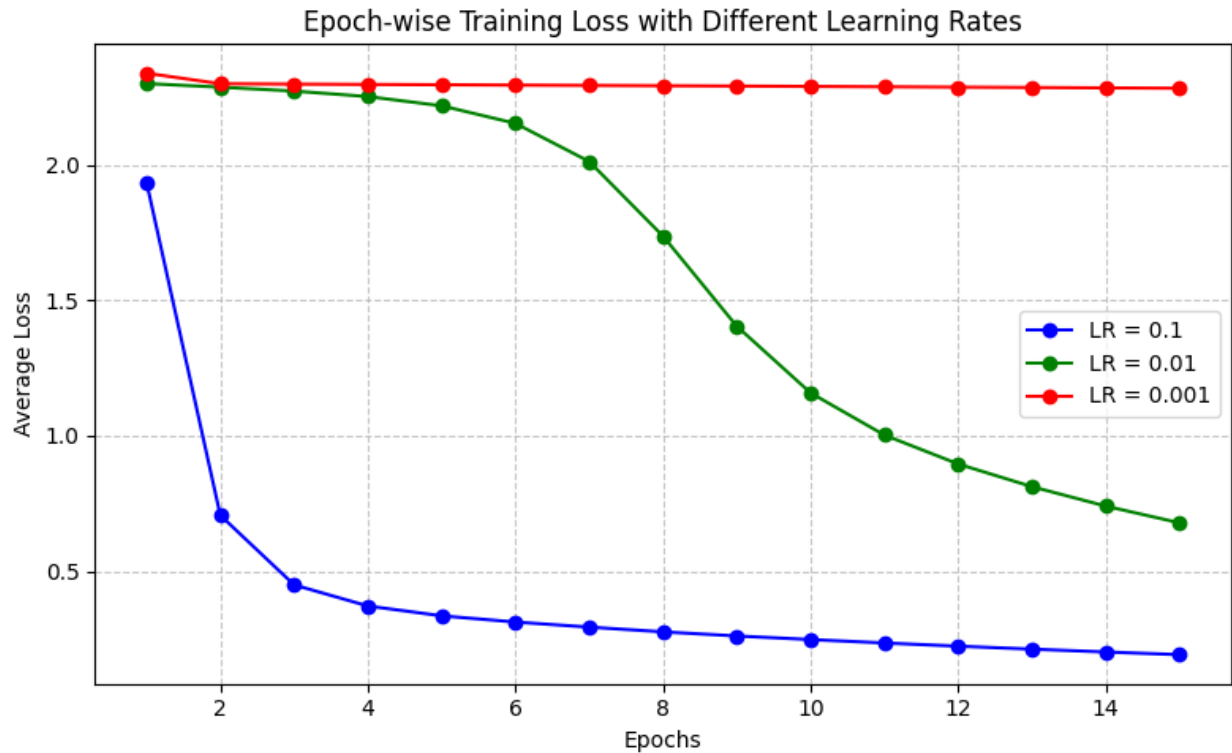
```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_
_classification.py:1565: UndefinedMetricWarning: Precision is ill-
defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classificatio
n.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classificatio
n.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))

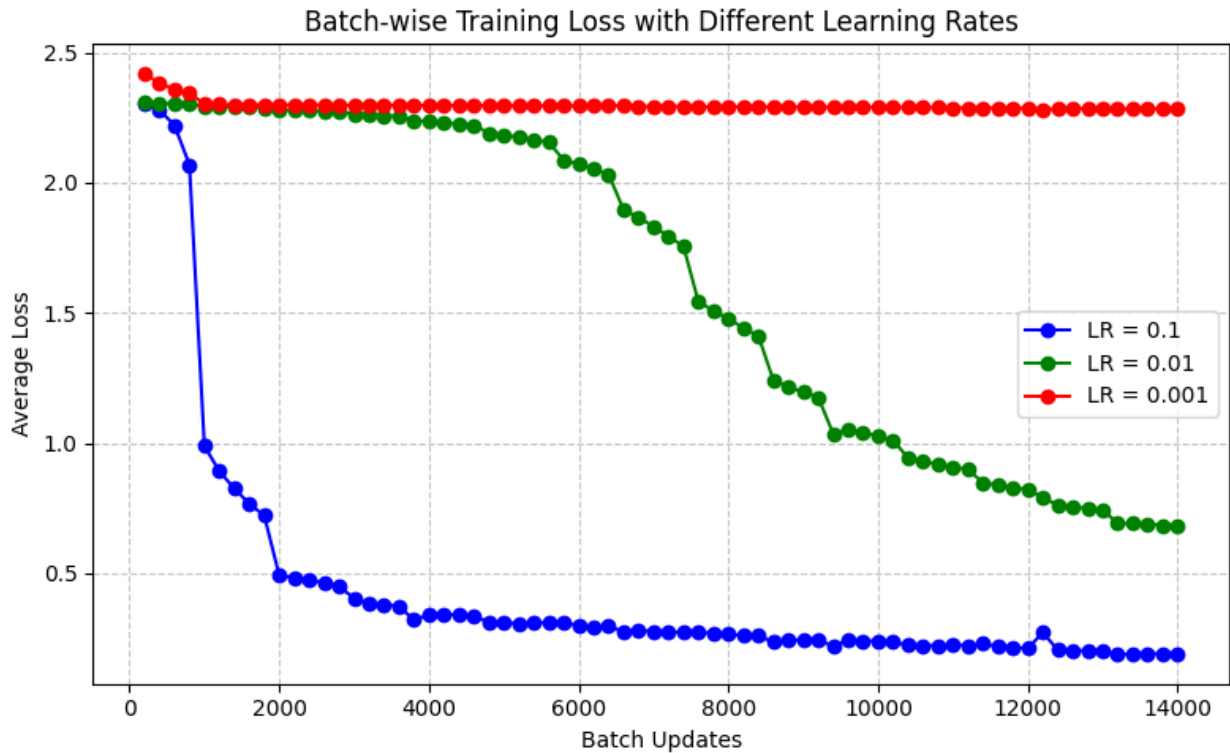
np.float64(0.1156)

plt.figure(figsize=(8,5))
plt.plot(range(1, len(loss_history_sigmoid_1)+1),
loss_history_sigmoid_1, marker='o', color='b', label="LR = 0.1")
plt.plot(range(1, len(loss_history)+1), loss_history, marker='o',
color='g', label="LR = 0.01")
plt.plot(range(1, len(loss_history_sigmoid_2)+1),
loss_history_sigmoid_2, marker='o', color='r', label="LR = 0.001")
plt.title("Epoch-wise Training Loss with Different Learning Rates")
plt.xlabel("Epochs")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()

```



```
plt.figure(figsize=(8,5))
plt.plot(iter_sigmoid_1, iter_loss_sigmoid_1, marker='o', color='b',
label="LR = 0.1")
plt.plot(iter_tanh, iter_loss, marker='o', color='g', label="LR =
0.01")
plt.plot(iter_sigmoid_2, iter_loss_sigmoid_2, marker='o', color='r',
label="LR = 0.001")
plt.title("Batch-wise Training Loss with Different Learning Rates")
plt.xlabel("Batch Updates")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()
```



Observations and Conclusions for different learning rates for the baseline model using Sigmoid Activation

- Test Accuracy with learning rate = 0.001 -> **11.56%**
- Test Accuracy with learning rate = 0.01 -> **82.30%** (Baseline model)
- Test Accuracy with learning rate = 0.1 -> **94.25%**

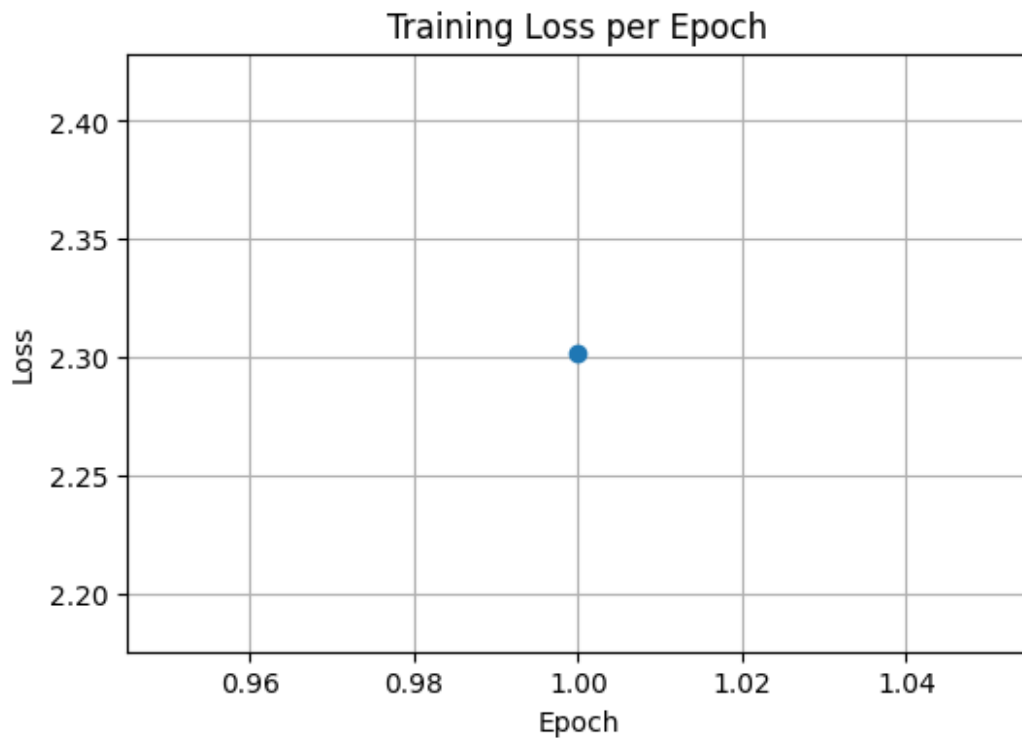
Higher Learning rate leads to better test accuracy and lesser loss in training. Increasing the learning rate helped in improving the baseline model quite drastically

Analyzing the MLP baseline model for different num of epochs using Sigmoid activation function

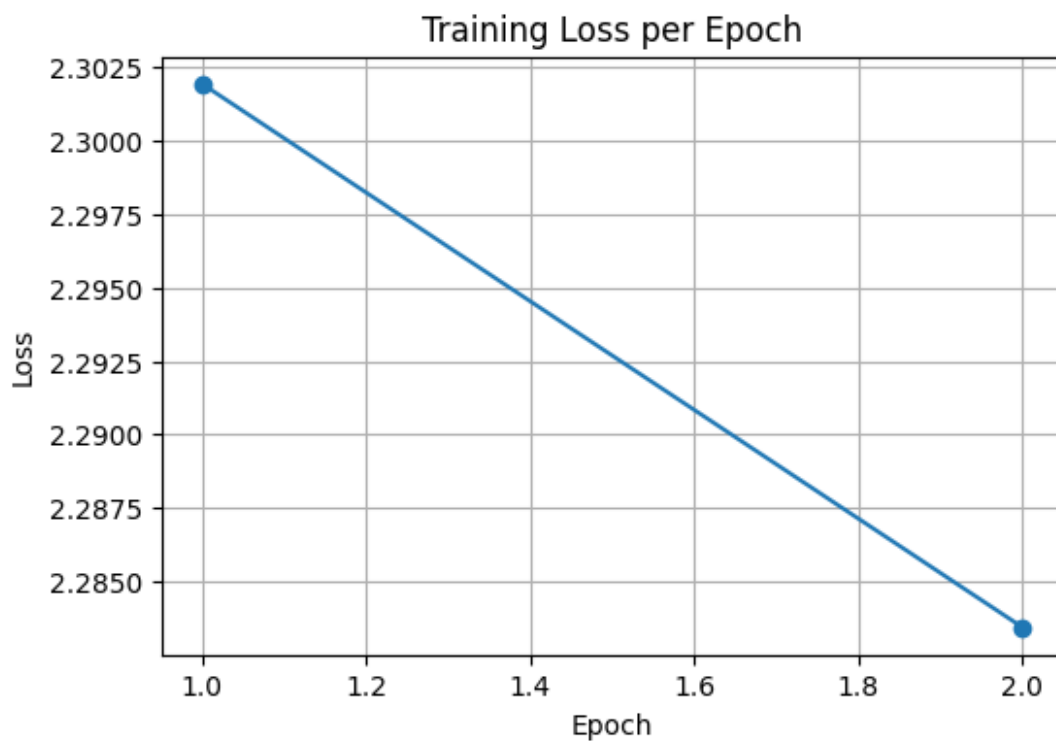
Number of Epochs = 10

```
model_3 = MLP()  
loss_history_sigmoid_3, iter_sigmoid_3, iter_loss_sigmoid_3 =  
train(model_3, train_loader, 10, 0.01)
```

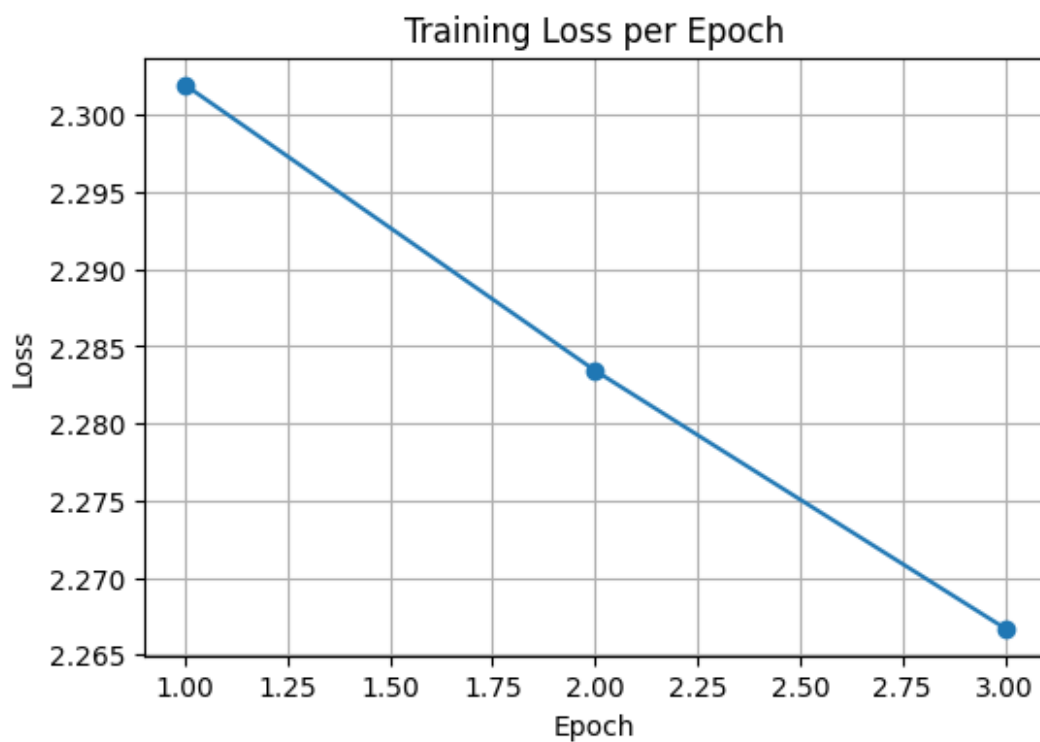
Epoch 1/10, Loss: 2.3019



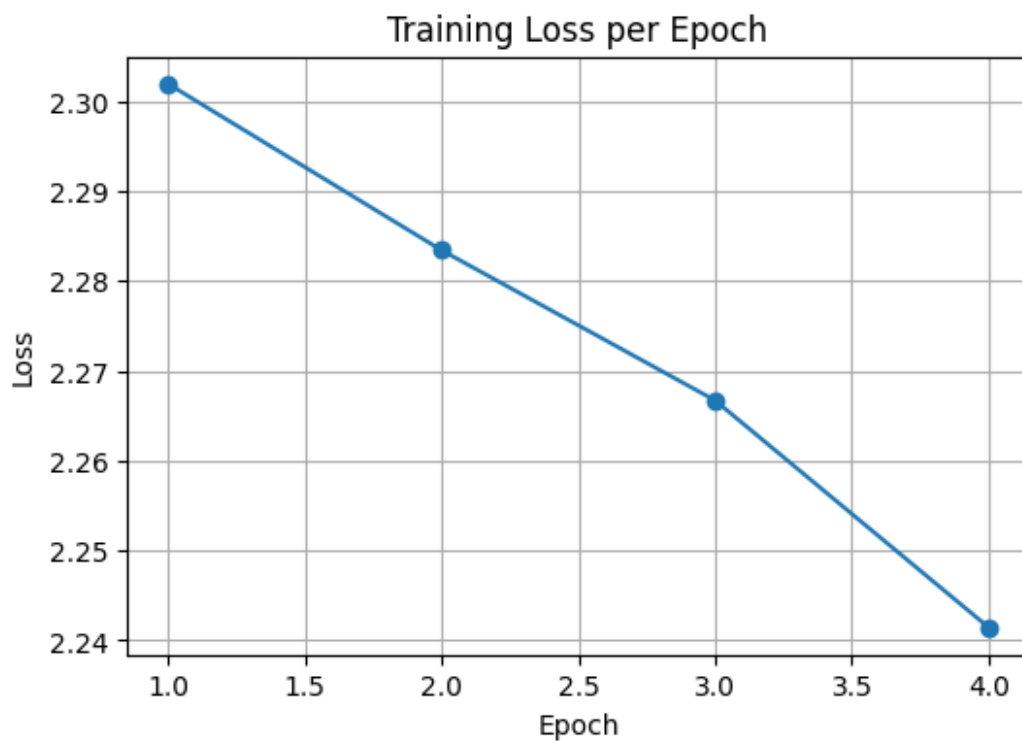
Epoch 2/10, Loss: 2.2834



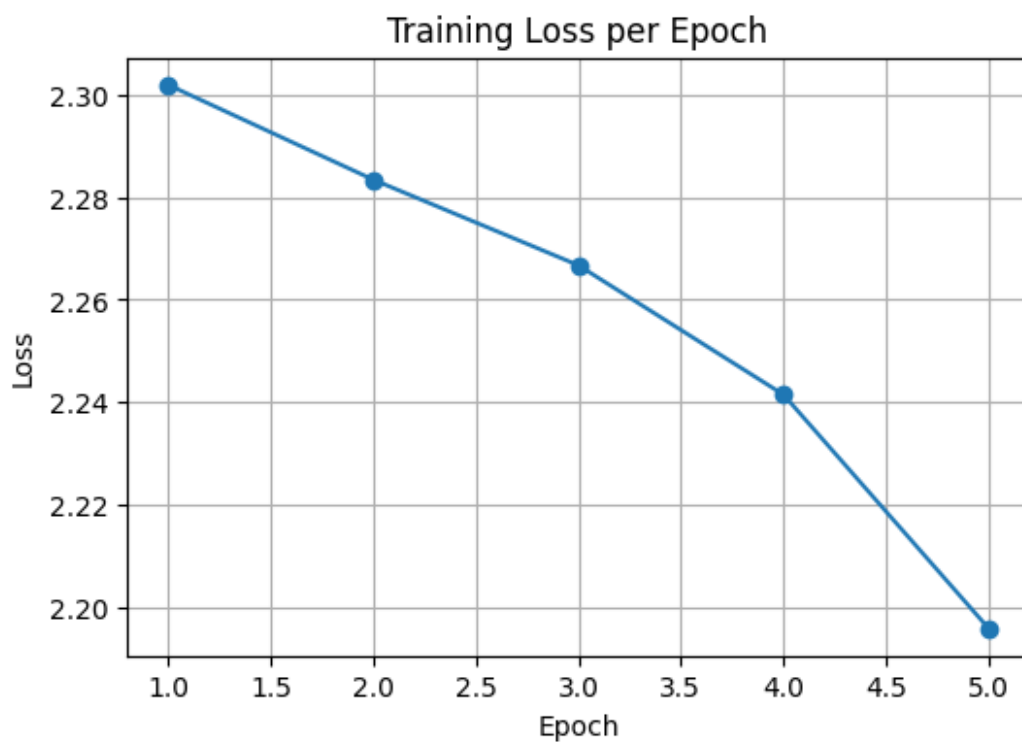
Epoch 3/10, Loss: 2.2667



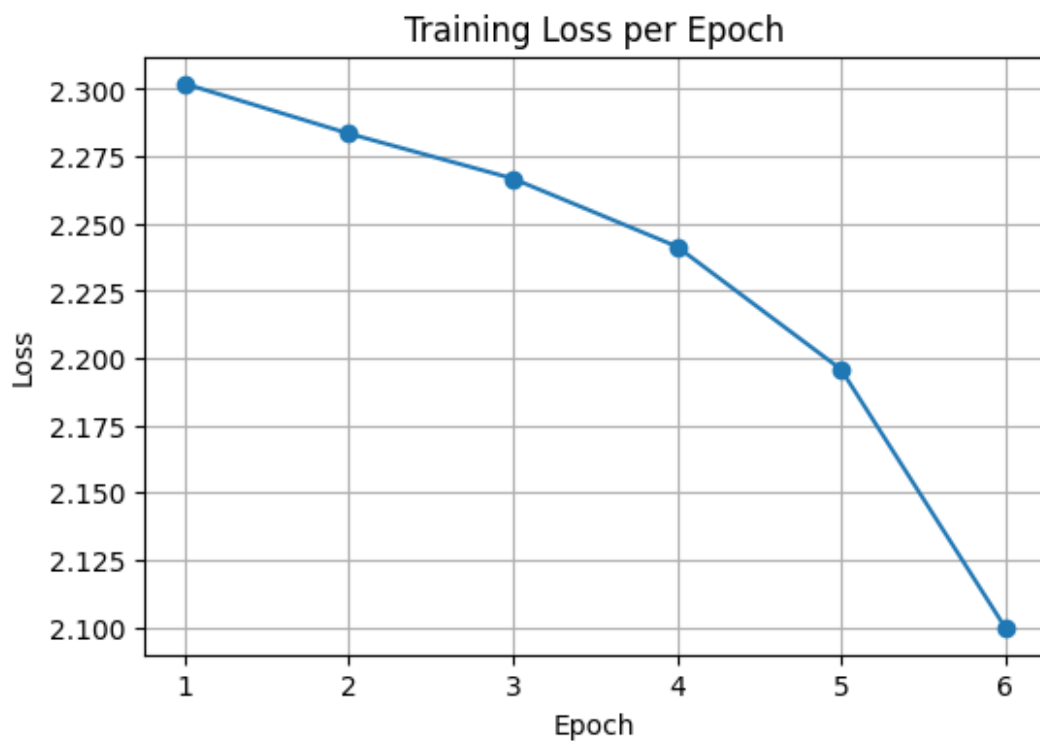
Epoch 4/10, Loss: 2.2415



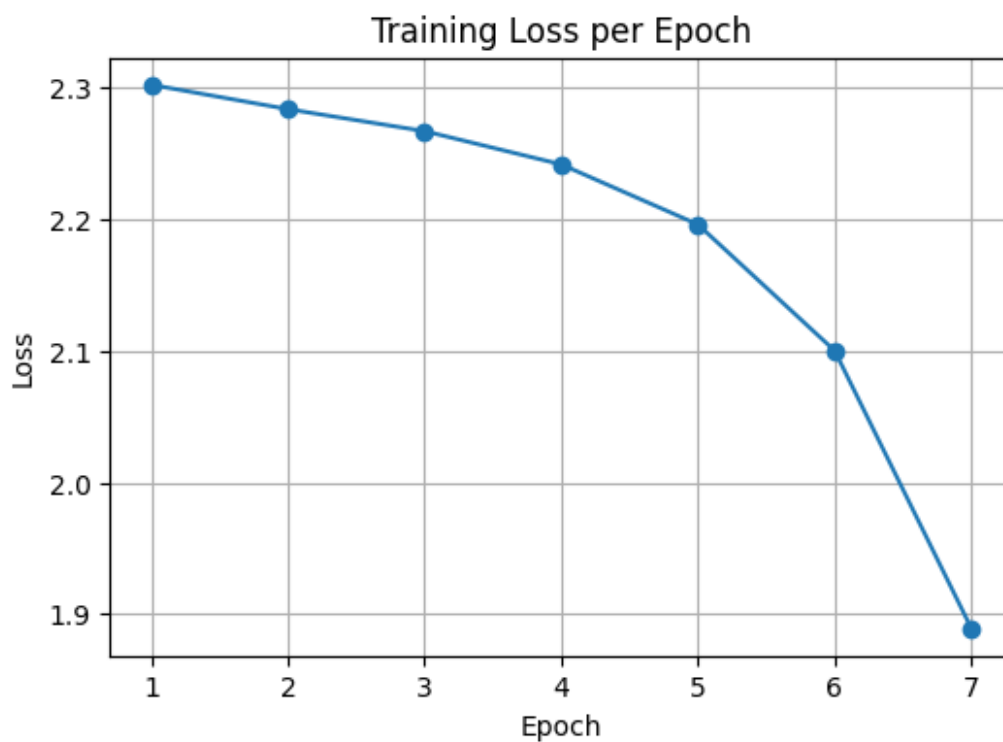
Epoch 5/10, Loss: 2.1958



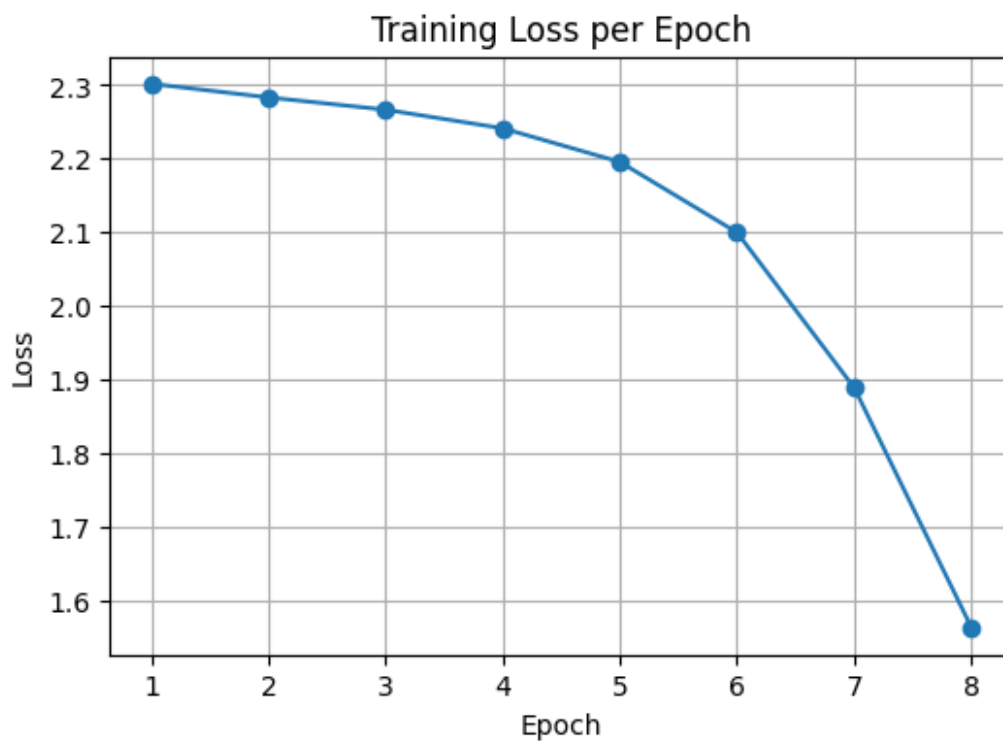
Epoch 6/10, Loss: 2.1000



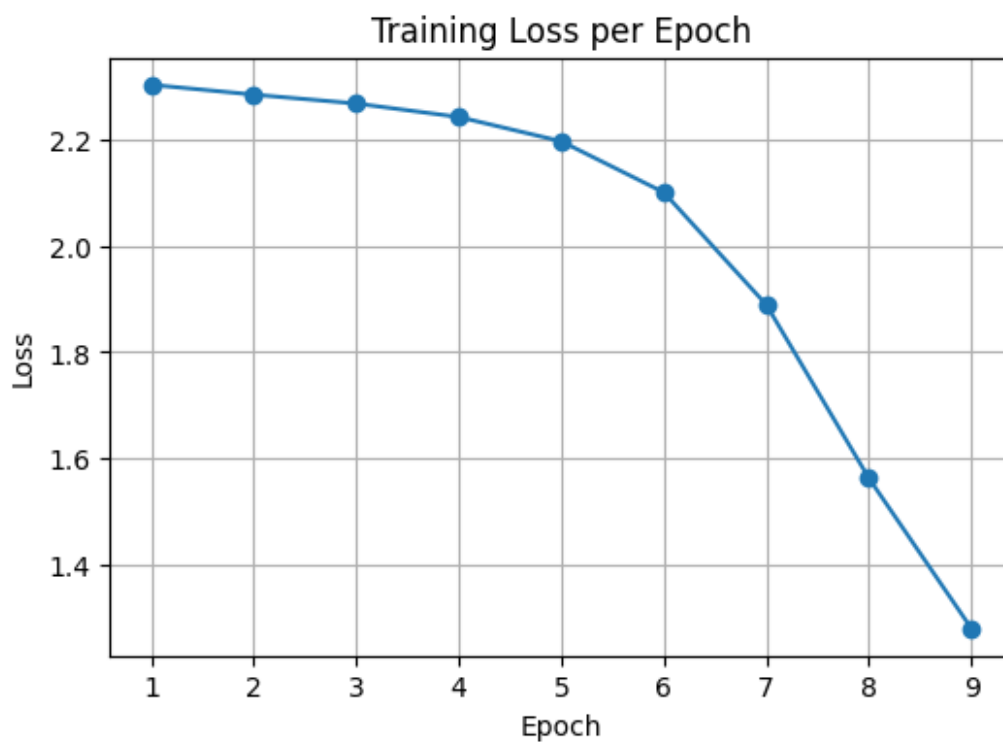
Epoch 7/10, Loss: 1.8891



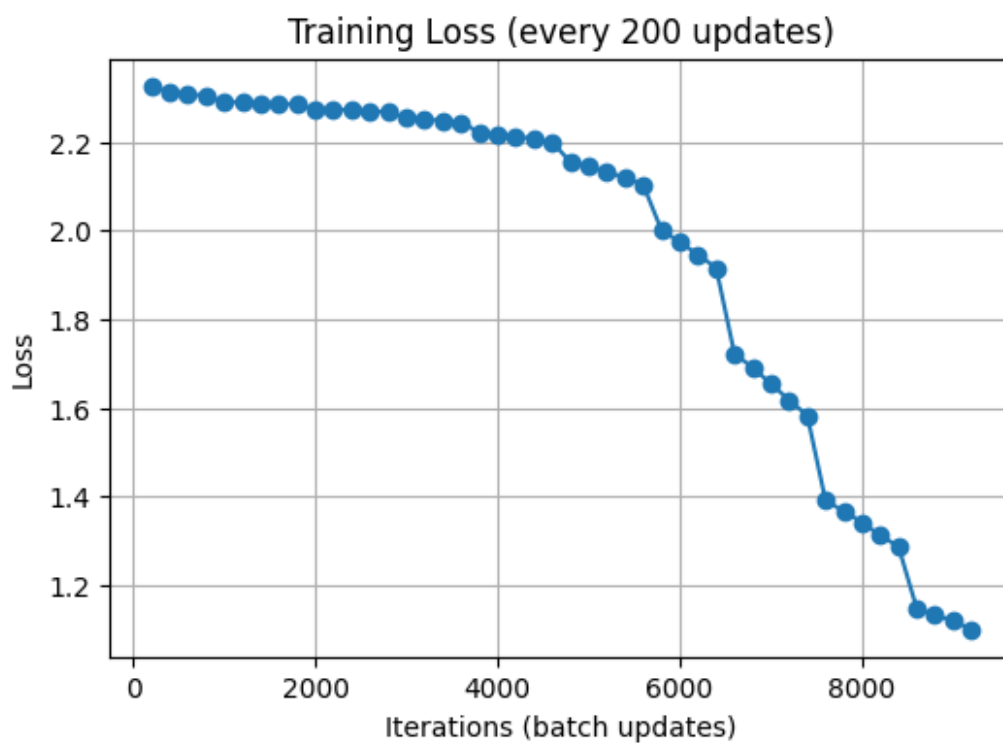
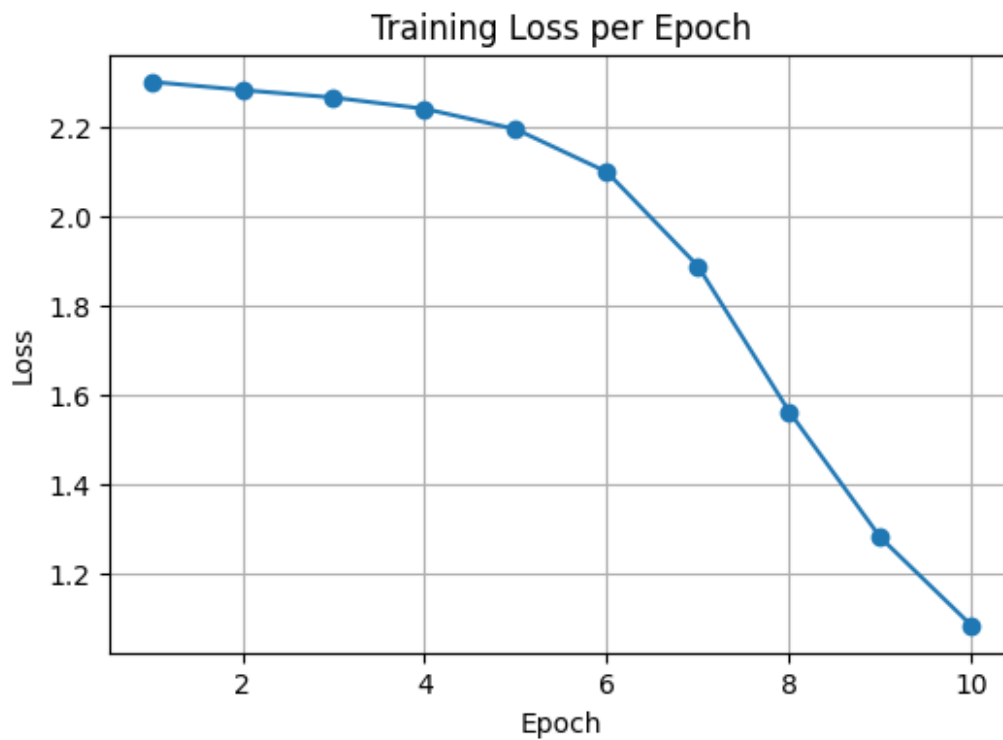
Epoch 8/10, Loss: 1.5635



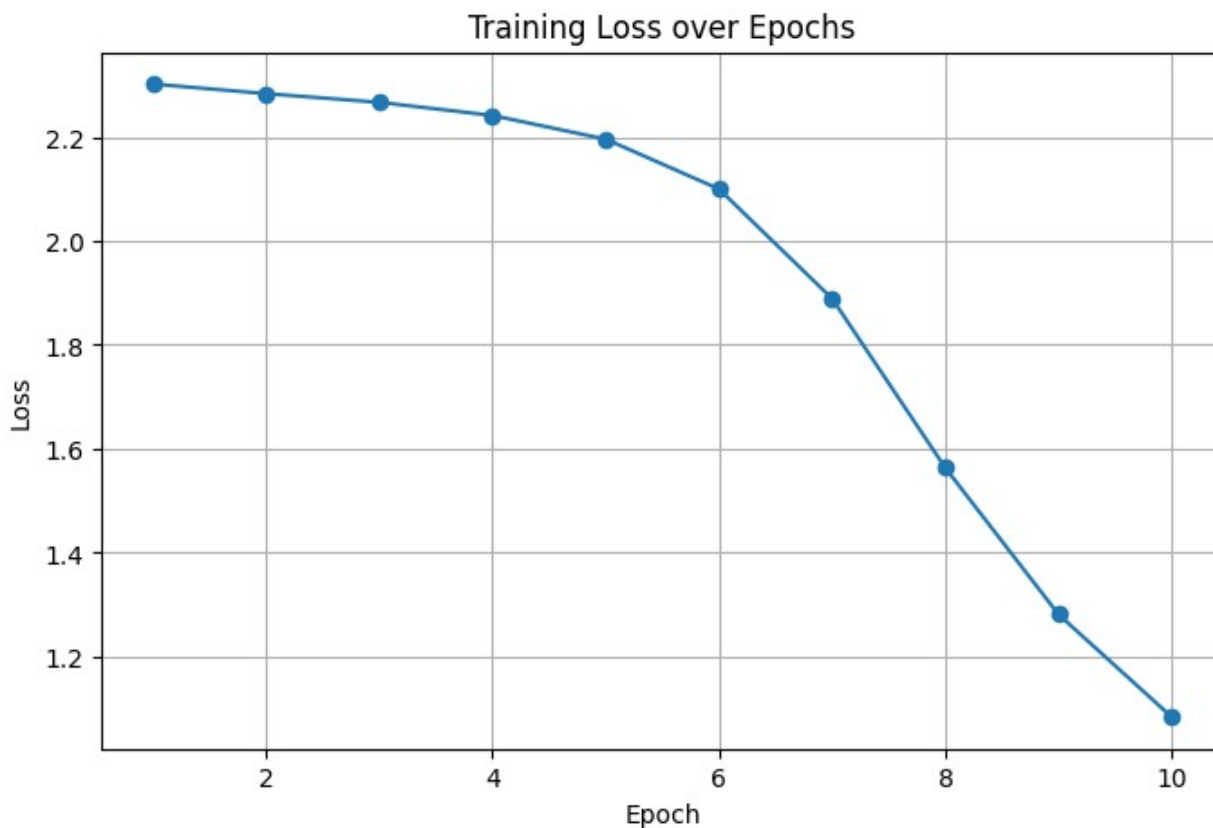
Epoch 9/10, Loss: 1.2815



Epoch 10/10, Loss: 1.0842



```
plot_training_curves(loss_history_sigmoid_3)
```



```
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("Test Set Evaluation for 10 epochs")
```

```
evaluate(model_3, X_test_np, y_test_np)
```

Test Set Evaluation for 10 epochs

Accuracy: 69.7600%

Confusion Matrix:

```
[ [ 942    0    1    4    1   21    8    1    2    0]
 [    0 1121    4    5    0    0    1    0    3    1]
 [   34   42   705   67   30   10  114    7   20    3]
 [    8   32   53  829    3   41    1   17   18    8]
 [    1   28    2    0  607    0   50   20    1  273]
 [   57   16   35  269   18  353   34   56   36   18]
 [   39    7   61    0   18    8  815    0   10    0]
 [    4   71    8    2   14    9    1  852    8   59]
 [   40  127  114  292   15   42   21   24  262   37]
 [   16   26    1   11  185    9    8  263    0  490]]
```

Classification Report:

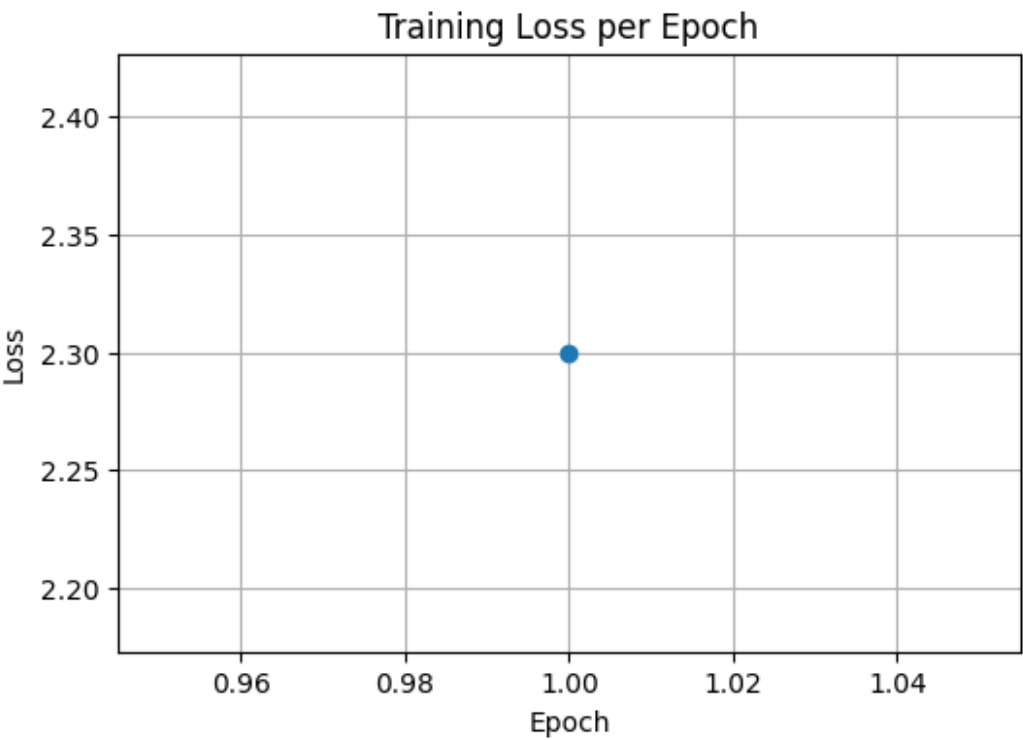
	precision	recall	f1-score	support
0	0.83	0.96	0.89	980

1	0.76	0.99	0.86	1135
2	0.72	0.68	0.70	1032
3	0.56	0.82	0.67	1010
4	0.68	0.62	0.65	982
5	0.72	0.40	0.51	892
6	0.77	0.85	0.81	958
7	0.69	0.83	0.75	1028
8	0.73	0.27	0.39	974
9	0.55	0.49	0.52	1009
accuracy			0.70	10000
macro avg	0.70	0.69	0.67	10000
weighted avg	0.70	0.70	0.68	10000

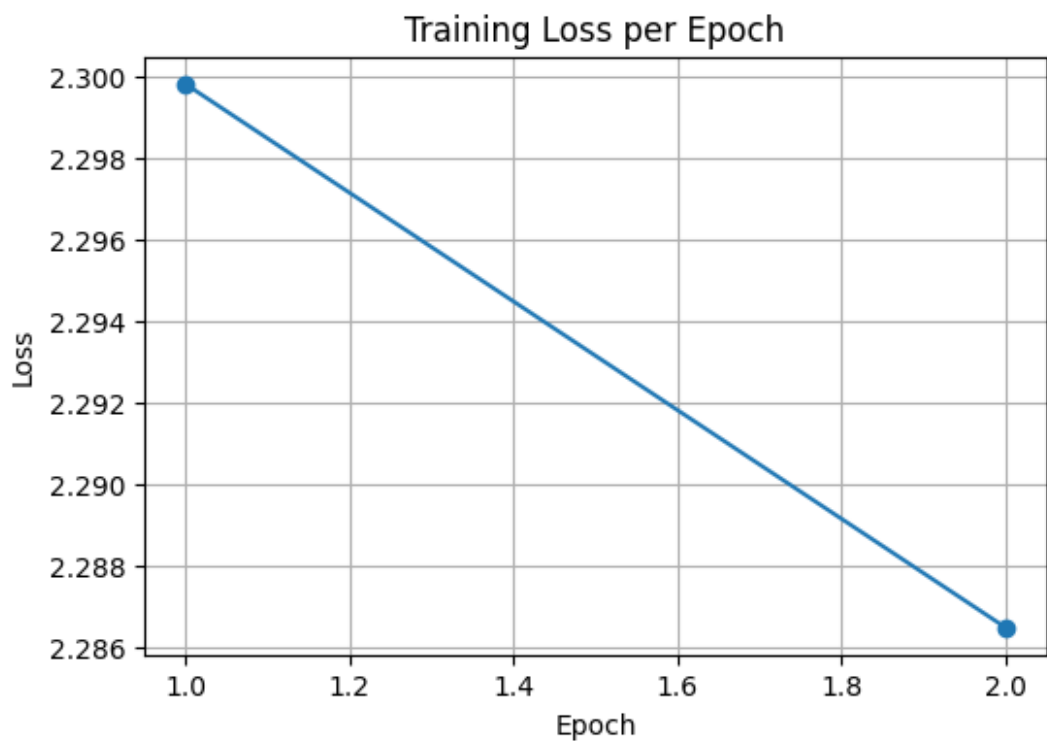
np.float64(0.6976)

Number of Epochs = 20

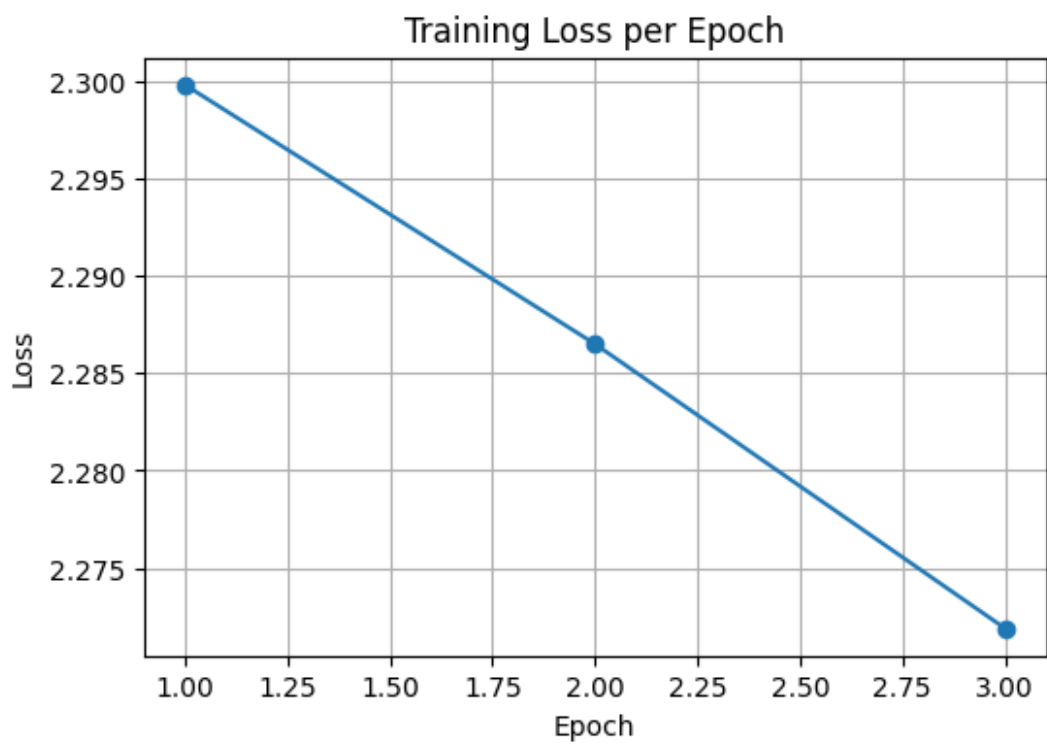
```
model_4 = MLP()
loss_history_sigmoid_4, iter_sigmoid_4, iter_loss_sigmoid_4 =
train(model_4, train_loader, 20, 0.01)
Epoch 1/20, Loss: 2.2998
```



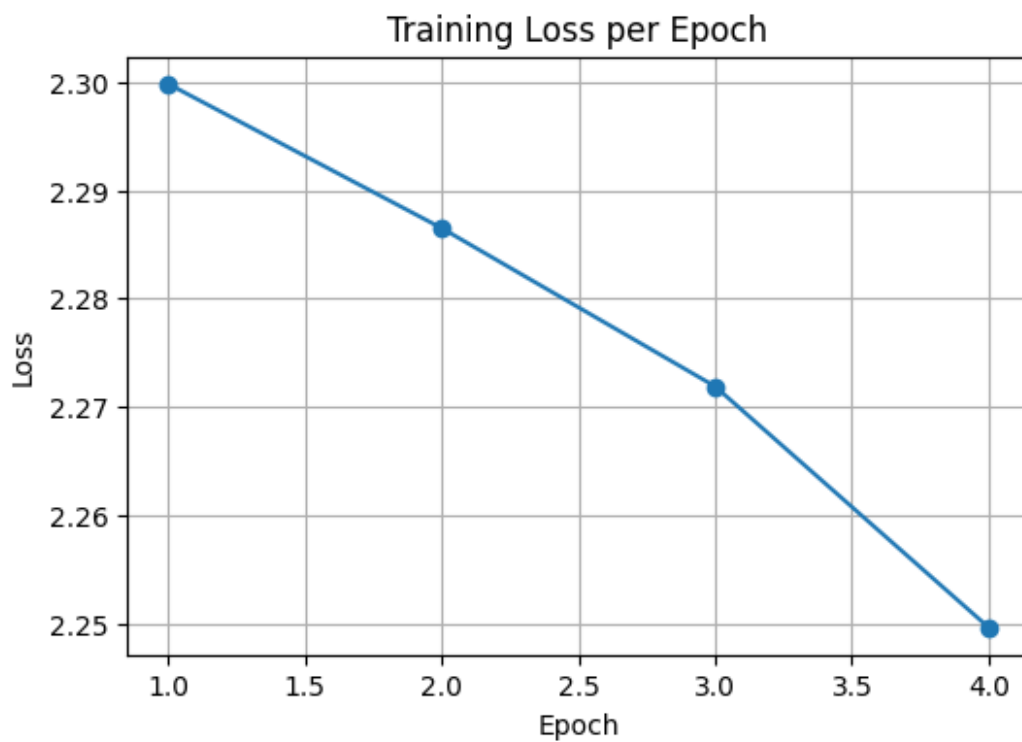
Epoch 2/20, Loss: 2.2865



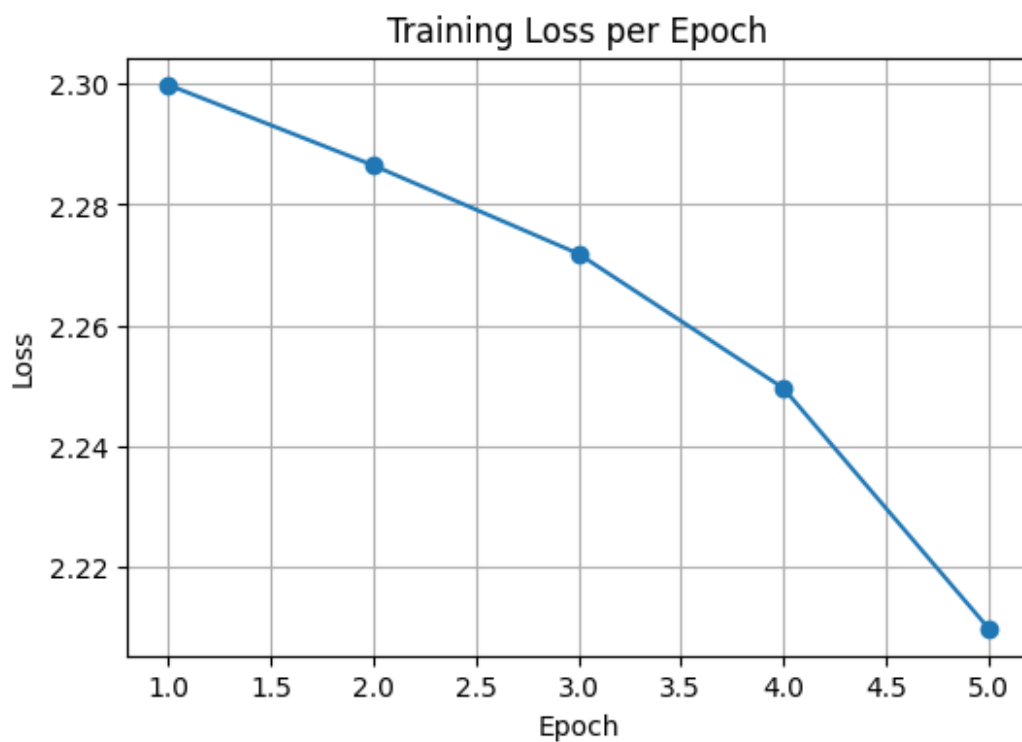
Epoch 3/20, Loss: 2.2719



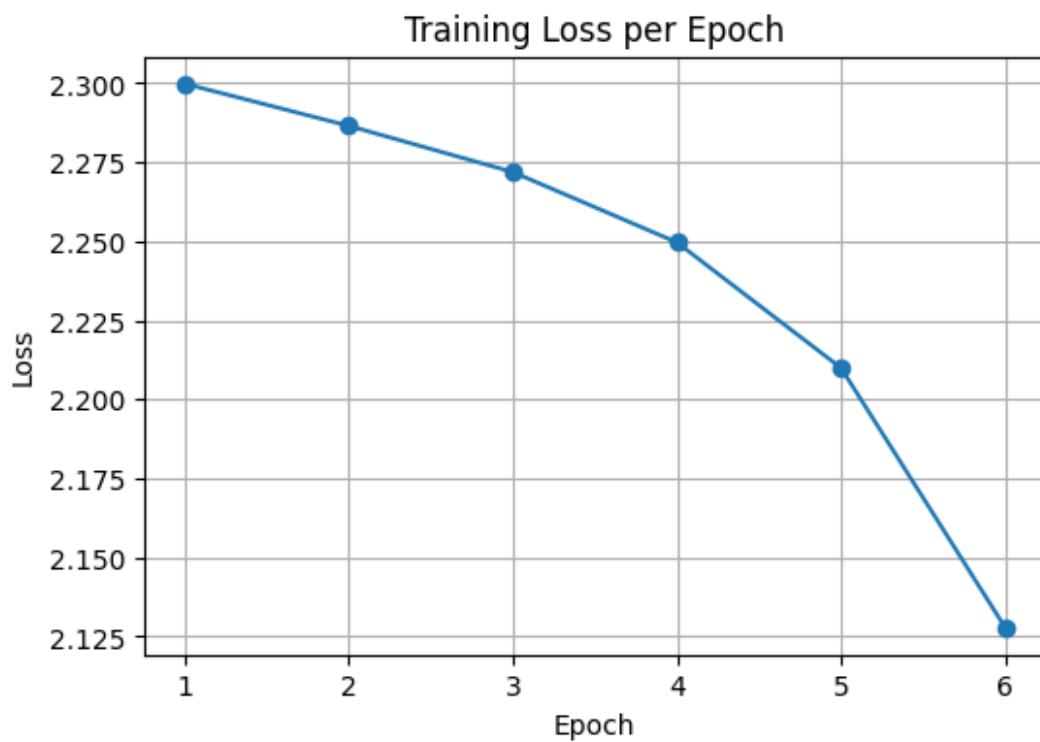
Epoch 4/20, Loss: 2.2496



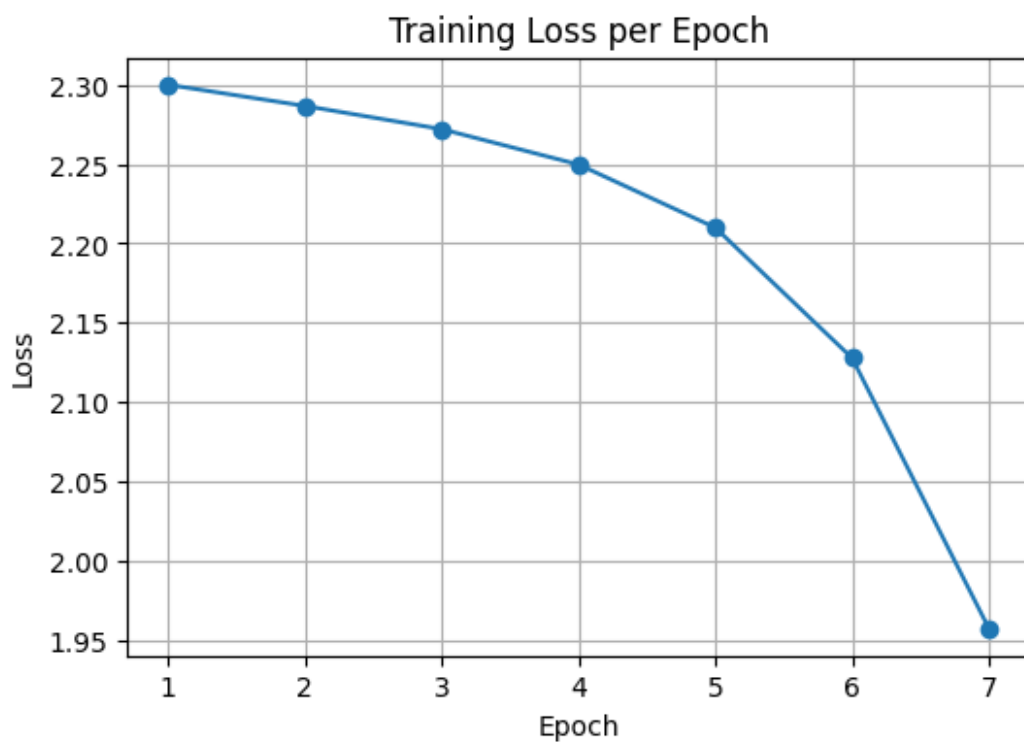
Epoch 5/20, Loss: 2.2099



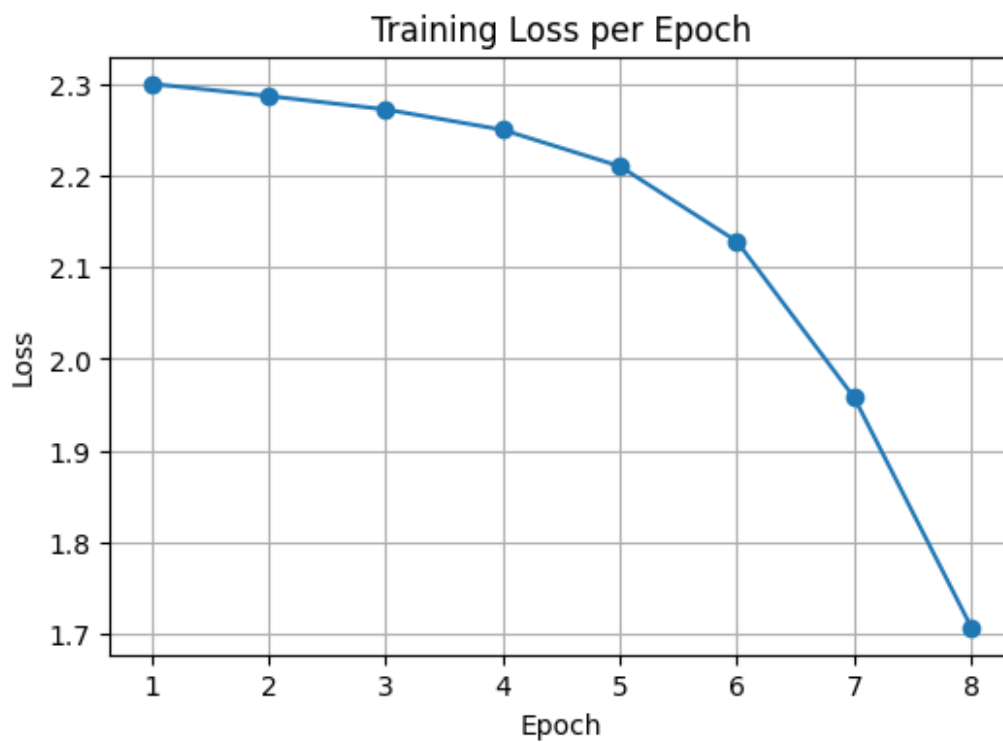
Epoch 6/20, Loss: 2.1279



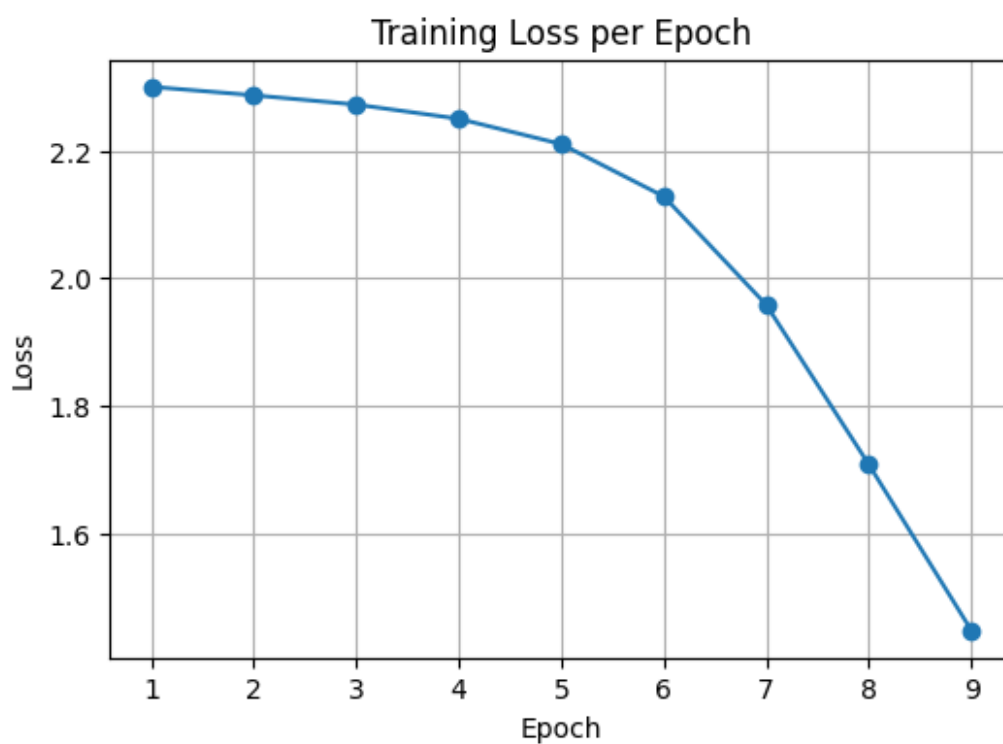
Epoch 7/20, Loss: 1.9578



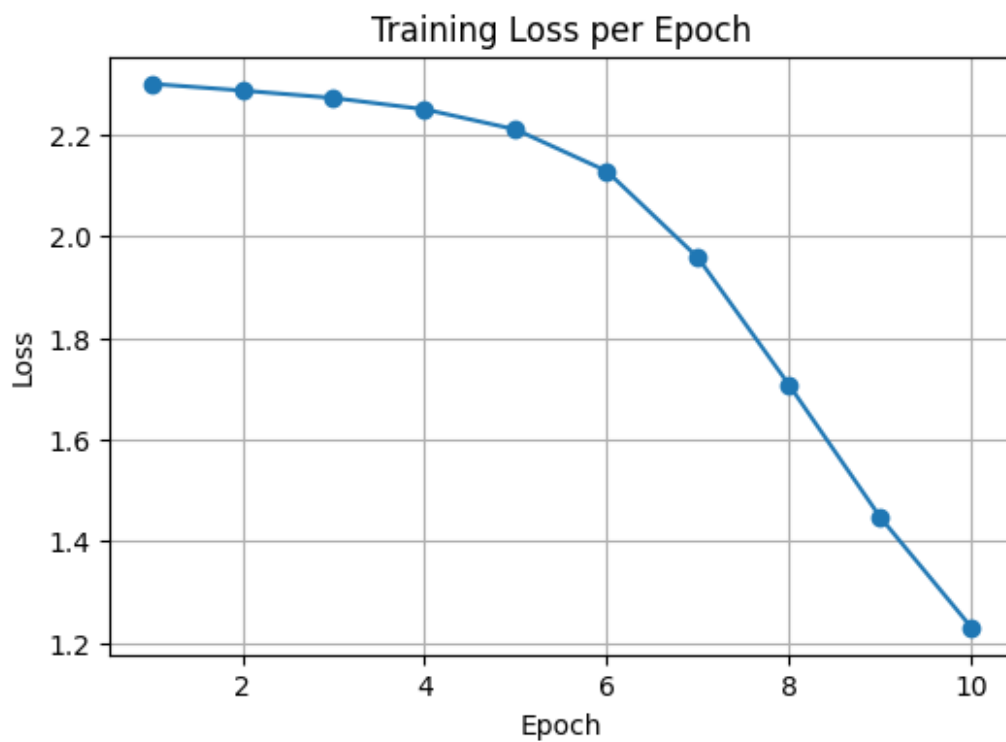
Epoch 8/20, Loss: 1.7073



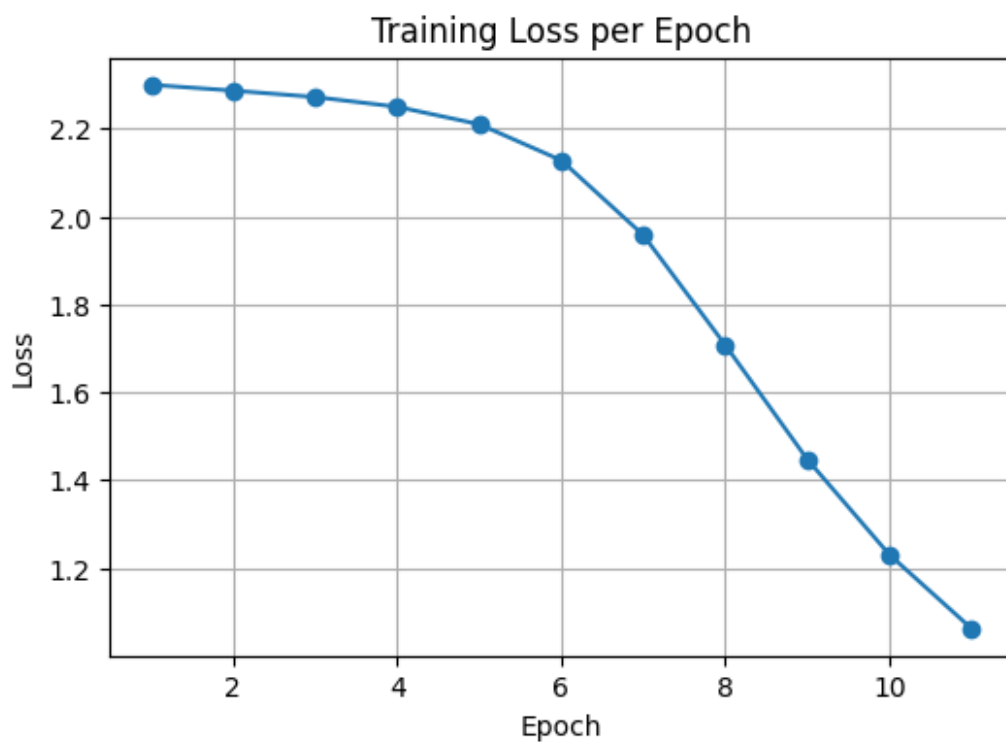
Epoch 9/20, Loss: 1.4472



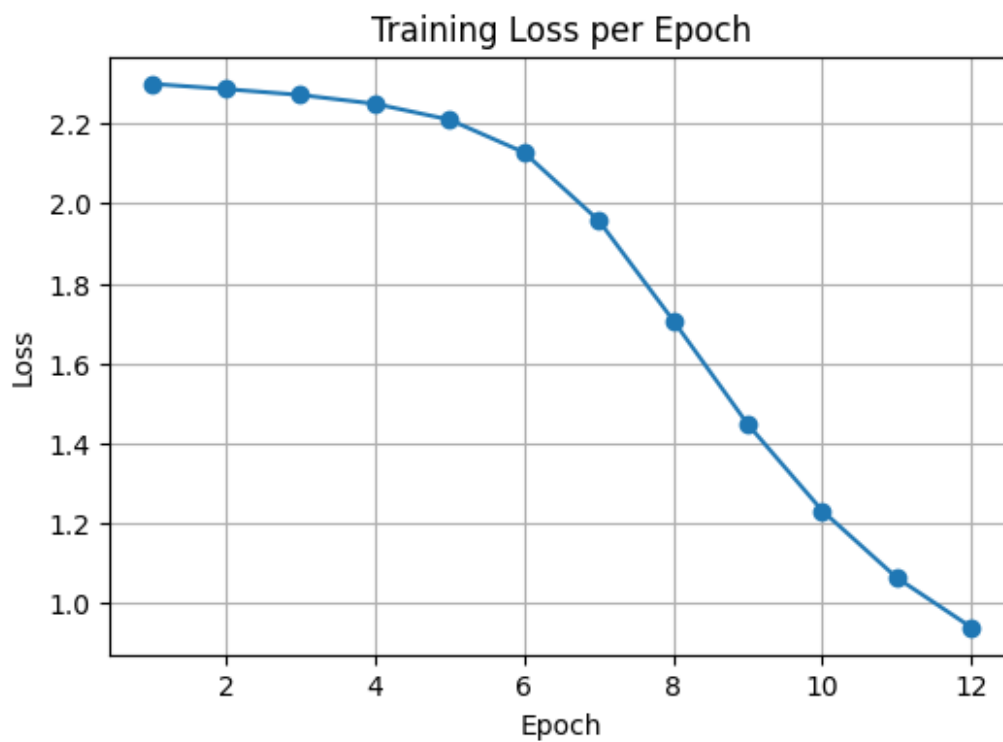
Epoch 10/20, Loss: 1.2305



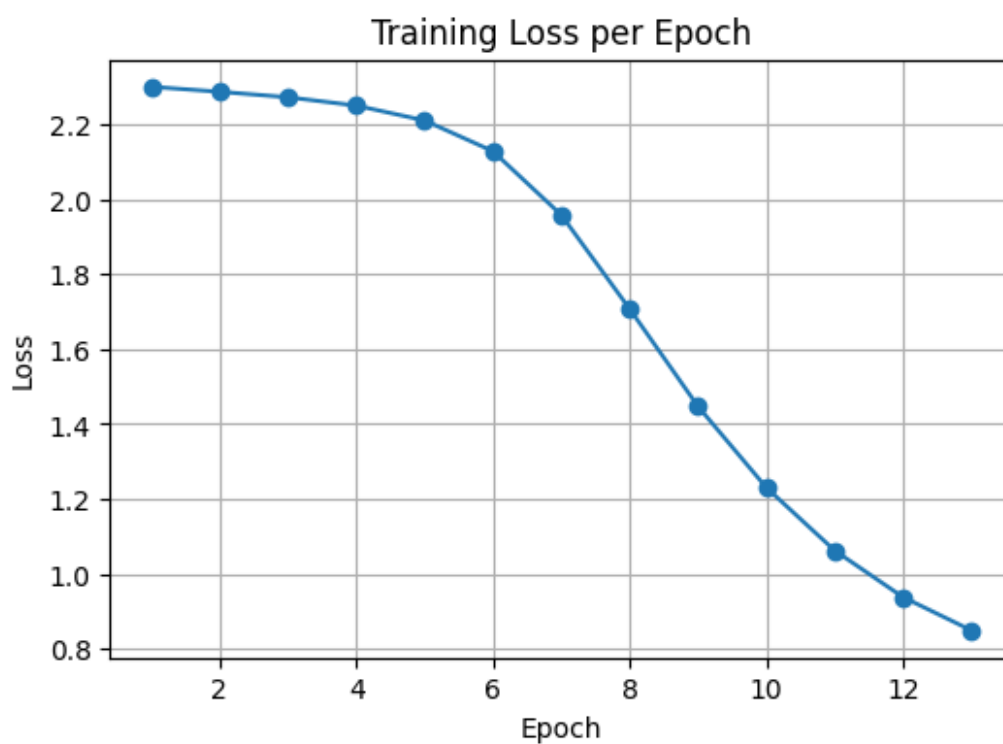
Epoch 11/20, Loss: 1.0625



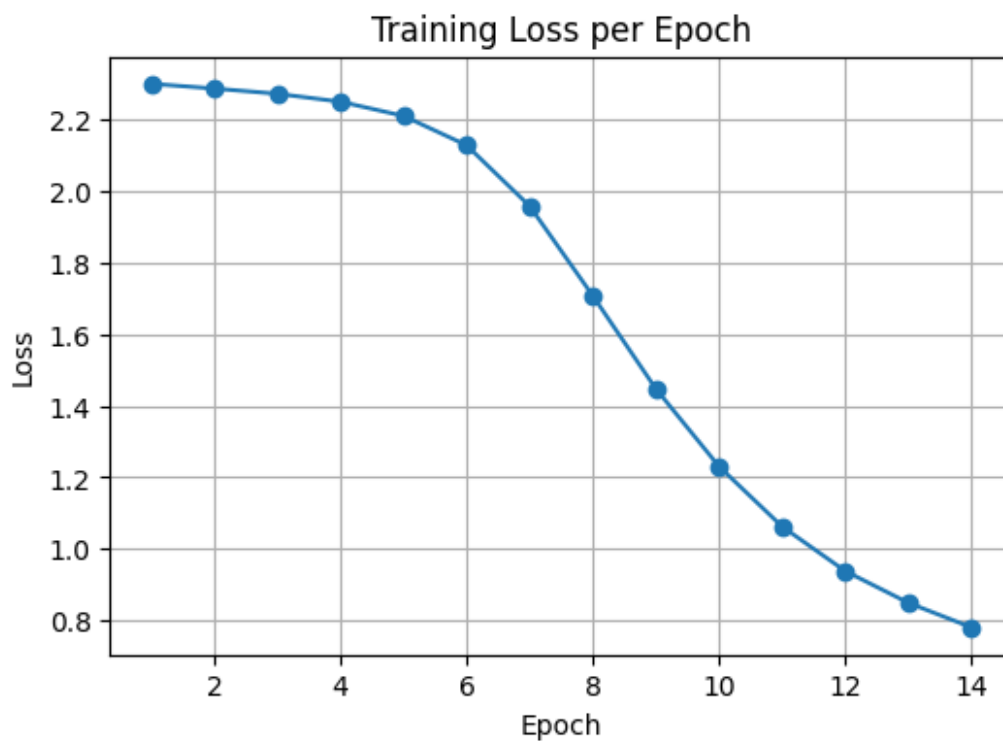
Epoch 12/20, Loss: 0.9382



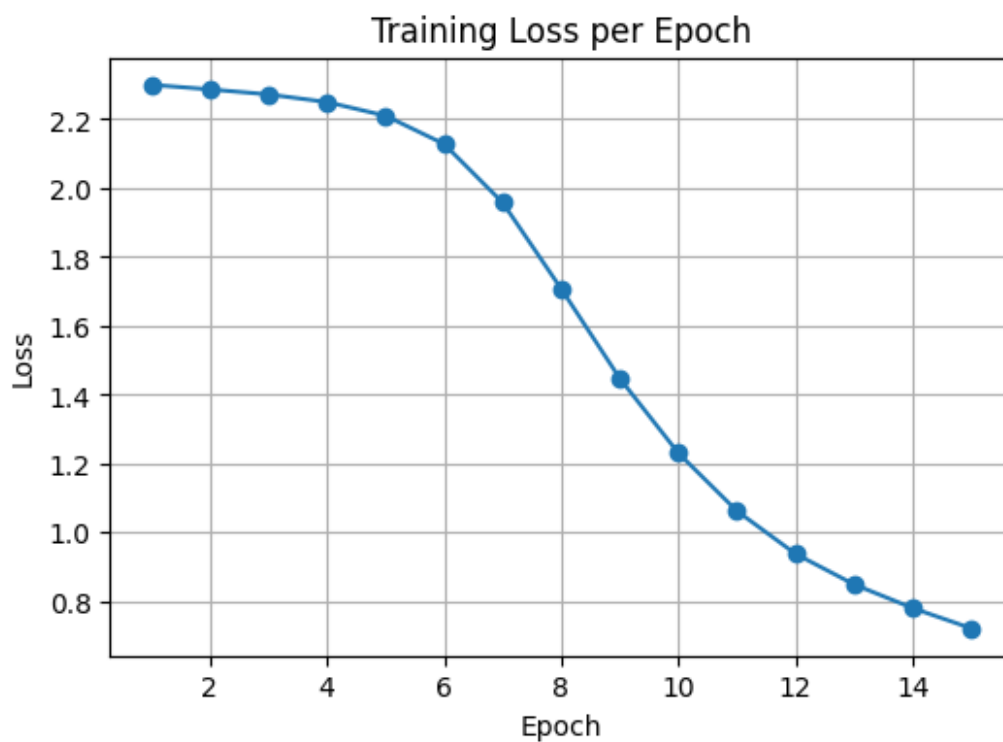
Epoch 13/20, Loss: 0.8494



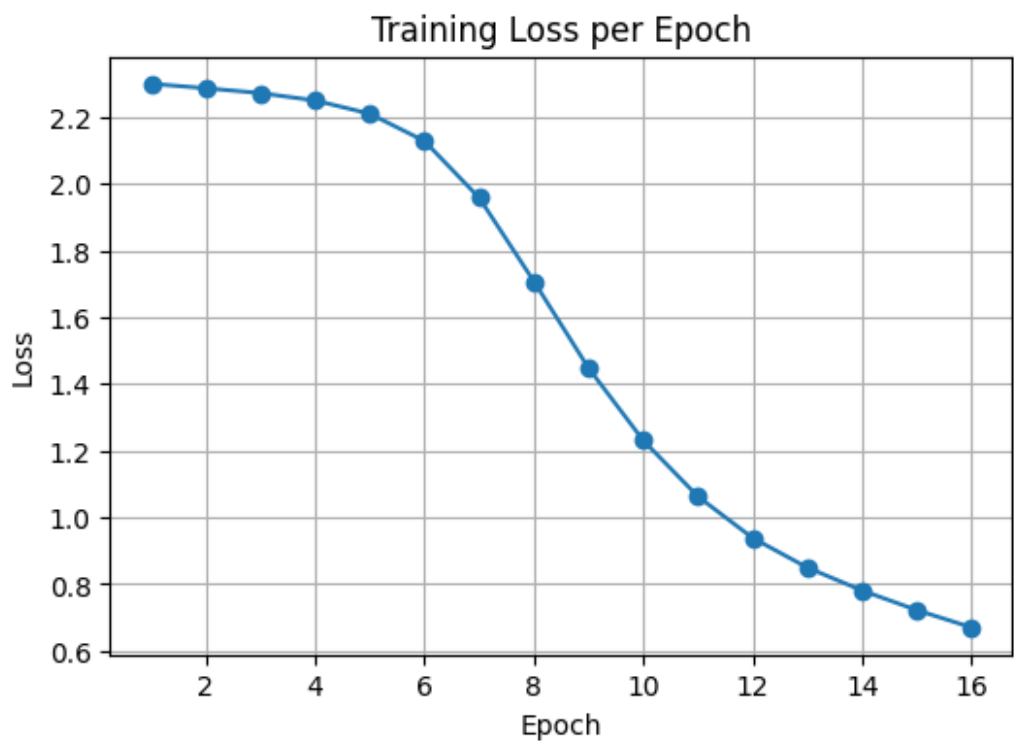
Epoch 14/20, Loss: 0.7807



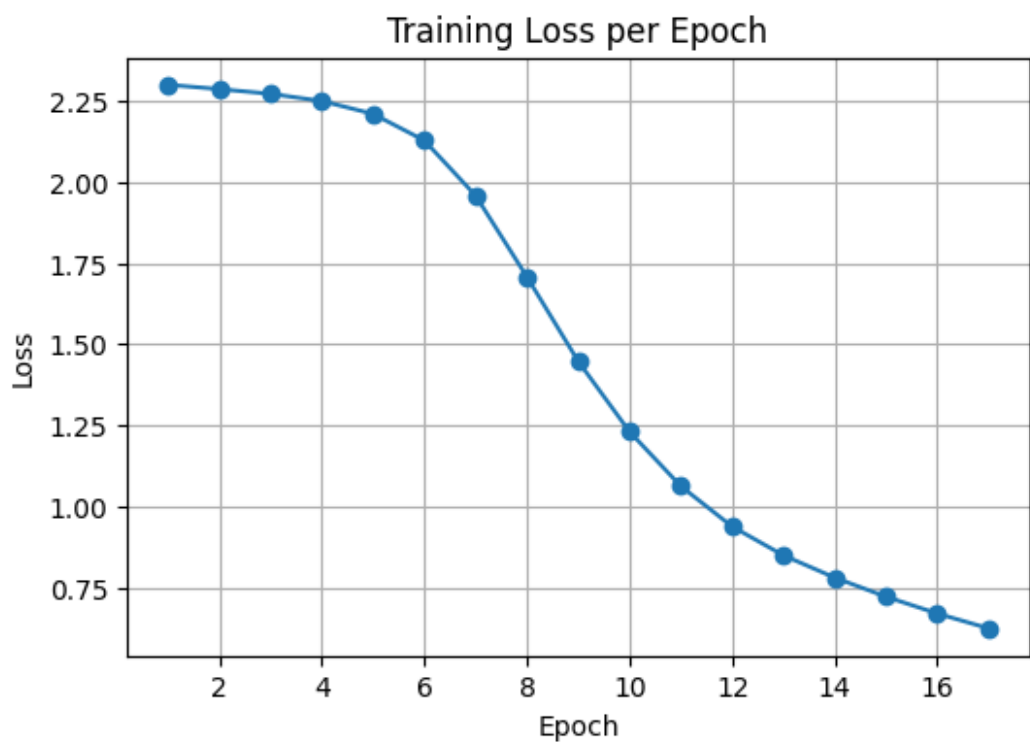
Epoch 15/20, Loss: 0.7219



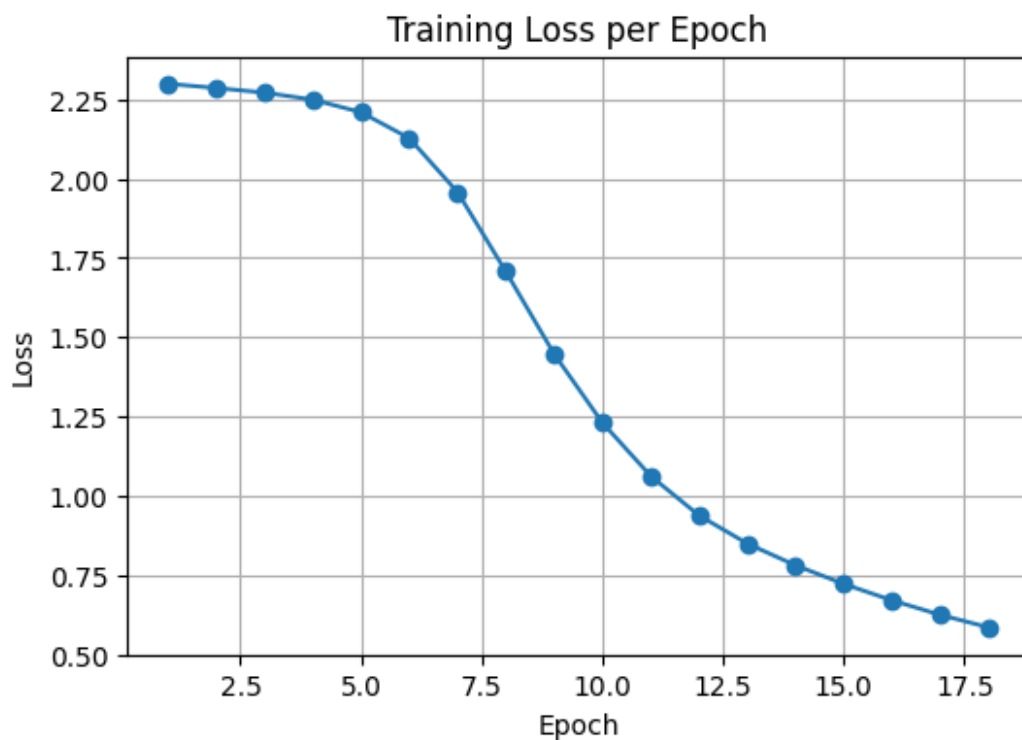
Epoch 16/20, Loss: 0.6701



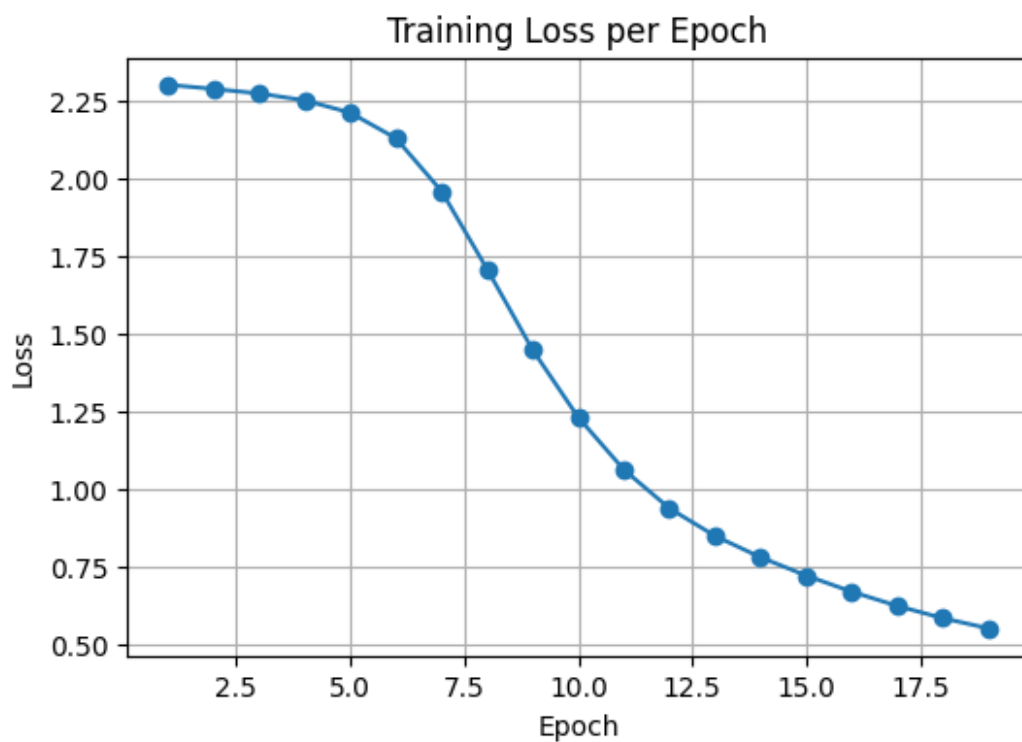
Epoch 17/20, Loss: 0.6244



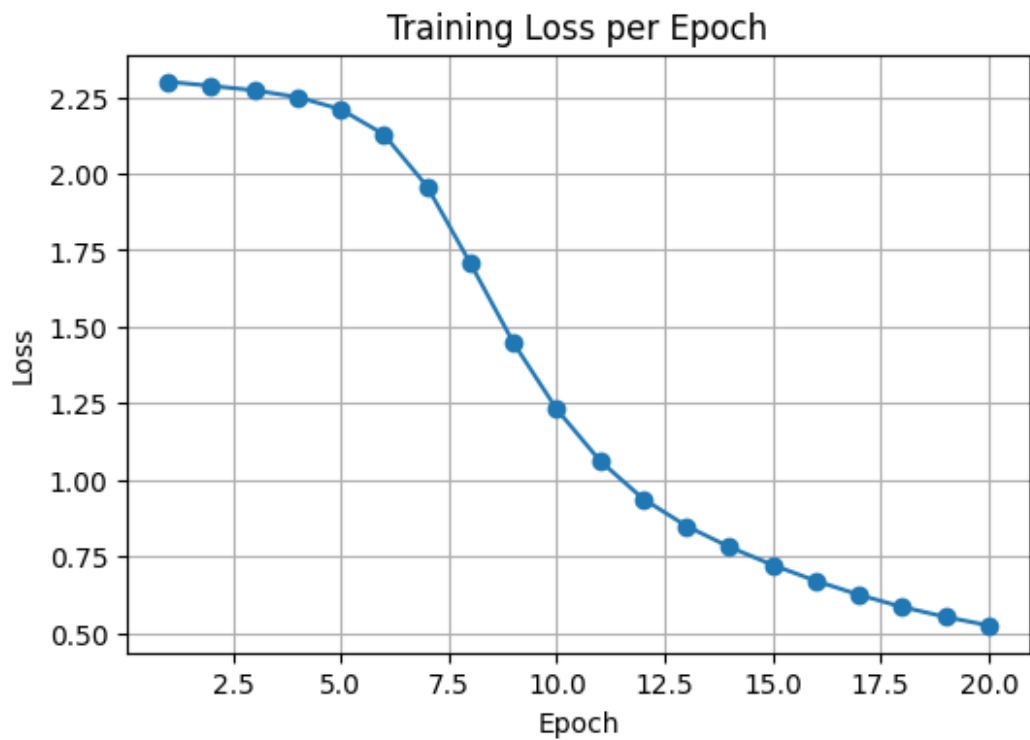
Epoch 18/20, Loss: 0.5850



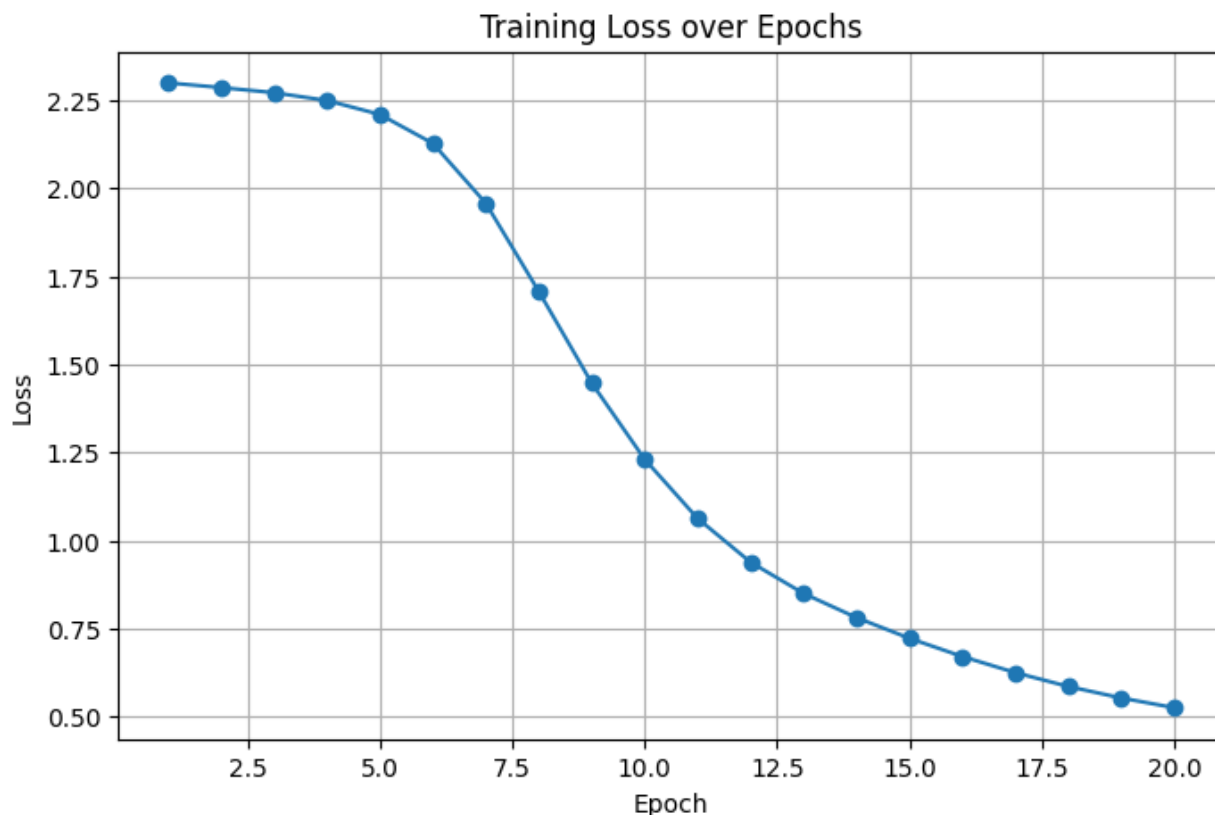
Epoch 19/20, Loss: 0.5523



Epoch 20/20, Loss: 0.5247



```
plot_training_curves(loss_history_sigmoid_4)
```



```
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("Test Set Evaluation for 20 epochs")
evaluate(model_4, X_test_np, y_test_np)
```

Test Set Evaluation for 20 epochs

Accuracy: 85.3900%

Confusion Matrix:

```
[ [ 951    0    2    3    0   13   10    1    0    0]
  [    0 1112    1    6    0    2    3    1    9    1]
  [   14   11  866   58   24    2   23   10   19    5]
  [    3    7   46  842    0   48    1   18   39    6]
  [    1    8    5    0  804    0   23    0    8  133]
  [   22    3   12   90   11  656   29    8   46   15]
  [   24    4   12    1   15   21  878    0    3    0]
  [    6   32   27    7    1    1    0  888    5   61]
  [    9   21   21   45   27   51   17    9  743   31]
  [   13    8    7    5   99   15    0   51   12  799]]
```

Classification Report:

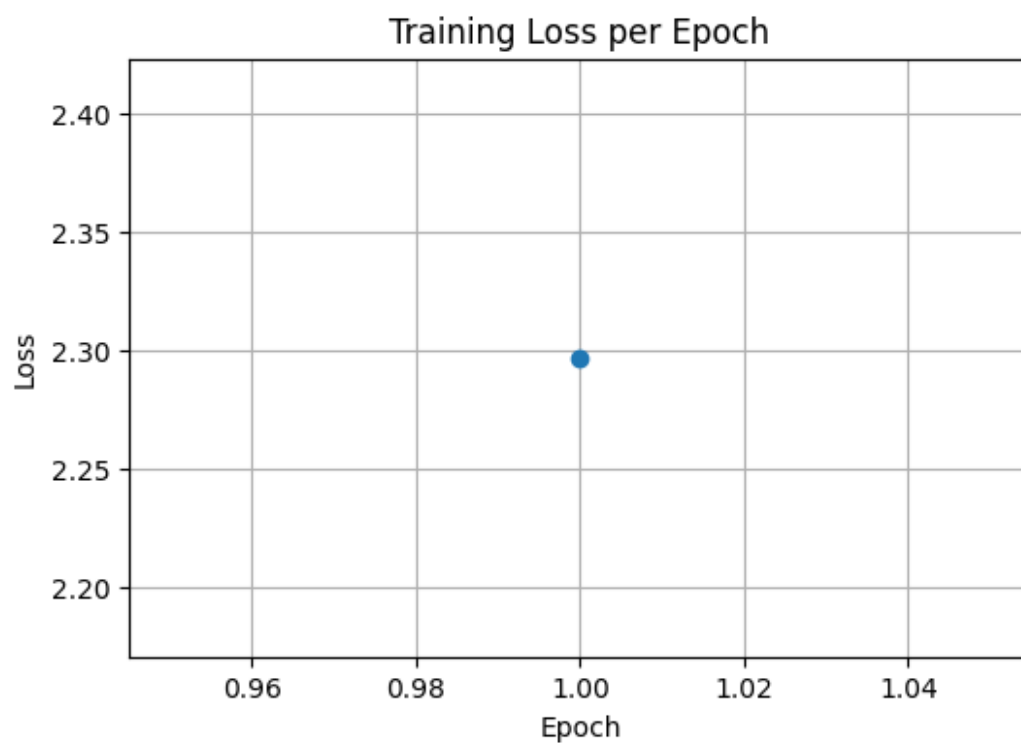
	precision	recall	f1-score	support
0	0.91	0.97	0.94	980
1	0.92	0.98	0.95	1135

2	0.87	0.84	0.85	1032
3	0.80	0.83	0.81	1010
4	0.82	0.82	0.82	982
5	0.81	0.74	0.77	892
6	0.89	0.92	0.90	958
7	0.90	0.86	0.88	1028
8	0.84	0.76	0.80	974
9	0.76	0.79	0.78	1009
accuracy			0.85	10000
macro avg	0.85	0.85	0.85	10000
weighted avg	0.85	0.85	0.85	10000
np.float64(0.8539)				

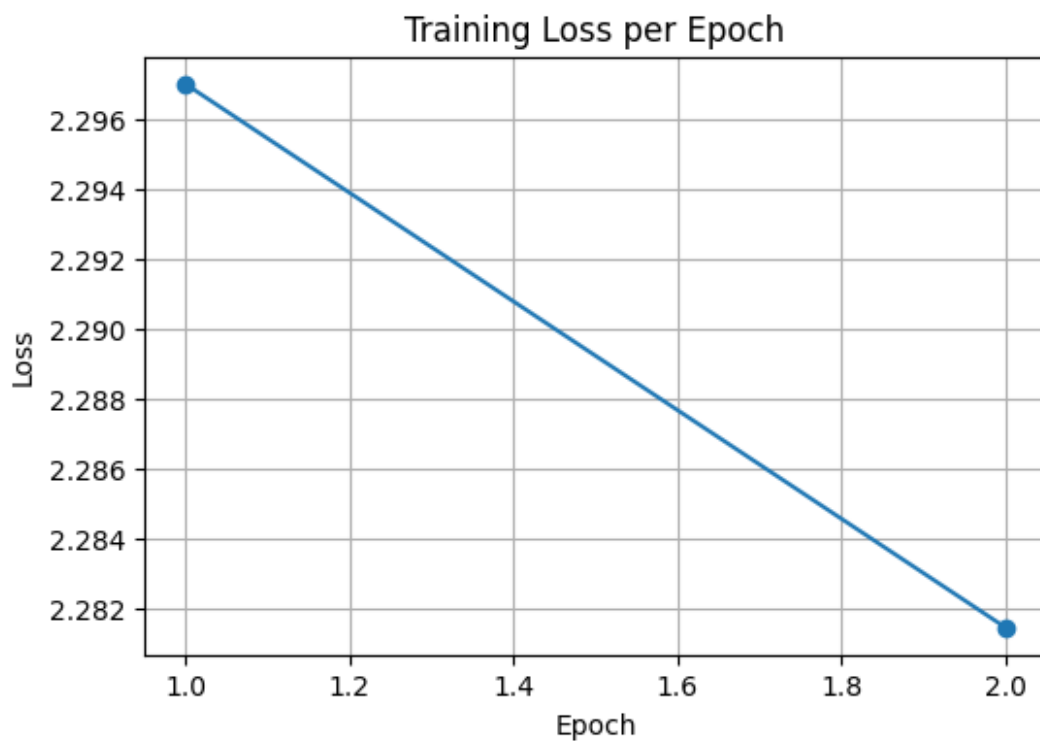
Number of Epochs = 25

```
model_5 = MLP()
loss_history_sigmoid_5, iter_sigmoid_5, iter_loss_sigmoid_5 =
train(model_5, train_loader, 25, 0.01)
```

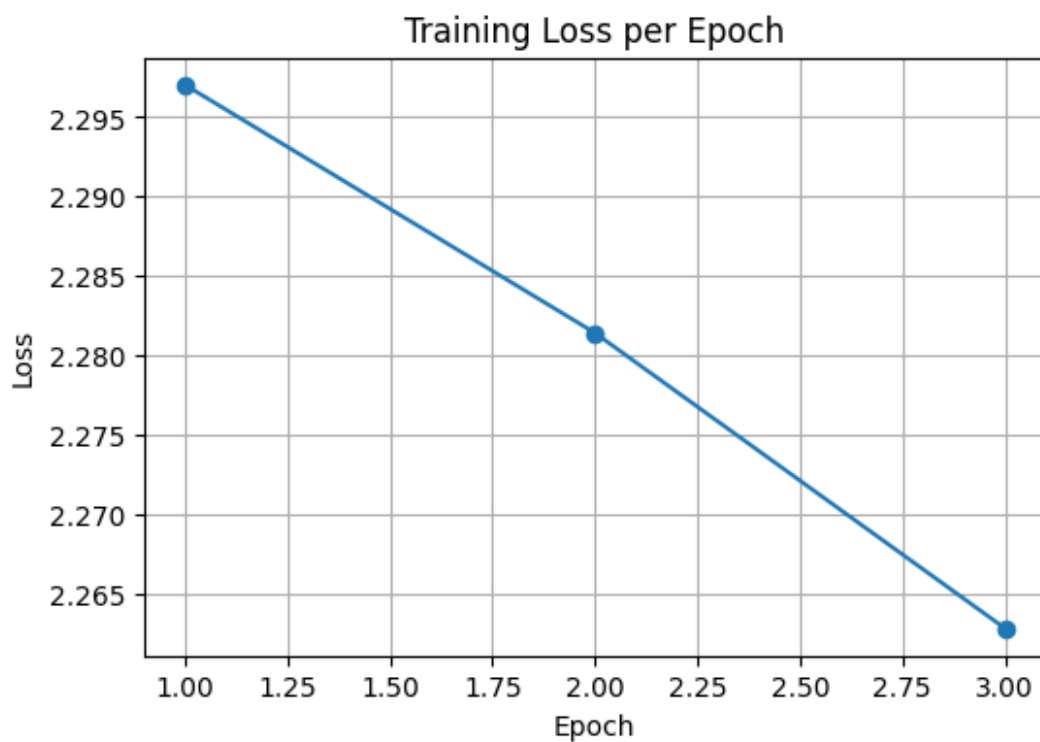
Epoch 1/25, Loss: 2.2970



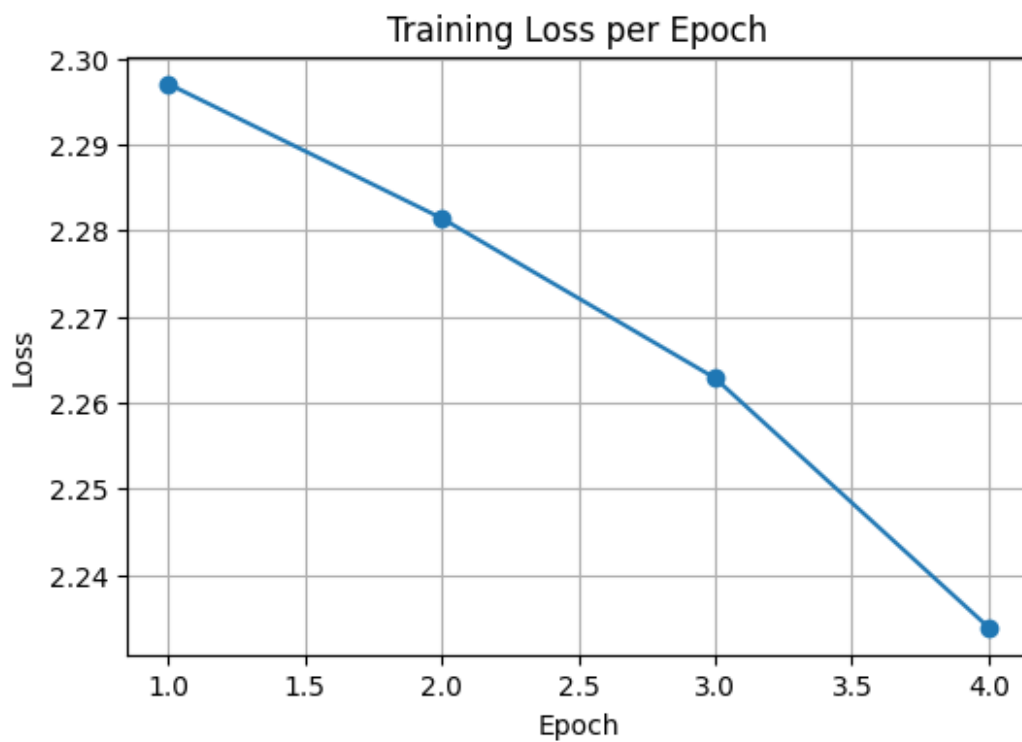
Epoch 2/25, Loss: 2.2814



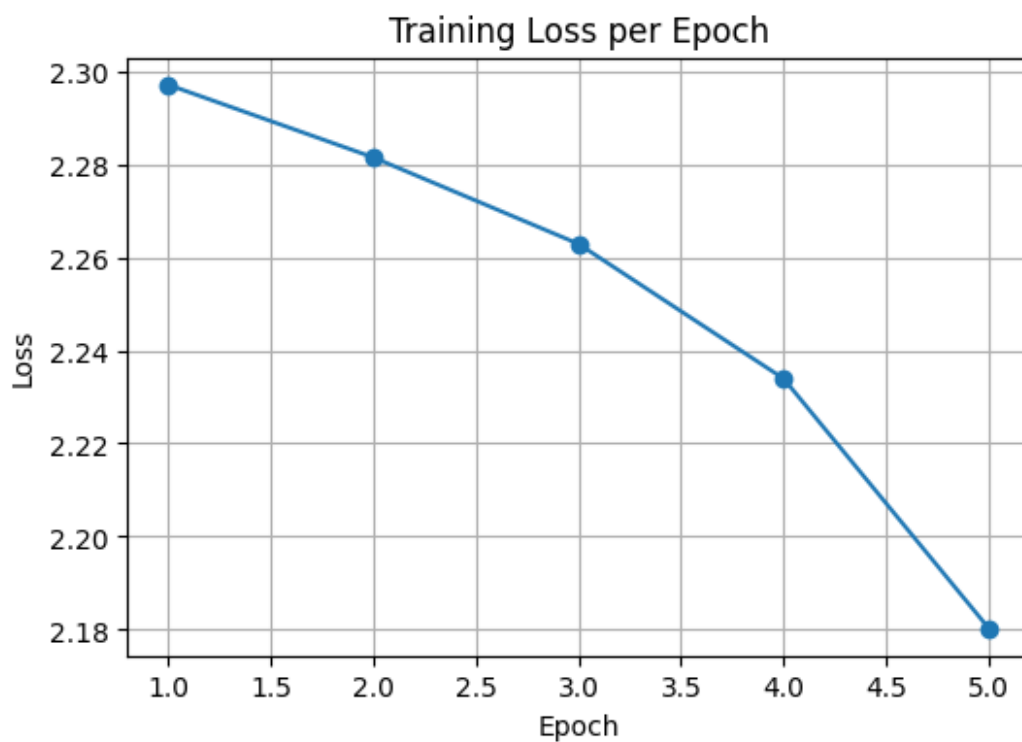
Epoch 3/25, Loss: 2.2629



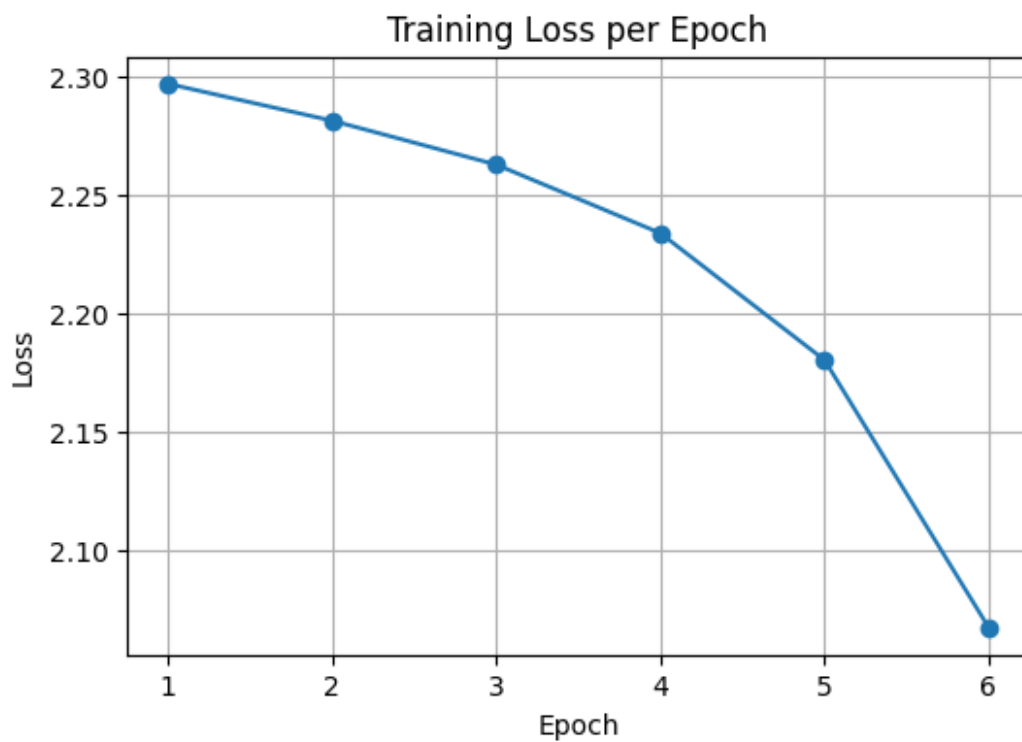
Epoch 4/25, Loss: 2.2339



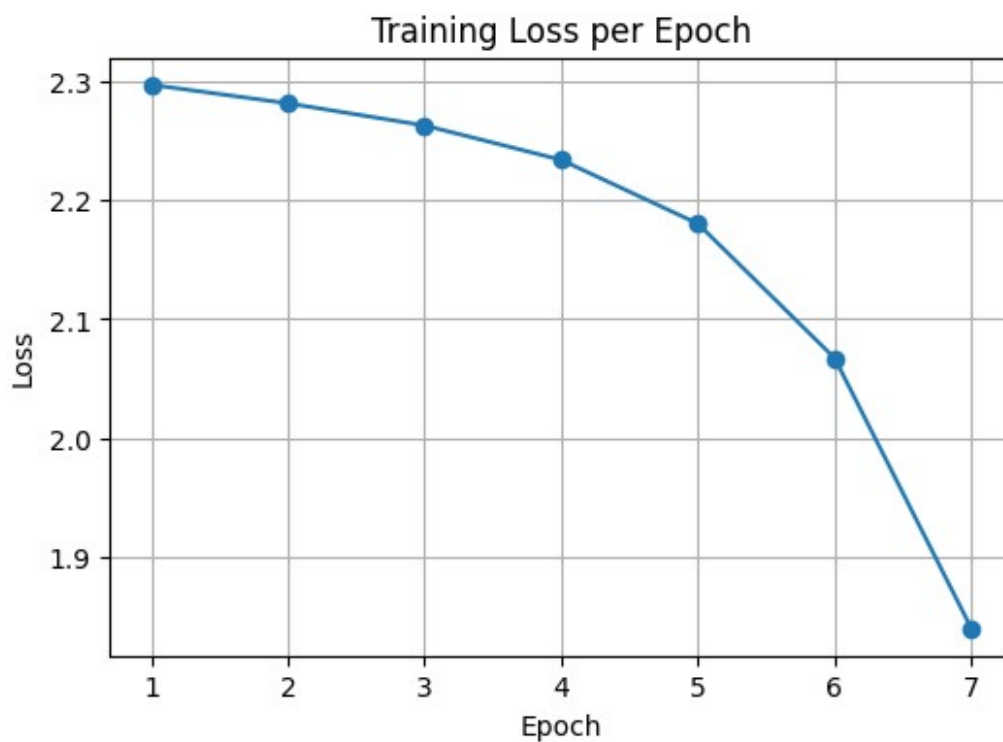
Epoch 5/25, Loss: 2.1803



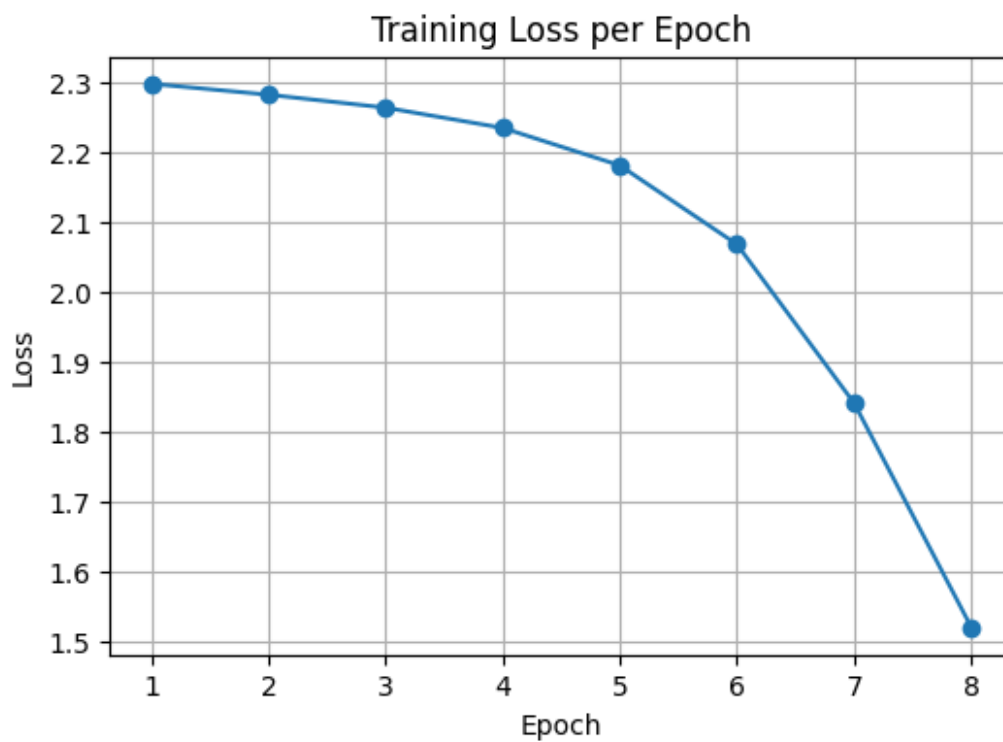
Epoch 6/25, Loss: 2.0673



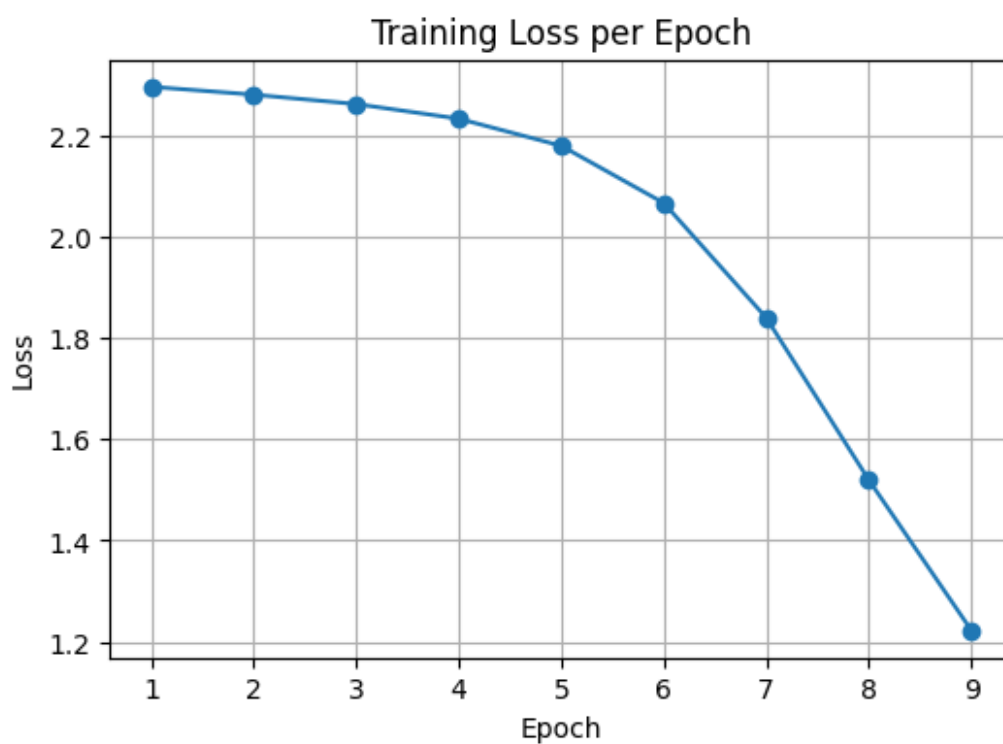
Epoch 7/25, Loss: 1.8402



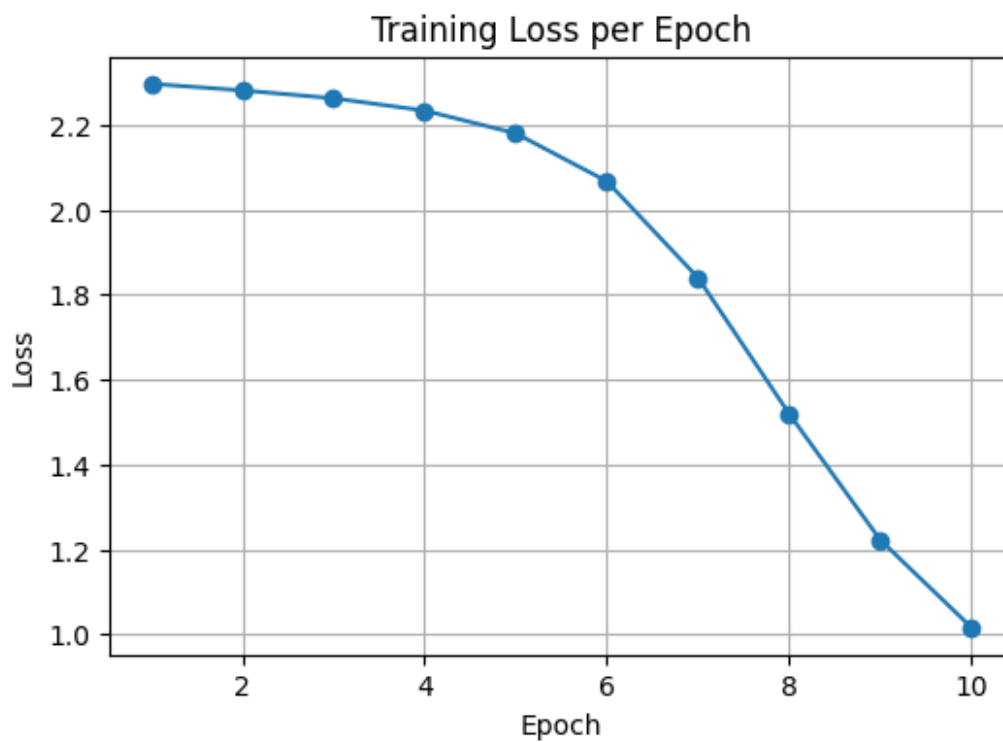
Epoch 8/25, Loss: 1.5197



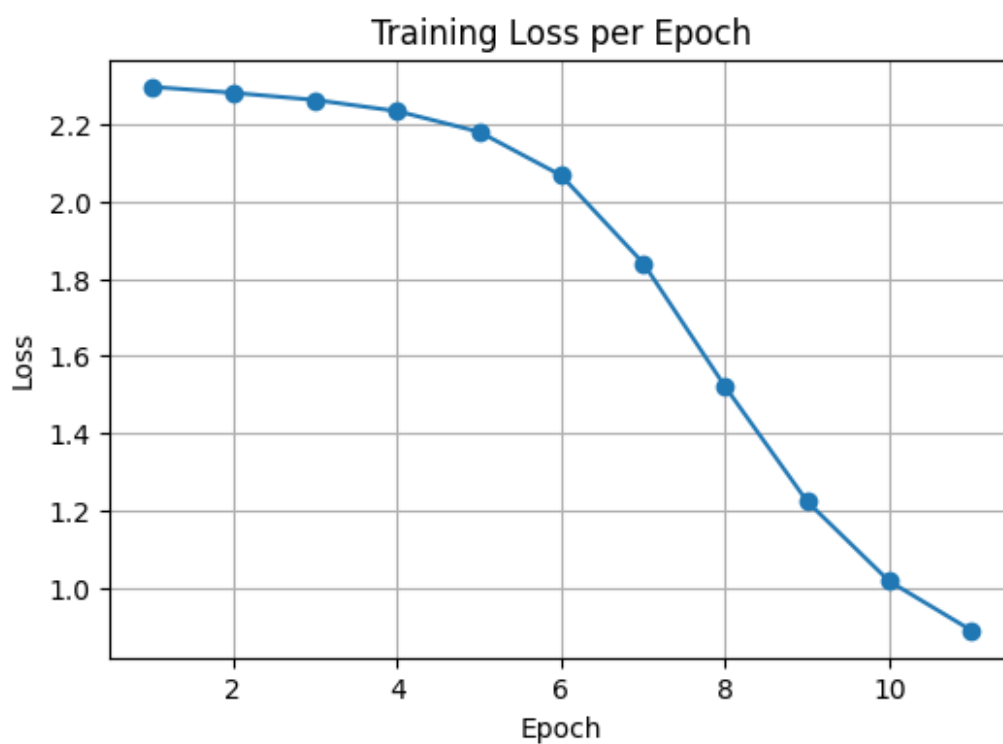
Epoch 9/25, Loss: 1.2224



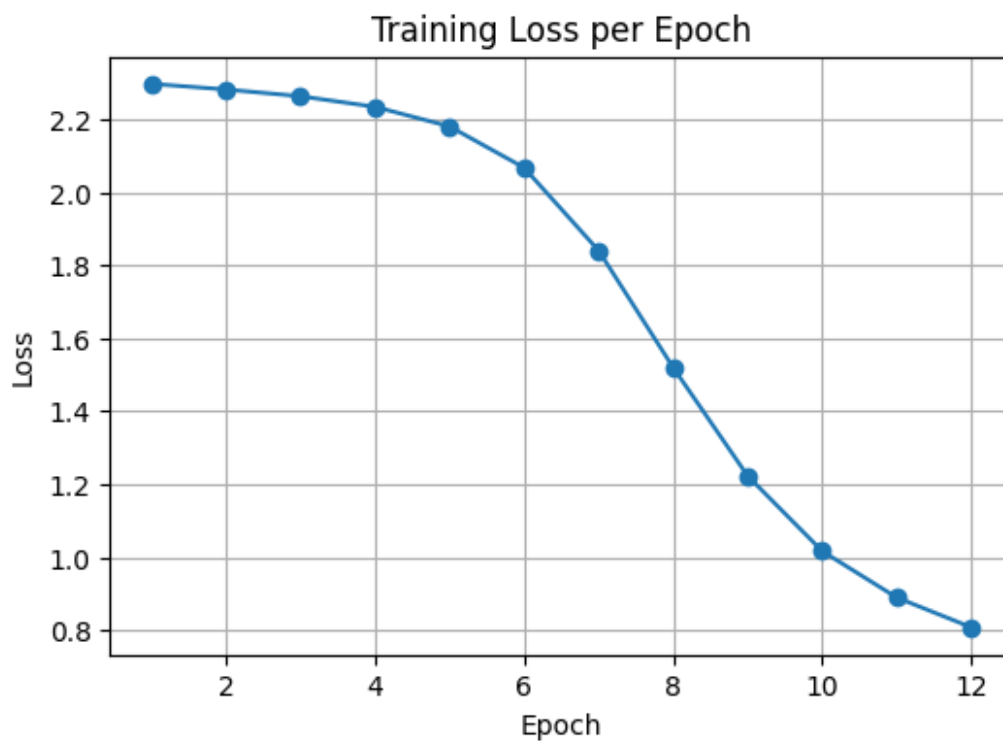
Epoch 10/25, Loss: 1.0162



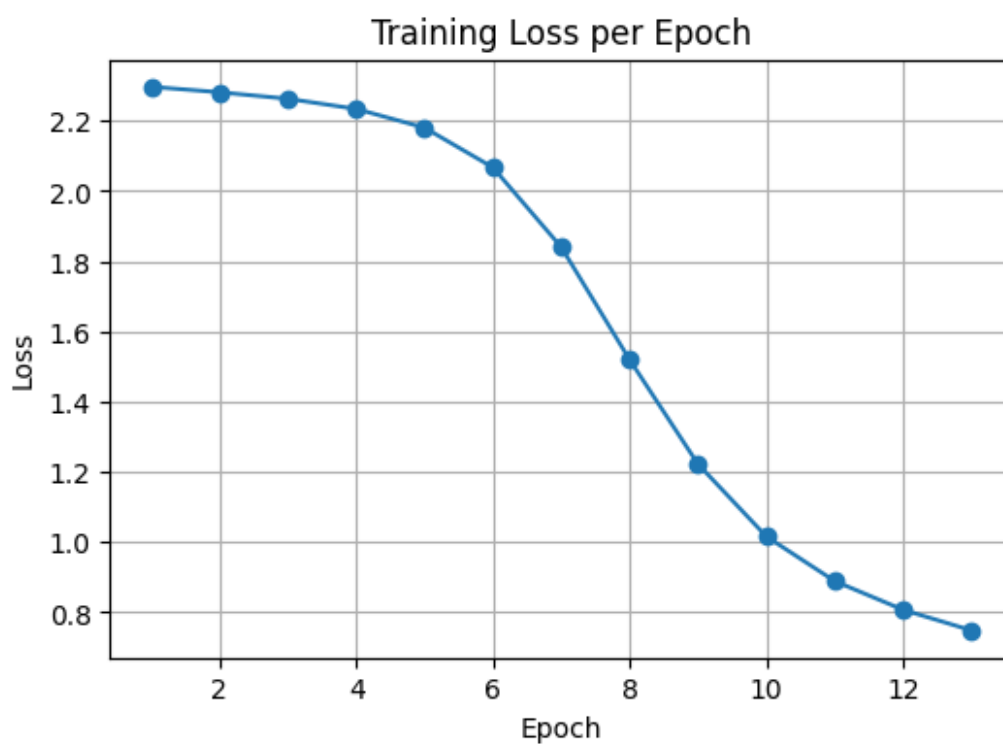
Epoch 11/25, Loss: 0.8891



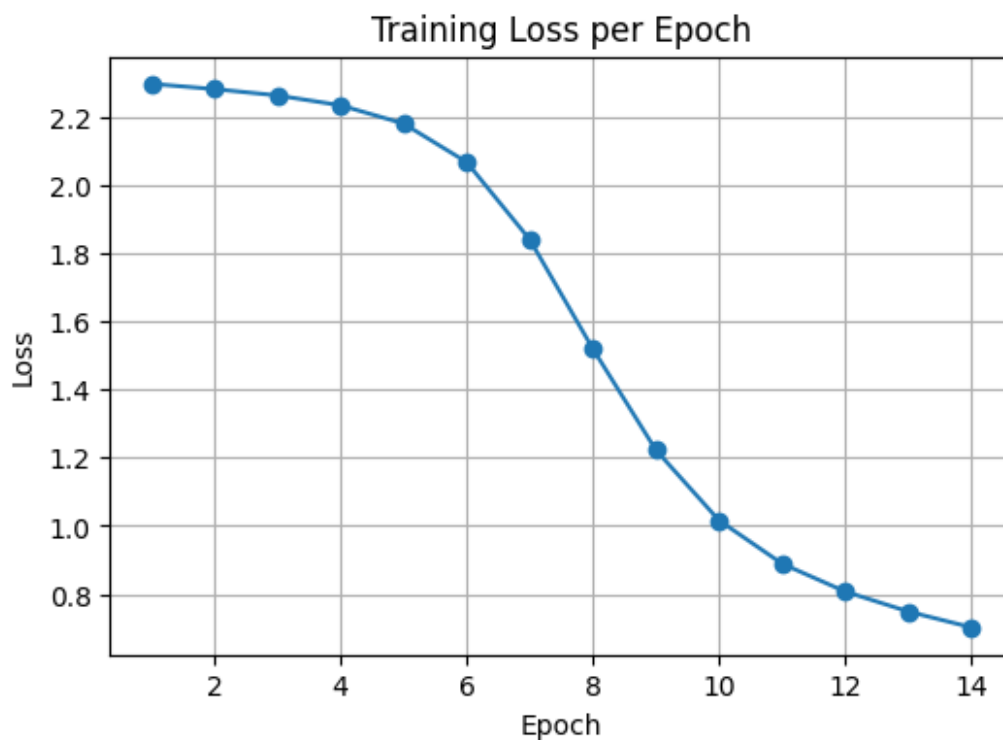
Epoch 12/25, Loss: 0.8073



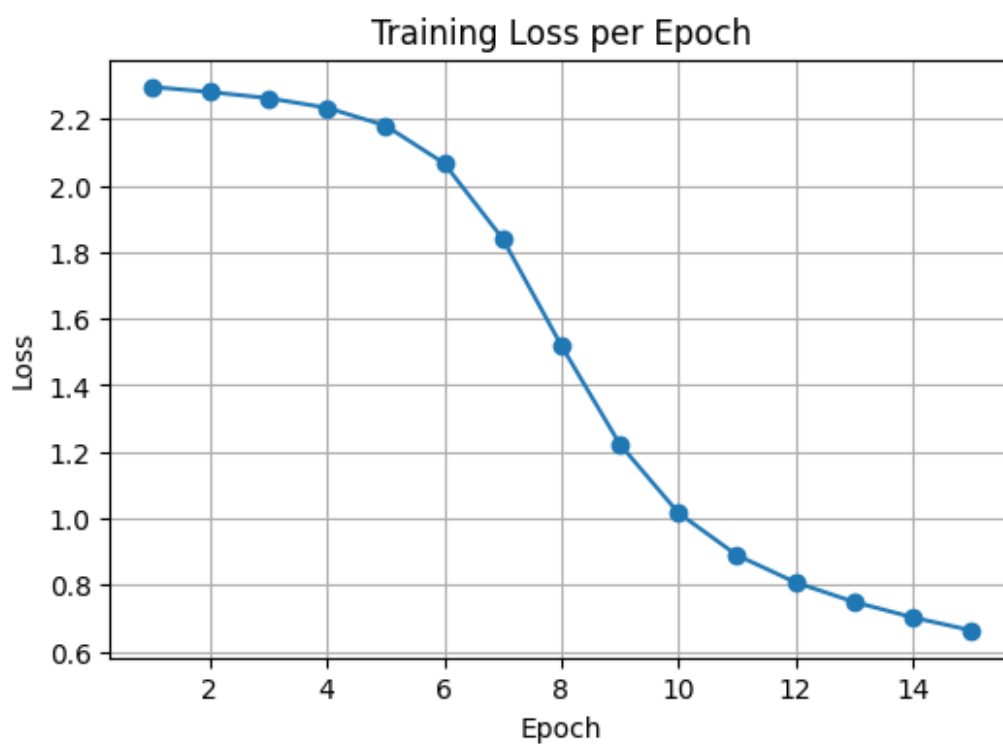
Epoch 13/25, Loss: 0.7487



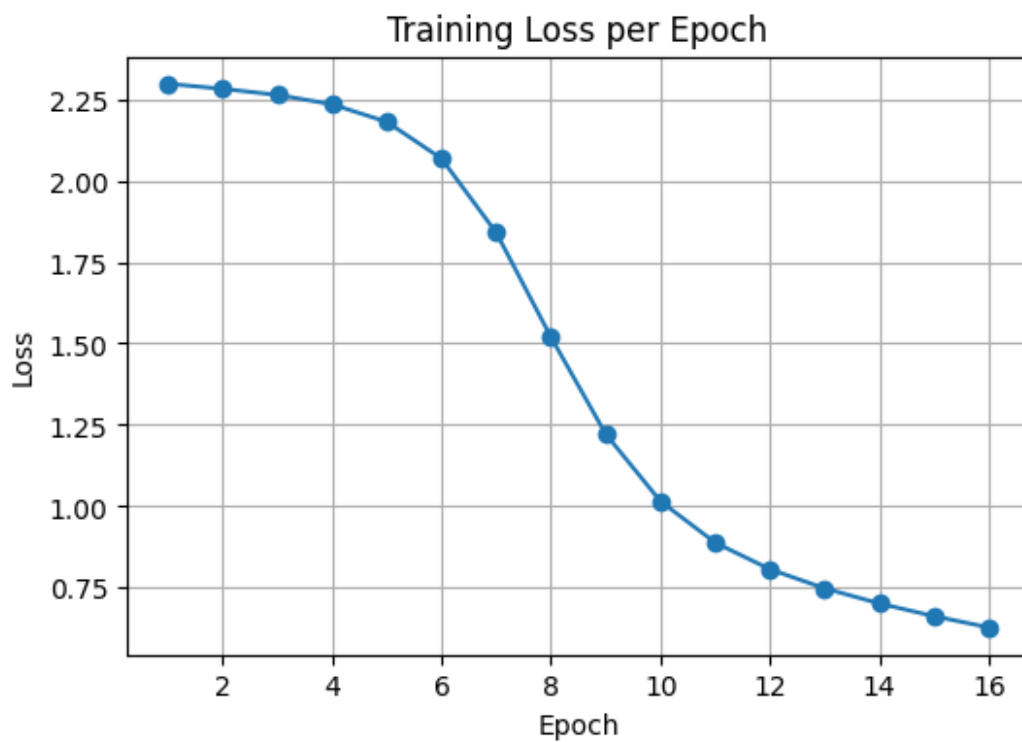
Epoch 14/25, Loss: 0.7022



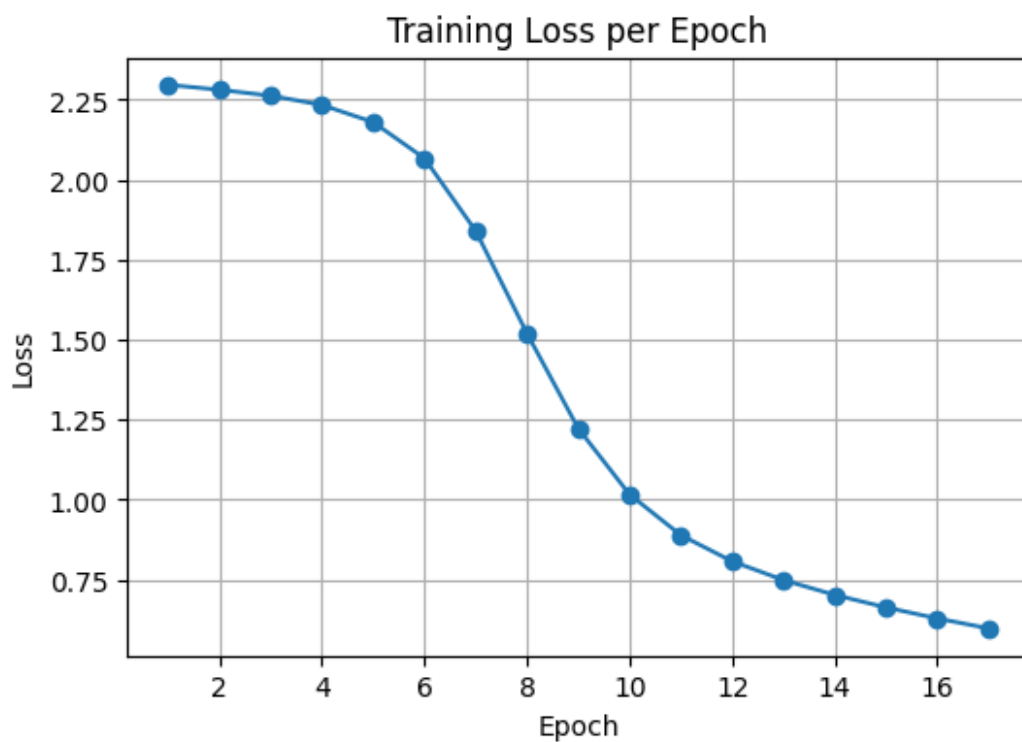
Epoch 15/25, Loss: 0.6632



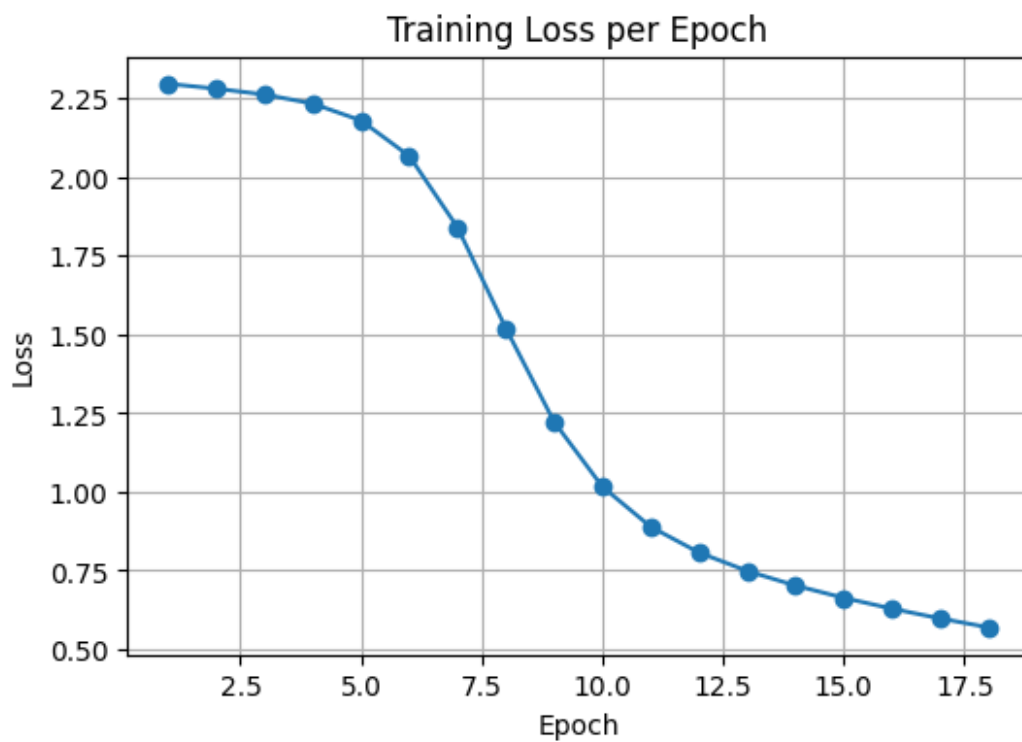
Epoch 16/25, Loss: 0.6288



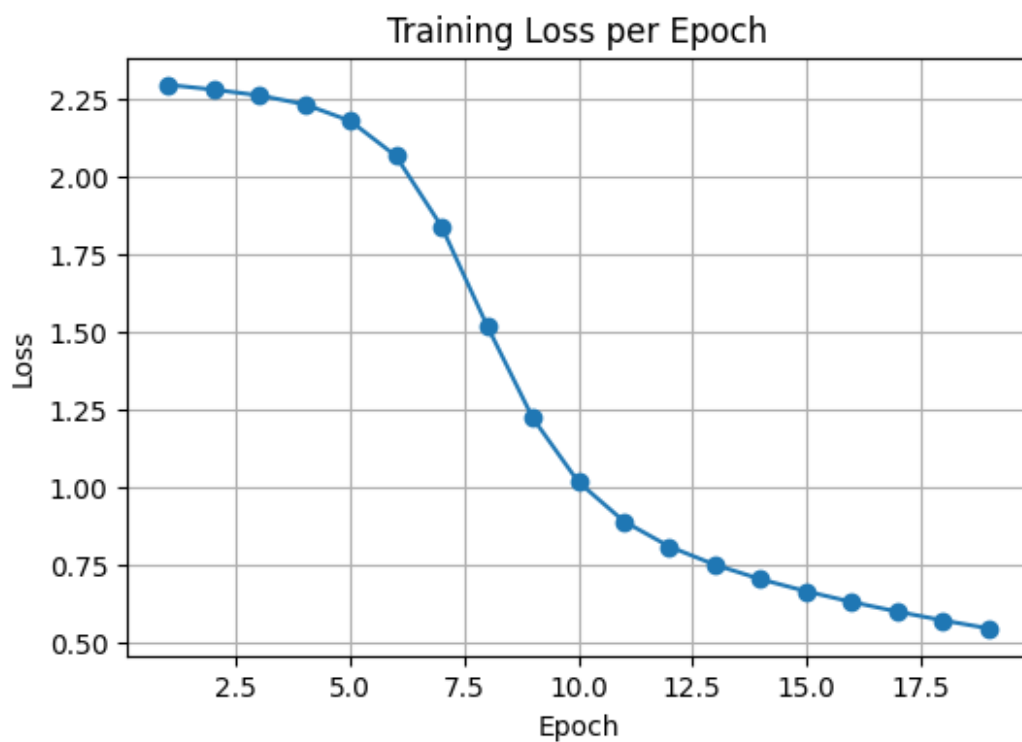
Epoch 17/25, Loss: 0.5977



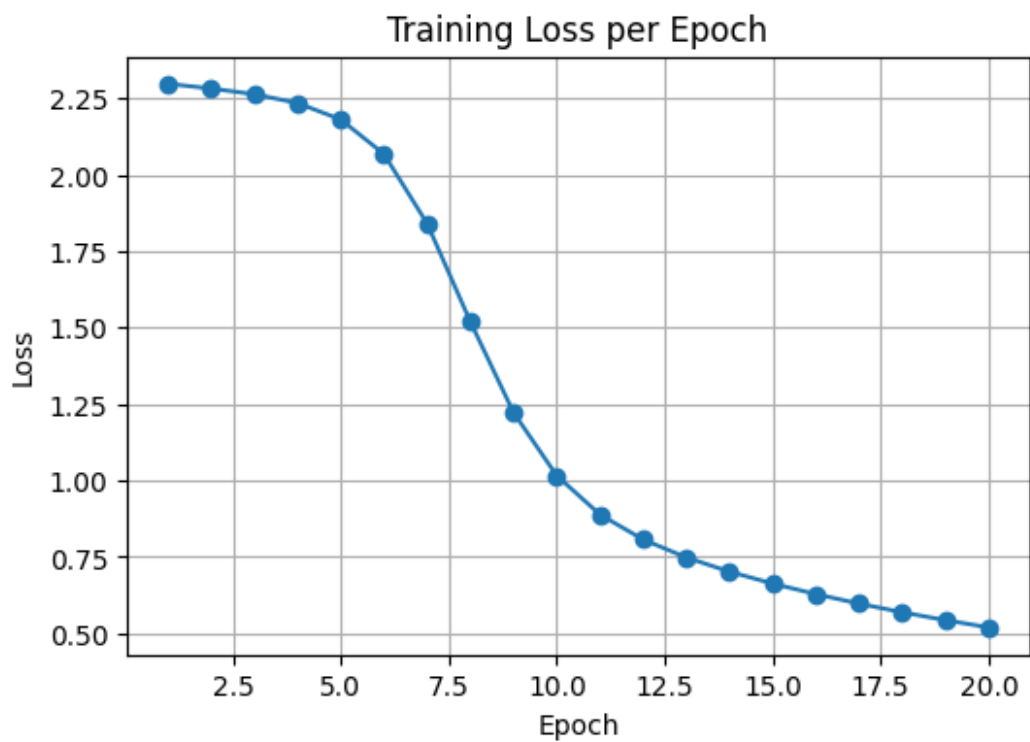
Epoch 18/25, Loss: 0.5694



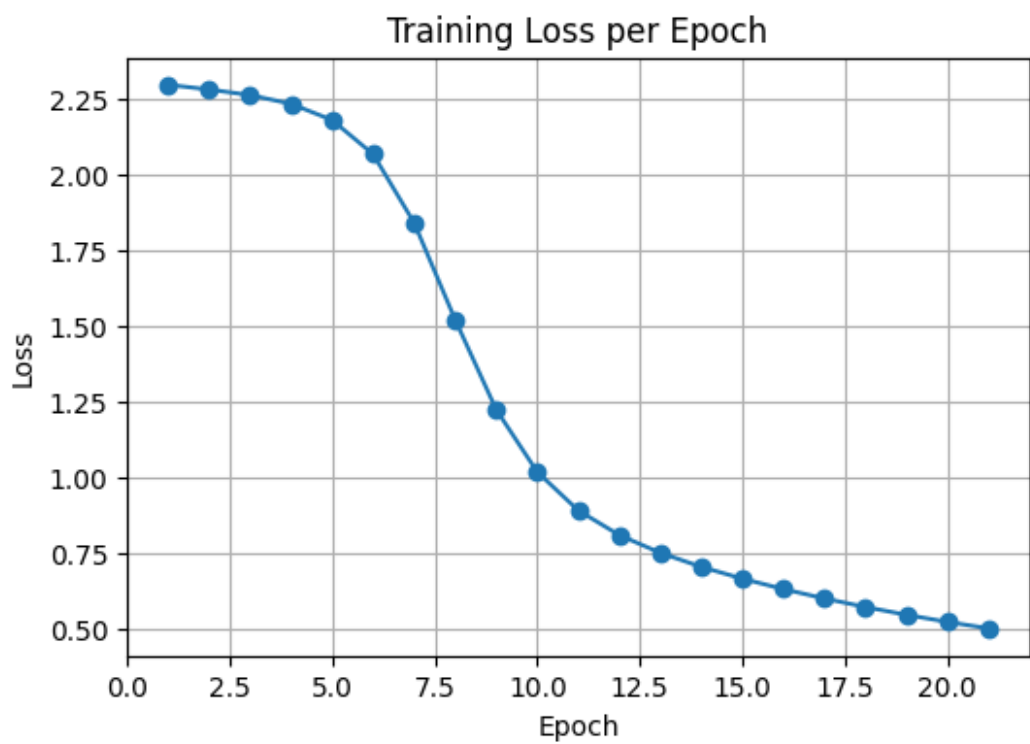
Epoch 19/25, Loss: 0.5437



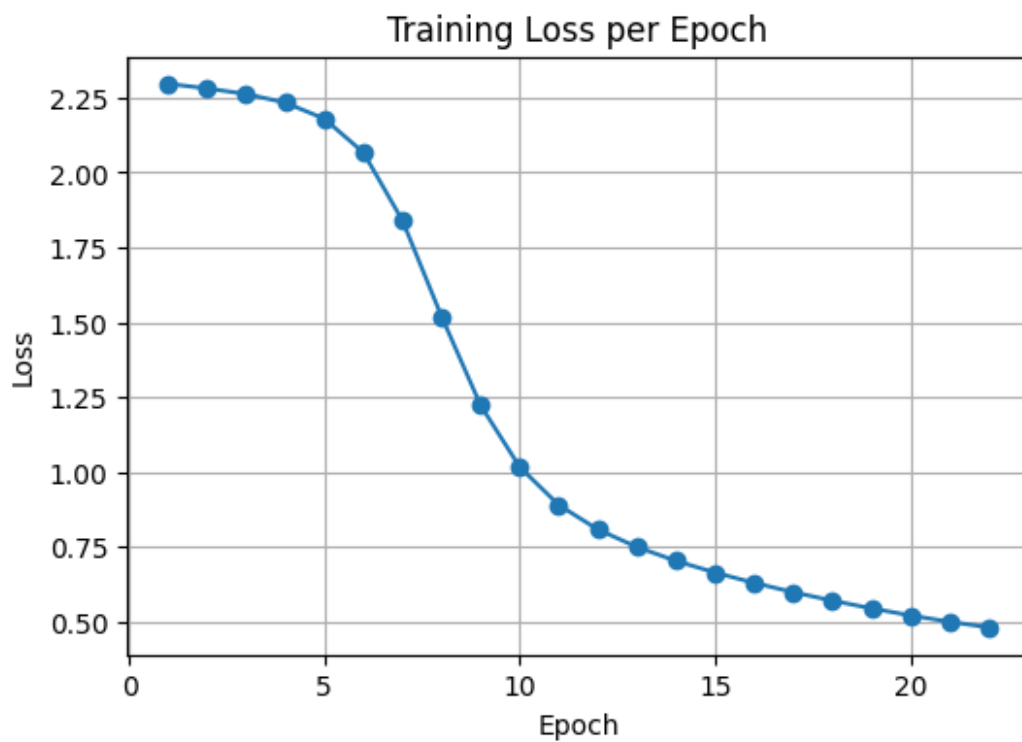
Epoch 20/25, Loss: 0.5201



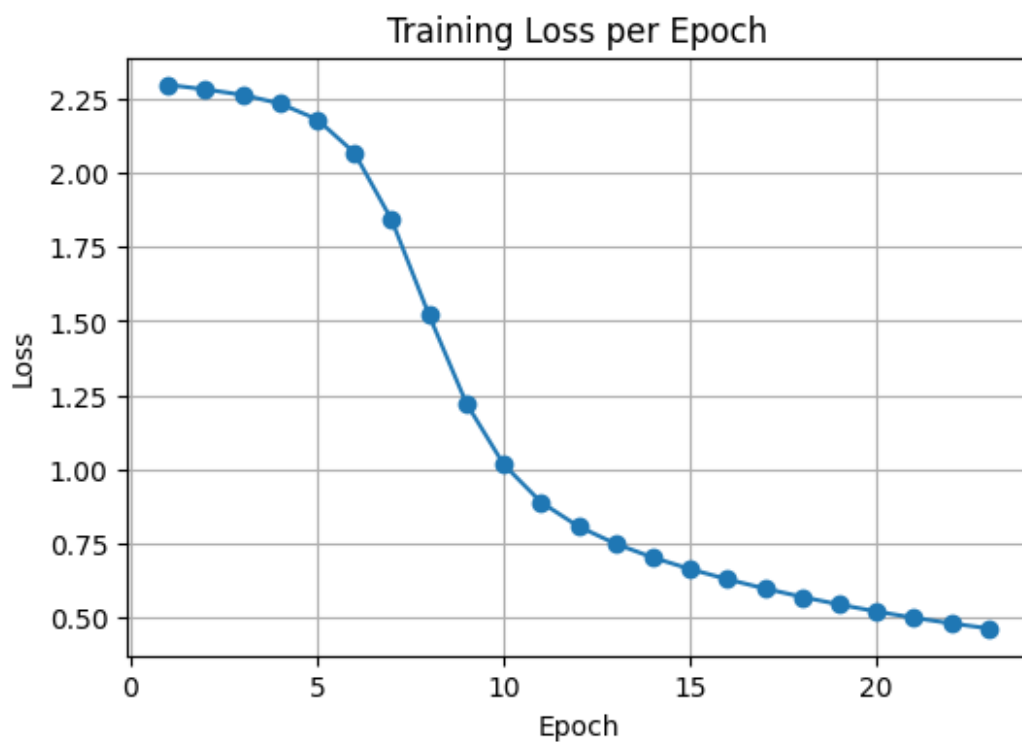
Epoch 21/25, Loss: 0.4989



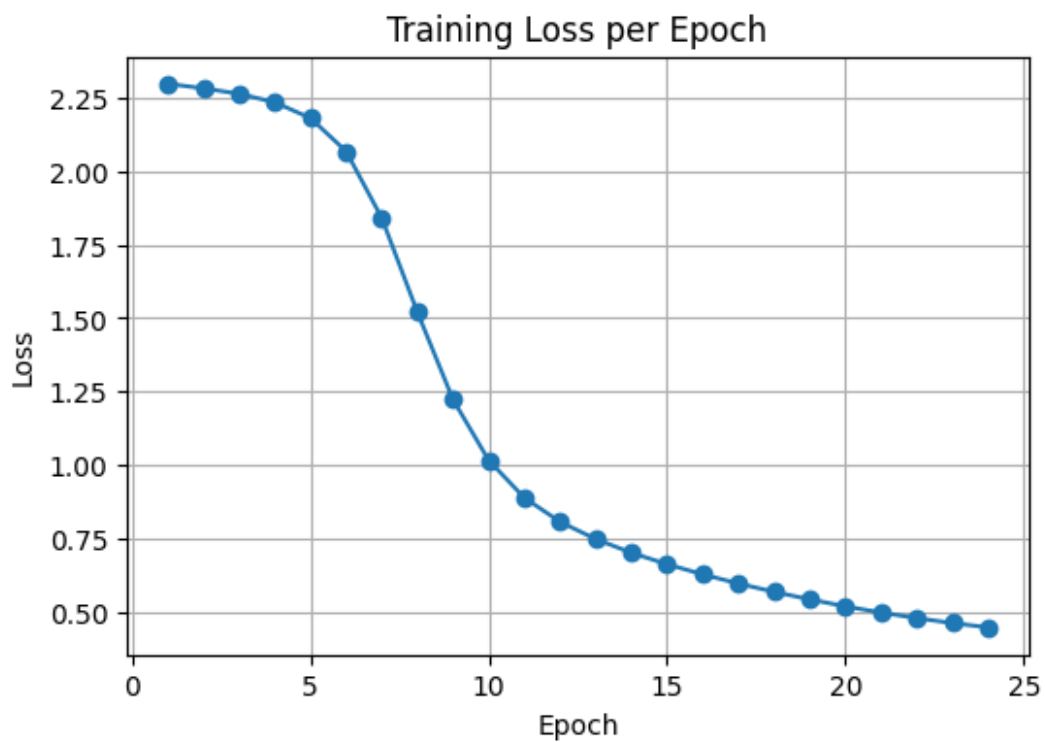
Epoch 22/25, Loss: 0.4801



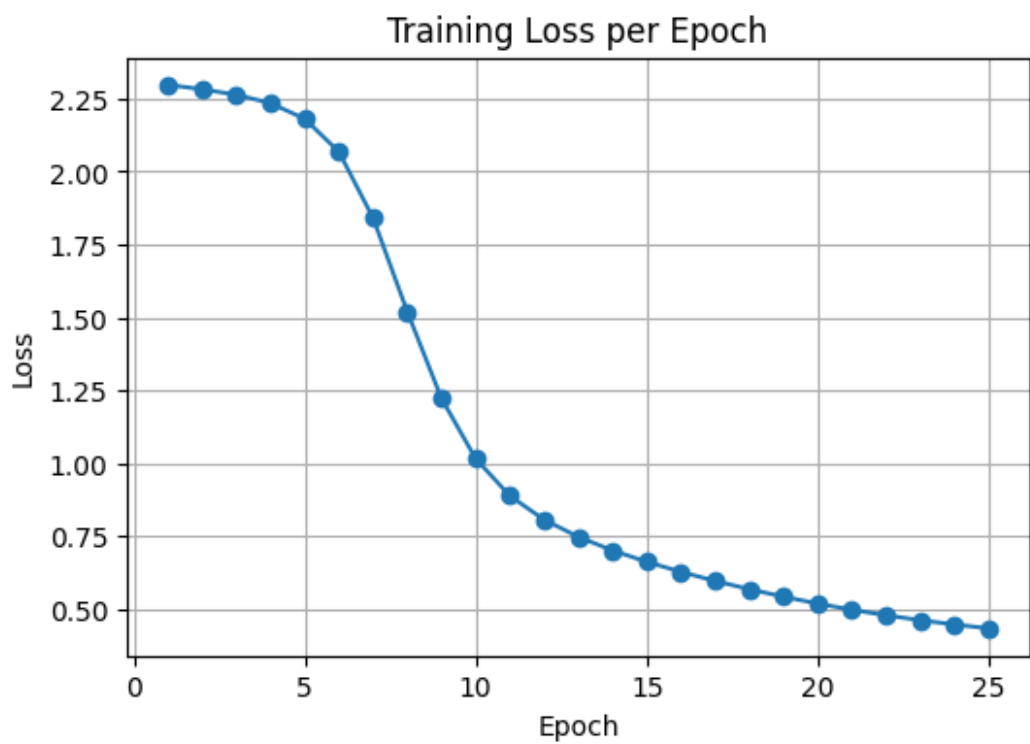
Epoch 23/25, Loss: 0.4632



Epoch 24/25, Loss: 0.4486

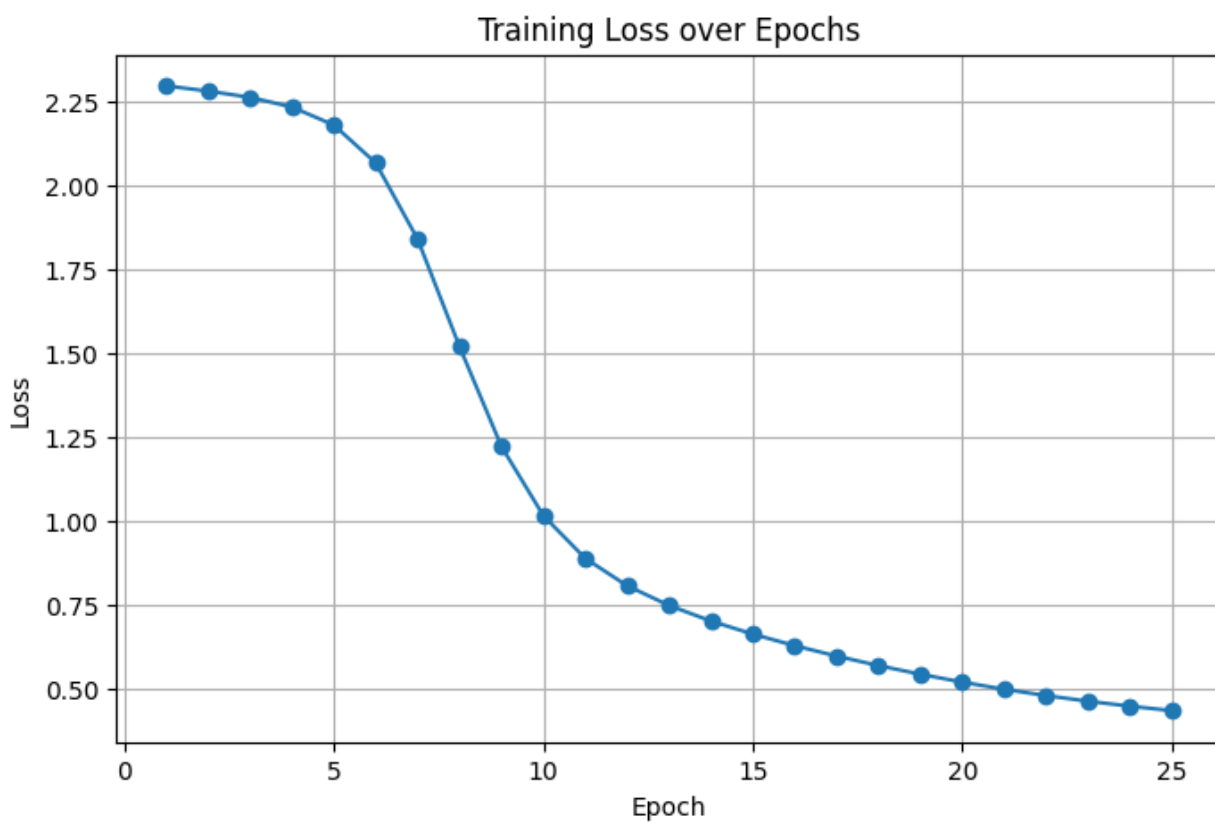


Epoch 25/25, Loss: 0.4352





```
plot_training_curves(loss_history_sigmoid_5)
```




```
X_test_np = mnist_testset.data.numpy().reshape(-1, 28*28) / 255.0
y_test_np = mnist_testset.targets.numpy()
```

```
print("Test Set Evaluation for 25 epochs")
evaluate(model_5, X_test_np, y_test_np)
```

Test Set Evaluation for 25 epochs

Accuracy: 87.7500%

Confusion Matrix:

```
[[ 948    0    2    3    0   18    6    2    1    0]
 [    0 1094    1    6    1    3    3    1   26    0]
 [   16   13  861   41   21    4   21    9   38    8]
 [    4    1   35  849    1   58    0   24   32    6]
 [    1    2    6    0  885    1   12    1    9   65]
 [   14    3    8   49   11  723   26   14   36    8]
 [   25    2   12    0   24   22  869    0    4    0]
 [    7   16   20   11    5    1    0  907    5   56]
 [    5    5   12   28   16   46   19    6  818   19]
 [   10    3    3    6   93   16    1   42   14  821]]
```

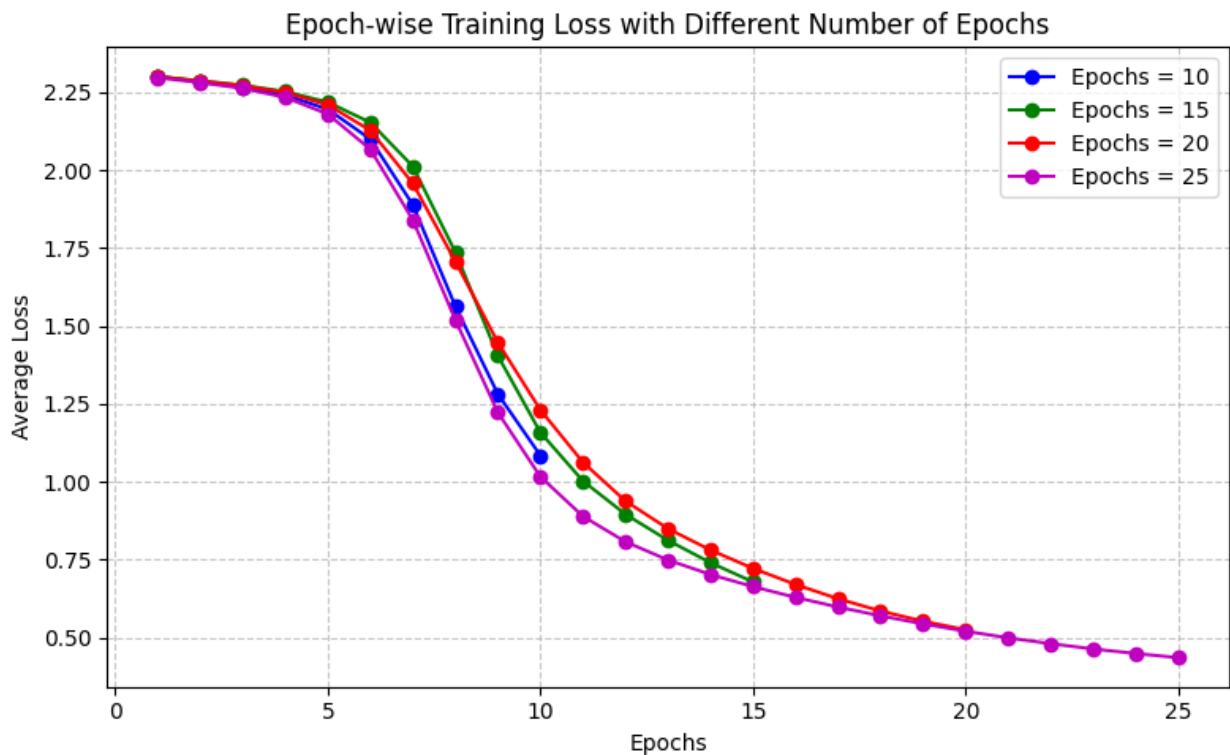
Classification Report:

	precision	recall	f1-score	support
0	0.92	0.97	0.94	980
1	0.96	0.96	0.96	1135
2	0.90	0.83	0.86	1032
3	0.85	0.84	0.85	1010
4	0.84	0.90	0.87	982
5	0.81	0.81	0.81	892
6	0.91	0.91	0.91	958
7	0.90	0.88	0.89	1028
8	0.83	0.84	0.84	974
9	0.84	0.81	0.82	1009
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000

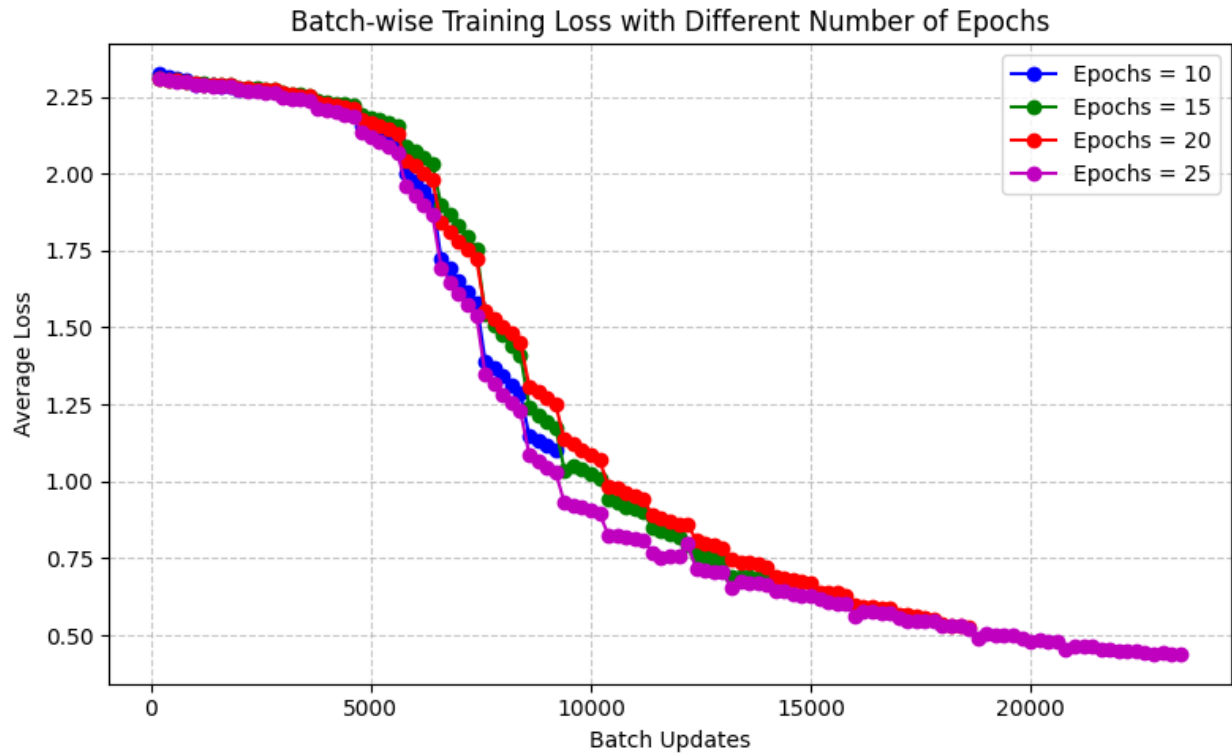
```
np.float64(0.8775)
```

```
plt.figure(figsize=(8,5))
plt.plot(range(1, len(loss_history_sigmoid_3)+1),
loss_history_sigmoid_3, marker='o', color='b', label="Epochs = 10")
plt.plot(range(1, len(loss_history)+1), loss_history, marker='o',
color='g', label="Epochs = 15")
plt.plot(range(1, len(loss_history_sigmoid_4)+1),
loss_history_sigmoid_4, marker='o', color='r', label="Epochs = 20")
plt.plot(range(1, len(loss_history_sigmoid_5)+1),
loss_history_sigmoid_5, marker='o', color='m', label="Epochs = 25")
```

```
plt.title("Epoch-wise Training Loss with Different Number of Epochs")
plt.xlabel("Epochs")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(8,5))
plt.plot(iter_sigmoid_3, iter_loss_sigmoid_3, marker='o', color='b',
label="Epochs = 10")
plt.plot(iter_tanh, iter_loss, marker='o', color='g', label="Epochs = 15")
plt.plot(iter_sigmoid_4, iter_loss_sigmoid_4, marker='o', color='r',
label="Epochs = 20")
plt.plot(iter_sigmoid_5, iter_loss_sigmoid_5, marker='o', color='m',
label="Epochs = 25")
plt.title("Batch-wise Training Loss with Different Number of Epochs")
plt.xlabel("Batch Updates")
plt.ylabel("Average Loss")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()
```



Observation and Conclusions based on variation in number of epochs

- Test accuracy for 10 epochs = **69.76%**
- Test accuracy for 15 epochs = **82.30%**
- Test accuracy for 20 epochs = **85.39%**
- Test accuracy for 25 epochs = **87.75%**

Increasing the number of epochs gives more time for the network to learn and hence positively helps in improving the accuracy in testing phase.

The overall training loss per epoch is roughly the same and independent of number of epochs, that is the training at the nth epoch is roughly same in each case of having different number of epochs for the overall training